

# C++ AMP : Language and Programming Model

**Version 0.99, May 2012**

© 2012 Microsoft Corporation. All rights reserved.

This specification reflects input from NVIDIA Corporation (Nvidia) and Advanced Micro Devices, Inc. (AMD).

**Copyright License.** Microsoft grants you a license under its copyrights in the specification to (a) make copies of the specification to develop your implementation of the specification, and (b) distribute portions of the specification in your implementation or your documentation of your implementation.

**Patent Notice.** Microsoft provides you certain patent rights for implementations of this specification under the terms of Microsoft's Community Promise, available at <http://www.microsoft.com/openspecifications/en/us/programs/community-promise/default.aspx>.

**THIS SPECIFICATION IS PROVIDED "AS IS." MICROSOFT MAY CHANGE THIS SPECIFICATION OR ITS OWN IMPLEMENTATIONS AT ANY TIME AND WITHOUT NOTICE. MICROSOFT MAKES NO REPRESENTATIONS OR WARRANTIES, EXPRESS, IMPLIED, OR STATUTORY, (1) AS TO THE INFORMATION IN THIS SPECIFICATION, INCLUDING ANY WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NON-INFRINGEMENT, OR TITLE; OR (2) THAT THE IMPLEMENTATION OF SUCH CONTENTS WILL NOT INFRINGE ANY THIRD PARTY PATENTS OR OTHER RIGHTS.**

## ABSTRACT

C++ AMP (Accelerated Massive Parallelism) is a native programming model that contains elements that span the C++ programming language and its runtime library. It provides an easy way to write programs that compile and execute on data-parallel hardware, such as graphics cards (GPUs).

The syntactic changes introduced by C++ AMP are minimal, but additional restrictions are enforced to reflect the limitations of data parallel hardware.

Data parallel algorithms are supported by the introduction of multi-dimensional array types, array operations on those types, indexing, asynchronous memory transfer, shared memory, synchronization and tiling/partitioning techniques.

<b>1</b>	<b>Overview .....</b>	<b>1</b>
1.1	Conformance .....	1
1.2	Definitions.....	2
1.3	Error Model.....	4
1.4	Programming Model .....	5
<b>2</b>	<b>C++ Language Extensions for Accelerated Computing .....</b>	<b>5</b>
2.1	Syntax.....	6
2.1.1	Function Declarator Syntax.....	6
2.1.2	Lambda Expression Syntax.....	7
2.1.3	Type Specifiers .....	7
2.2	Meaning of Restriction Specifiers .....	7
2.2.1	Function Definitions .....	8
2.2.2	Constructors and Destructors .....	8
2.2.3	Lambda Expressions.....	9
2.3	Expressions Involving Restricted Functions .....	9
2.3.1	Function pointer conversions .....	9
2.3.2	Function Overloading.....	9
2.3.2.1	Overload Resolution.....	11
2.3.2.2	Name Hiding .....	12
2.3.3	Casting.....	12
2.4	amp Restriction Modifier.....	12
2.4.1	Restrictions on Types .....	12
2.4.1.1	Type Qualifiers .....	12
2.4.1.2	Fundamental Types .....	13
2.4.1.3	Compound Types .....	13
2.4.2	Restrictions on Function Declarators .....	13
2.4.3	Restrictions on Function Scopes .....	14

2.4.3.1	Literals .....	14
2.4.3.2	Primary Expressions (C++11 5.1) .....	14
2.4.3.3	Lambda Expressions .....	14
2.4.3.4	Function Calls (C++11 5.2.2) .....	14
2.4.3.5	Local Declarations .....	14
2.4.3.5.1	tile_static Variables .....	14
2.4.3.6	Type-Casting Restrictions .....	15
2.4.3.7	Miscellaneous Restrictions .....	15
<b>3</b>	<b>Device Modeling .....</b>	<b>15</b>
3.1	The concept of a compute accelerator .....	15
3.2	accelerator .....	15
3.2.1	Default Accelerator .....	16
3.2.2	Synopsis .....	16
3.2.3	Static Members .....	17
3.2.4	Constructors .....	17
3.2.5	Members .....	18
3.2.6	Properties .....	20
3.3	accelerator_view .....	20
3.3.1	Synopsis .....	20
3.3.2	Queuing Mode .....	21
3.3.3	Constructors .....	21
3.3.4	Members .....	21
3.4	Device enumeration and selection API .....	23
3.4.1	Synopsis .....	23
<b>4</b>	<b>Basic Data Elements .....</b>	<b>24</b>
4.1	index<N> .....	24
4.1.1	Synopsis .....	24
4.1.2	Constructors .....	26
4.1.3	Members .....	26
4.1.4	Operators .....	26
4.2	extent<N> .....	28
4.2.1	Synopsis .....	28
4.2.2	Constructors .....	30
4.2.3	Members .....	30
4.2.4	Operators .....	31
4.3	tiled_extent<D0,D1,D2> .....	32
4.3.1	Synopsis .....	32
4.3.2	Constructors .....	34
4.3.3	Members .....	34
4.3.4	Operators .....	35

4.4	tiled_index<D0,D1,D2> .....	35
4.4.1	Synopsis .....	36
4.4.2	Constructors.....	38
4.4.3	Members.....	39
4.5	tile_barrier .....	39
4.5.1	Synopsis .....	40
4.5.2	Constructors.....	40
4.5.3	Members.....	40
4.5.4	Other Memory Fences and Barriers.....	40
4.6	completion_future .....	41
4.6.1	Synopsis .....	41
4.6.2	Constructors.....	42
4.6.3	Members.....	42
<b>5</b>	<b>Data Containers .....</b>	<b>43</b>
5.1	array<T,N> .....	43
5.1.1	Synopsis .....	44
5.1.2	Constructors.....	51
5.1.2.1	Staging Array Constructors.....	53
5.1.3	Members.....	55
5.1.4	Indexing.....	56
5.1.5	View Operations .....	57
5.2	array_view<T,N> .....	58
5.2.1	Synopsis .....	59
5.2.1.1	array_view<T,N> .....	59
5.2.1.2	array_view<const T,N>.....	63
5.2.2	Constructors.....	66
5.2.3	Members.....	67
5.2.4	Indexing.....	68
5.2.5	View Operations .....	69
5.3	Copying Data.....	70
5.3.1	Synopsis .....	71
5.3.2	Copying between array and array_view .....	72
5.3.3	Copying from standard containers to arrays or array_views.....	73
5.3.4	Copying from arrays or array_views to standard containers.....	74
<b>6</b>	<b>Atomic Operations.....</b>	<b>75</b>
6.1	Synopsis .....	75
6.2	Atomically Exchanging Values.....	75
6.3	Atomically Applying an Integer Numerical Operation .....	76
<b>7</b>	<b>Launching Computations: parallel_for_each .....</b>	<b>77</b>
7.1	Capturing Data in the Kernel Function Object .....	80

7.2	Exception Behaviour .....	80
<b>8</b>	<b>Correctly Synchronized C++ AMP Programs .....</b>	<b>80</b>
8.1	Concurrency of sibling threads launched by a <code>parallel_for_each</code> call.....	81
8.1.1	Correct usage of tile barriers .....	82
8.1.2	Establishing order between operations of concurrent <code>parallel_for_each</code> threads .....	83
8.1.2.1	Barrier-incorrect programs .....	84
8.1.2.2	Compatible memory operations .....	84
8.1.2.3	Concurrent memory operations.....	84
8.1.2.4	Racy programs.....	85
8.1.2.5	Race-free programs.....	85
8.2	Cumulative effects of a <code>parallel_for_each</code> call .....	86
8.3	Effects of copy and <code>copy_async</code> operations.....	88
8.4	Effects of <code>array_view::synchronize</code> , <code>synchronize_async</code> and <code>refresh</code> functions.....	88
<b>9</b>	<b>Math Functions .....</b>	<b>89</b>
9.1	<code>fast_math</code> .....	89
9.2	<code>precise_math</code> .....	92
<b>10</b>	<b>Graphics (Optional).....</b>	<b>98</b>
10.1	<code>texture&lt;T,N&gt;</code> .....	98
10.1.1	Synopsis .....	99
10.1.2	Introduced typedefs .....	101
10.1.3	Constructing an uninitialized texture .....	101
10.1.4	Constructing a texture from a host side iterator .....	102
10.1.5	Constructing a texture from a host-side data source .....	103
10.1.6	Constructing a texture by cloning another .....	104
10.1.7	Assignment operator .....	104
10.1.8	Copying textures.....	104
10.1.9	Moving textures.....	105
10.1.10	Querying texture's physical characteristics .....	105
10.1.11	Querying texture's logical dimensions .....	105
10.1.12	Querying the <code>accelerator_view</code> where the texture resides.....	105
10.1.13	Reading and writing textures .....	106
10.1.14	Global texture copy functions .....	106
10.1.14.1	Global async texture copy functions.....	107
10.1.15	Direct3d Interop Functions.....	107
10.2	<code>writeonly_texture_view&lt;T,N&gt;</code> .....	107
10.2.1	Synopsis .....	108
10.2.2	Introduced typedefs .....	108
10.2.3	Construct a writeonly view over a texture .....	108
10.2.4	Copy constructors and assignment operators.....	108

10.2.5	Destructor .....	109
10.2.6	Querying underlying texture's physical characteristics .....	109
10.2.7	Querying the underlying texture's accelerator_view .....	109
10.2.7.1	Querying underlying texture's logical dimensions (through a view) .....	109
10.2.7.2	Writing a write-only texture view .....	109
10.2.8	Global writeonly_texture_view copy functions.....	110
10.2.8.1	Global async writeonly_texture_view copy functions .....	110
10.2.9	Direct3d Interop Functions.....	110
10.3	norm and unorm .....	110
10.3.1	Synopsis .....	110
10.3.2	Constructors and Assignment.....	112
10.3.3	Operators.....	112
10.4	Short Vector Types.....	112
10.4.1	Synopsis .....	113
10.4.2	Constructors .....	114
10.4.2.1	Constructors from components .....	114
10.4.2.2	Explicit conversion constructors .....	115
10.4.3	Component Access (Swizzling) .....	115
10.4.3.1	Single-component access .....	115
10.4.3.2	Two-component access.....	116
10.4.3.3	Three-component access .....	116
10.4.3.4	Four-component access .....	117
10.5	Template class short_vector_traits .....	117
10.5.1	Synopsis .....	117
10.5.2	Typedefs .....	120
10.5.3	Members .....	121
<b>11</b>	<b>D3D interoperability (Optional) .....</b>	<b>122</b>
<b>12</b>	<b>Error Handling.....</b>	<b>124</b>
12.1	static_assert .....	124
12.2	Runtime errors .....	124
12.2.1	runtime_exception .....	124
12.2.1.1	Specific Runtime Exceptions .....	125
12.2.2	out_of_memory.....	125
12.2.3	invalid_compute_domain.....	126
12.2.4	unsupported_feature .....	126
12.2.5	accelerator_view_removed.....	127
12.3	Error handling in device code (amp-restricted functions) (Optional) .....	127
<b>13</b>	<b>Appendix: C++ AMP Future Directions (Informative).....</b>	<b>129</b>
13.1	Versioning Restrictions .....	129

13.1.1	<i>auto</i> restriction.....	129
13.1.2	Automatic restriction deduction .....	129
13.1.3	<i>amp</i> Version.....	130
13.2	Projected Evolution of <i>amp</i> -Restricted Code.....	130

# 1 Overview

C++ AMP is a compiler and programming model extension to C++ that enables the acceleration of C++ code on data-parallel hardware.

One example of data-parallel hardware today is the discrete graphics card (GPU), which is becoming increasingly relevant for general purpose parallel computations, in addition to its main function as a graphics accelerator. A GPU is conceptually (and usually physically) separate from the CPU, having discrete memory address space, and incurring high cost for transferring data between CPU and GPU memory. The programmer must carefully balance the cost of this data transfer overhead against the computational acceleration achievable by parallel execution on the device.

Another example of data-parallel hardware is the SIMD vector instruction set, and associated registers, found in all modern processors.

For the remainder of this specification, we shall refer to the data-parallel hardware as the *accelerator*. In the few places where the distinction matters, we shall refer to a GPU or a VectorCPU.

The C++ AMP programming model gives the developer explicit control over all of the above aspects of interaction with the accelerator. The developer may explicitly manage all communication between the CPU and the accelerator, and this communication can be either synchronous or asynchronous. The data parallel computations performed on the accelerator are expressed using high-level abstractions, such as multi-dimensional arrays, high level array manipulation functions, and multi-dimensional indexing operations, all based on a large subset of the C++ programming language.

The programming model contains multiple layers, allowing developers to trade off ease-of-use with maximum performance.

C++ AMP is composed of three broad categories of functionality:

1. C++ language and compiler
  - a. Kernel functions are compiled into code that is specific to the accelerator.
2. Runtime
  - a. The runtime contains a C++ AMP abstraction of lower-level accelerator APIs, as well as support for multiple host threads and processors, and multiple accelerators.
  - b. Asynchronous execution is supported through an eventing model.
3. Programming model
  - a. A set of classes describing the shape and extent of data.
  - b. A set of classes that contain or refer to data used in computations
  - c. A set of functions for copying data to and from accelerators
  - d. A math library
  - e. An atomic library
  - f. A set of miscellaneous intrinsic functions

## 1.1 Conformance

All text in this specification falls into one of the following categories:

- *Informative: shown in this style.*  
Informative text is non-normative; for background information only; not required to be implemented in order to conform to this specification.
- *Microsoft-specific: shown in this style.*



Microsoft-specific text is non-normative; for background information only; not required to be implemented in order to conform to this specification; explains features that are specific to the Microsoft implementation of the C++ AMP programming model. However, implementers are free to implement these feature, or any subset thereof.

- Normative: all text, unless otherwise marked (see previous categories) is normative. Normative text falls into the following two sub-categories:
  - Optional: each section of the specification that falls into this sub-category includes the suffix “(Optional)” in its title. A conforming implementation of C++ AMP may choose to support such features, or not. (Microsoft-specific portions of the text are also Optional.)
  - Required: unless otherwise stated, all Normative text falls into the sub-category of Required. A conforming implementation of C++ AMP *must* support *all* Required features.

Conforming implementations shall provide all normative features and any number of optional features. Implementations may provide additional features so long as these features are exposed in namespaces other than those listed in this specification. Implementation may provide additional language support for amp-restricted functions (section 2.1) by following the rules set forth in section 13.

The programming model utilizes Microsoft’s Visual C++ syntax for *properties*. Any such property shall be considered optional. An implementation is free to use equivalent mechanisms for introducing such properties as long as they provide the same functionality of indirection to a member function as Microsoft’s Visual C++ properties do.

## 1.2 Definitions

This section introduces terms used within the body of this specification.

- **Accelerator**  
A hardware device or capability that enables accelerated computation on data-parallel workloads. Examples include:
  - Graphics Processing Unit, or GPU, other coprocessor, accessible through the PCIe bus.
  - SIMD units of the host node exposed through software emulation of a hardware accelerator.
- **Array**  
A dense N-dimensional data structure.
- **Array View**  
A view into a contiguous piece of memory that adds array-like dimensionality.
- **Compressed texture format.**  
A format that divides a texture into blocks that allow the texture to be reduced in size by a fixed ratio; typically 4:1 or 6:1. Compressed textures are useful when perfect image/texel fidelity is not necessary but where minimizing memory storage and bandwidth are critical to application performance.
- **Extent**  
A vector of integers that describes lengths of N-dimensional array-like objects.
- **Global memory**  
On a GPU, global memory is the main off-chip memory store,  
*Informative: Typically, on current-generation GPUs, global memory is implemented in DRAM, with access times of 400-1000 cycles; the GPU clock speed is around 1 Ghz; and is non-cached. Global memory is accessed in a coalesced pattern with a granularity of 128 bytes, so when accessing 4 bytes of global memory, 32 successive threads need to read the 32 successive 4-byte addresses, to be fully coalesced.*

*Informative: The memory space of current GPUs is almost always disjoint from its host system.*

- **GPGPU:** A General Purpose GPU, which is a GPU capable of running non-graphics computations.
- **GPU:** A specialized (co)processor that offloads graphics computation and rendering from the host. As GPUs have evolved, they have become increasingly able to offload non-graphics computations as well (see GPGPU).
- **Heterogenous programming**  
A workload that combines kernels executing on data-parallel compute nodes with algorithms running on CPUs.
- **Host**  
The operating system process and the CPU(s) that it is running on.
- **Host thread**  
The operating system thread and the CPU(s) that it is running on. A host thread may initiate a copy operation or parallel loop operation that may run on an accelerator.
- **Index**  
A vector of integers that describes an N-dimensional point in iteration space or index space.
- **Kernel; Kernel function**  
A program designed to be executed at a C++ AMP call-site. More generally, a kernel is a unit of computation that executes on an accelerator. A kernel function is a special case; it is the root of a logical call graph of functions that execute on an accelerator. A C++ analogy is that it is the “`main()`” function for an accelerator program
- **Perfect loop nest**  
A loop nest in which the body of each outer loop consists of a single statement that is a loop.
- **Pixel**  
A pixel, or *picture element*, represents a single element in a digital image. Typically pixels are composed of multiple color components such as a red, green and blue values. Other color representation exist, including single channel images that just represent intensity or black and white values.
- **SIMD unit**  
Single Instruction Multiple Data. A machine programming model where a single instruction operates over multiple pieces of data. Translating a program to use SIMD is known as vectorization. GPUs have multiple SIMD units, which are the streaming multiprocessors.  
*Informative: An SSE (Nehalem, Phenom) or AVX (Sandy Bridge) or LRBni (Larrabee) vector unit is a SIMD unit or vector processor.*
- **SMP**  
Symmetric Multi-Processor – standard PC multiprocessor architecture.
- **Texel**  
A texel or *texture element* represents a single element of a texture space. Texel elements are mapped to 1D, 2D or 3D surfaces during sampling, rendering and/or rasterization and end up as pixel elements on a display.
- **Texture**  
A texture is a 1, 2 or 3 dimensional logical array of texels which is optimized in hardware for spacial access using texture caches. Textures typically are used to represent image, volumetric or other visual information, although they are efficient for many data arrays which need to be optimized for spacial access or need to interpolate

between adjacent elements. Textures provide virtualization of storage, whereby shader code can sample a texture object as if it contained logical elements of one type (e.g., float4) whereas the concrete physical storage of the texture is represented in terms of a second type (e.g., four 8-bit channels). This allows the application of the same shader algorithms on different types of concrete data.

- **Texture Format**

Texture formats define the type and arrangement of the underlying bytes representing a texel value.

*Informative: Direct3D supports many types of formats, which are described under the DXGI\_FORMAT enumeration.*

- **Texture memory**

Texture memory space resides in GPU memory and is cached in texture cache. A texture fetch costs one memory read from GPU memory only on a cache miss, otherwise it just costs one read from texture cache. The texture cache is optimized for 2D spatial locality, so threads of the same scheduling unit that read texture addresses that are close together in 2D will achieve best performance. Also, it is designed for streaming fetches with a constant latency; a cache hit reduces global memory bandwidth demand but not fetch latency.

- **Thread group; Thread tile**

A set of threads that are scheduled together, can share tile-local memory, and can participate in barrier synchronization.

- **Tile-local memory**

User-defined cache on streaming multiprocessors on GPUs. Shared memory is local to a multiprocessor and shared across threads executing on the same multiprocessor. Shared memory allocations per thread group will affect the total number of thread groups that are in-flight per multiprocessor. For example, if each thread uses 4KB of thread memory and there is a limit of 16KB per multiprocessor, then the number of thread groups in-flight is limited to 4 and perhaps less depending upon register allocation patterns.

- **Tiling**

Tiling is the partitioning of an N-dimensional array into same sized 'tiles' which are N-dimensional rectangles with sides parallel to the coordinate axes. Tiling is essentially the process of recognizing the current thread group as being a cooperative gang of threads, with the decomposition of a global index into a local index plus a tile offset. In C++ AMP it is viewing a global index as a local index and a tile ID described by the canonical correspondence:

*compute grid ~ dispatch grid x thread group*

In particular, tiling provides the local geometry with which to take advantage of shared memory and barriers whose usage patterns enable coalescing of global memory access.

- **Restricted function**

A function that is declared to obey the restrictions of a particular C++ AMP subset. A function can be CPU-restricted, in which case it can run on a host CPU. A function can be amp-restricted, in which case it can run on an amp-capable accelerator, such as a GPU or VectorCPU. A function can carry more than one restriction.

## 1.3 Error Model

Host-side runtime library code for C++ AMP has a different error model than device-side code. For more details, examples and exception categorization see Error Handling.

**Host-Side Error Model:** On a host, C++ exceptions and assertions will be used to present semantic errors and hence will be categorized and listed as error states in API descriptions.

**Device-Side Error Model:** On a device, error state is conveyed through the `assert` intrinsic. The `debug_printf` intrinsic is additionally supported for logging messages from within the accelerator code.

**Compile-time asserts:** The C++ intrinsic `static_assert` is often used to handle error states that are detectable at compile time. In this way `static_assert` is a technique for conveying static semantic errors and as such they will be categorized similar to exception types.

## 1.4 Programming Model

Here are the types and patterns that comprise C++ AMP.

- **Indexing level**
  - `index<N>`
  - `extent<N>`
  - `tilde_extent<D0,D1,D2>`
  - `tilde_index<D0,D1,D2>`
- **Data level**
  - `array<T,N>`
  - `array_view<T,N>`, `array_view<const T,N>`
- **Runtime level**
  - `accelerator`
  - `accelerator_view`
- **Call-site level**
  - `parallel_for_each`
  - `copy` – various commands to move data between compute nodes
- **Kernel level**
  - `tile_barrier`
  - `restrict()` clause
  - `tile_static`
  - Atomic functions
  - Math functions (precise and fast)
- **Graphics (optional)**
  - `texture<T,N>`
  - `writeonly_texture_view<T,N>`
  - Short vector types
- **direct3d interop (optional and Microsoft-specific)**
  - Data interoperation on arrays and textures
  - Scheduling interoperation accelerators and accelerator views
  - **Direct3d intrinsic functions for clamping, bit counting, and other special arithmetic operations.**

## 2 C++ Language Extensions for Accelerated Computing

C++ AMP adds a closed set<sup>1</sup> of restriction specifiers to the C++ type system, with new syntax, as well as rules for how they behave with respect to conversion rules and overloading.

Restriction specifiers apply to function declarators only. The restriction specifiers perform the following functions:

1. They become part of the signature of the function.
2. They enforce restrictions on the content and/or behaviour of that function.
3. They may designate a particular subset of the C++ language

For example, an “amp” restriction would imply that a function must conform to the defined subset of C++ such that it is amenable for use on a typical GPU device.

---

<sup>1</sup> There is no mechanism proposed here to allow developers to extend the set of restrictions.

## 2.1 Syntax

A new grammar production is added to represent a sequence of such restriction specifiers.

```
restriction-specifier-seq:
    restriction-specifier
    restriction-specifier-seq restriction-specifier
```

```
restriction-specifier:
    restrict ( restriction-seq )
```

```
restriction-seq:
    restriction
    restriction-seq , restriction
```

```
restriction:
    amp-restriction
    cpu
```

```
amp-restriction:
    amp
```

The **restrict** keyword is a contextual keyword. The restriction specifiers contained within a **restrict** clause are not reserved words.

Multiple restrict clauses, such as **restrict(A) restrict(B)**, behave exactly the same as **restrict(A,B)**. Duplicate restrictions are allowed and behave as if the duplicates are discarded.

The **cpu** restriction specifies that this function will be able to run on the host CPU.

If a declarator elides the restriction specifier, it behaves as if it were specified with **restrict(cpu)**, except when a restriction specifier is determined by the surrounding context as specified in section 2.2.1. If a declarator contains a restriction specifier, then it specifies the entire set of restrictions (in other words: **restrict(amp)** means will be able to run on the amp target, need not be able to run the CPU).

### 2.1.1 Function Declarator Syntax

The function declarator grammar (classic & trailing return type variation) are adjusted as follows:

```
D1 ( parameter-declaration-clause ) cv-qualifier-seqopt ref-qualifieropt restriction-specifier-seqopt
    exception-specificationopt attribute-specifieropt
```

```
D1 ( parameter-declaration-clause ) cv-qualifier-seqopt ref-qualifieropt restriction-specifier-seqopt
    exception-specificationopt attribute-specifieropt trailing-return-type
```

Restriction specifiers shall not be applied to other declarators (e.g.: arrays, pointers, references). They can be applied to all kinds of functions including free functions, static and non-static member functions, special member functions, and overloaded operators.

Examples:

```
auto grod() restrict(amp);
auto freedle() restrict(amp)-> double;
class Fred {
```

```

298 public:
299     Fred() restrict(amp)
300         : member-initializer
301     { }
302
303     Fred& operator=(const Fred&) restrict(amp);
304
305     int kreeble(int x, int y) const restrict(amp);
306     static void zot() restrict(amp);
307 };

```

*restriction-specifier-seq<sub>opt</sub>* applies to to all expressions between the *restriction-specifier-seq* and the end of the function-definition, lambda-expression, member-declarator, lambda-declarator or declarator.

### 2.1.2 Lambda Expression Syntax

The lambda expression syntax is adjusted as follows:

*lambda-declarator:*

( *parameter-declaration-clause* ) *attribute-specifier<sub>opt</sub>* *mutable<sub>opt</sub>* *restriction-specifier-seq<sub>opt</sub>*  
*exception-specification<sub>opt</sub>* *trailing-return-type<sub>opt</sub>*

When a restriction modifier is applied to a lambda expression, the behavior is as if all member functions of the generated functor are restriction-modified.

### 2.1.3 Type Specifiers

Restriction specifiers are not allowed anywhere in the type specifier grammar, even if it specifies a function type. For example, the following is not well-formed and will produce a syntax error:

```

323 typedef float FuncType(int);
324
325 restrict(cpu) FuncType* pf; // Illegal; restriction specifiers not allowed in type specifiers
326
327

```

The correct way to specify the previous example is:

```

329 typedef float FuncType(int) restrict(cpu);
330
331 FuncType* pf;
332
333

```

or simply

```

335 float (*pf)(int) restrict(cpu);
336
337

```

## 2.2 Meaning of Restriction Specifiers

The restriction specifiers on the declaration of a given function *F* must agree with those specified on the definition of function *F*.

Multiple restriction specifiers may be specified for a given function: the effect is that the function enforces the union of the restrictions defined by each restriction modifier.

*Informative: not for this release: It is possible to imagine two restriction specifiers that are intrinsically incompatible with each other (for example, **pure** and **elemental**). When this occurs, the compiler will produce an error.*

*Refer to section 13 for treatment of versioning of restrictions*

The restriction specifiers on a function become part of its signature, and thus can be used to overload.

Every expression (or sub-expression) that is evaluated in code that has multiple restriction specifiers must have the same type in the context of each restriction. It is a compile-time error if an expression can evaluate to different types under the different restriction specifiers. Function overloads should be defined with care to avoid a situation where an expression can evaluate to different types with different restrictions.

### 2.2.1 Function Definitions

The restriction specifiers applied to a function definition are recursively applied to all function declarators and type names defined within its body that do not have explicit restriction specifiers (i.e.: through nested classes that have member functions, and through lambdas.) For example:

```
void glorp() restrict(amp) {
    class Foo {
        void zot() {...} // "zot" is amp-restricted
    };

    auto f1 = [] (int y) { ... }; // Lambda is amp-restricted
    auto f2 = [] (int y) restrict(cpu) { ... }; // Lambda is cpu-restricted
    typedef int int_void_amp(); // int_void_amp is amp-restricted
    ...
}
```

This also applies to the function scope of a lambda body.

### 2.2.2 Constructors and Destructors

Constructors can have overloads that are differentiated by restriction specifiers.

Since destructors cannot be overloaded, the destructor must contain a restriction specifier that covers the union of restrictions on all the constructors. (A destructor can achieve the same effect of overloading by calling auxiliary cleanup functions that have different restriction specifiers.)

For example:

```
class Foo {
public:
    Foo() { ... }
    Foo() restrict(amp) { ... }

    ~Foo() restrict(cpu,amp);
};

void UnrestrictedFunction() {
    Foo a; // calls "Foo::Foo()"
    ...
    // a is destructed with "Foo::~~Foo()"
}

void RestrictedFunction() restrict(amp) {
    Foo b; // calls "Foo::Foo() restrict(amp)"
    ...
    // b is destructed with "Foo::~~Foo()"
}

class Bar {
public:
    Bar() { ... }
    Bar() restrict(amp) { ... }

    ~Bar(); // error: restrict(cpu,amp) required
};
```

A virtual function declaration in a derived class will override a virtual function declaration in a base class only if the derived class function has the same restriction specifiers as the base. E.g.:

```
class Base {
public:
    virtual void foo() restrict(R1);
};

class Derived : public Base {
public:
    virtual void foo() restrict(R2); // Does not override Base::foo
};
```

### 2.2.3 Lambda Expressions

When restriction specifiers are applied to a lambda declarator, the behavior is as if the restriction specifiers are applied to all member functions of the compiler-generated function object. For example:

```
Foo ambientVar;

auto functor = [ambientVar] (int y) restrict(amp) -> int { return y + ambientVar.z; };
```

is equivalent to:

```
Foo ambientVar;

class <lambdaName> {
public:
    <lambdaName>(const Foo& foo)
        : capturedFoo(foo)
    { }

    ~<lambdaName>() { }

    int operator()(int y) restrict(amp) { return y + ambientVar.z; }
};

<lambdaName> functor;
```

## 2.3 Expressions Involving Restricted Functions

### 2.3.1 Function pointer conversions

New implicit conversion rules must be added to account for restricted function pointers (and references). Given an expression of type “pointer to  $R_1$ -function”, this type can be implicitly converted to type “pointer to  $R_2$ -function” if and only if  $R_1$  has all the restriction specifiers of  $R_2$ . Stated more intuitively, it is okay for the target function to be more restricted than the function pointer that invokes it; it’s not okay for it to be less restricted. E.g.:

```
int func(int) restrict(R1,R2);
int (*pfn)(int) restrict(R1) = func; // ok, since func(int) restrict(R1,R2) is at least R1
```

(Note that C++ AMP does not support function pointers in the current *restrict(amp)* subset)

### 2.3.2 Function Overloading

Restriction specifiers become part of the function type to which they are attached. I.e.: they become part of the signature of the function. Functions can thus be overloaded by differing modifiers, and each unique set of modifiers forms a unique overload.

The restriction specifiers of a function shall not overlap with any restriction specifiers in another function within the same overload set.



```

467
468 int func(int x) restrict(cpu,amp);
469 int func(int x) restrict(cpu); // error, overlaps with previous declaration
470

```

The target of the function call operator must resolve to an overloaded set of functions that is *at least* as restricted as the body of the calling function (see Overload Resolution). E.g.:

```

473
474 void grod();
475 void glorp() restrict(amp);
476
477 void foo() restrict(amp) {
478     glorp(); // okay: glorp has amp restriction
479     grod();  // error: grod lacks amp restriction
480 }
481

```

It is permissible for a less-restrictive call-site to call a more-restrictive function.

Compiler-generated constructors and destructors (and other special member functions) behave as if they were declared with as many restrictions as possible while avoiding ambiguities and errors. For example:

```

486
487 struct Grod {
488     int a;
489     int b;
490
491     // compiler-generated default constructor: Grod() restrict(cpu,amp);
492
493     int frool() restrict(amp) {
494         return a+b;
495     }
496
497     int blarg() restrict(cpu) {
498         return a*b;
499     }
500
501     // compiler-generated destructor: ~Grod() restrict(cpu,amp);
502 };
503
504 void d3dCaller() restrict(amp) {
505     Grod g; // okay because compiler-generated default constructor is restrict(amp)
506
507     int x = g.frool();
508
509     // g.~Grod() called here; also okay
510 }
511
512 void d3dCaller() restrict(cpu) {
513     Grod g; // okay because compiler-generated default constructor is restrict(cpu)
514
515     int x = g.blarg();
516
517     // g.~Grod() called here; also okay
518 }
519

```

The compiler must behave this way since the local usage of “Grod” in this case should not affect other potential uses of it in other restricted or unrestricted scopes.

More specifically, the compiler follows the standard C++ rules, ignoring restrictions, to determine which special member functions to generate and how to generate them. Then the restrictions are set according to the following steps:

The compiler sets the restrictions of compiler-generated destructors to the intersection of the restrictions on all of the destructors of the data members [able to destroy all data members] and all of the base classes’ destructors [able to call all

base classes' destructors]. If there are no such destructors, then all possible restrictions are used [able to destroy in any context]. However, any restriction that would result in an error is not used.

More specifically, the compiler sets the restrictions of compiler-generated destructors to the intersection of the restrictions on all of the destructors of the member fields *[able to destroy all member fields]* and all of the base classes' destructors [able to call all base classes' destructors]. If there are no such destructors, then all possible restrictions are used [able to destroy in any context]. However, any restriction that would result in an error is not set.

The compiler sets the restrictions of compiler-generated default constructors to the intersection of the restrictions on all of the default constructors of the member fields [able to construct all member fields], all of the base classes' default constructors [able to call all base classes' default constructors], and the destructor of the class [able to destroy in any context constructed]. However, any restriction that would result in an error is not set.

The compiler sets the restrictions of compiler-generated copy constructors to the intersection of the restrictions on all of the copy constructors of the member fields [able to construct all member fields], all of the base classes' copy constructors [able to call all base classes' copy constructors], and the destructor of the class [able to destroy in any context constructed]. However, any restriction that would result in an error is not set.

The compiler sets the restrictions of compiler-generated assignment operators to the intersection of the restrictions on all of the assignment operators of the member fields [able to assign all member fields] and all of the base classes' assignment operators [able to call all base classes' assignment operators]. However, any restriction that would result in an error is not set.

### 2.3.2.1 Overload Resolution

Overload resolution depends on the set of restrictions (function modifiers) in force at the call site.

```
int func(int x) restrict(A);
int func(int x) restrict(B,C);
int func(int x) restrict(D);

void foo() restrict(B) {
    int x = func(5); // calls func(int x) restrict(B,C)
    ...
}
```

A call to function *F* is valid if and only if the overload set of *F* covers all the restrictions in force in the calling function. This rule can be satisfied by a single function *F* that contains all the require restrictions, or by a set of overloaded functions *F* that each specify a subset of the restrictions in force at the call site. For example:

```
void Z() restrict(amp,sse2,cpu) { }

void Z_caller() restrict(amp,sse,cpu) {
    Z(); // okay; all restrictions available in a single function
}

void X() restrict(amp) { }
void X() restrict(sse) { }
void X() restrict(cpu) { }

void X_caller() restrict(amp,sse,cpu) {
    X(); // okay; all restrictions available in separate functions
}

void Y() restrict(amp) { }
```

<sup>2</sup> Note that “sse” is used here for illustration only, and does not imply further meaning to it in this specification.

```

582
583 void Y_caller() restrict(cpu,amp) {
584     Y(); // error; no available Y() that satisfies CPU restriction
585 }

```

When a call to a restricted function is satisfied by more than one function, then the compiler must generate an as-if-runtime<sup>3</sup>-dispatch to the correctly restricted version.

### 2.3.2.2 Name Hiding

Overloading via restriction specifiers does not affect the name hiding rules. For example:

```

591
592 void foo(int x) restrict(amp) { ... }
593
594 namespace N1 {
595     void foo(double d) restrict(cpu) { .... }
596
597     void foo_caller() restrict(amp) {
598         foo(10); // error; global foo() is hidden by N1::foo
599     }
600 }

```

The name hiding rules in C++11 Section 3.3.10 state that within namespace N1, the global name “Foo” is hidden by the local name “Foo”, and is *not overloaded* by it.

### 2.3.3 Casting

A restricted function type can be cast to a more restricted function type using a normal C-style cast or *reinterpret\_cast*. (A cast is not needed when losing restrictions, only when gaining.) For example:

```

607
608 void unrestricted_func(int,int);
609
610 void restricted_caller() restrict(amp) {
611     ((void (*)(int,int) restrict(amp))unrestricted_func)(6, 7);
612     reinterpret_cast<(void (*)(int,int) restrict(amp))>(unrestricted_func)(6, 7);
613 }

```

A program which attempts to invoke a function expression after such unsafe casting can exhibit undefined behavior.

## 2.4 amp Restriction Modifier

The *amp* restriction modifier applies a relatively small set of restrictions that reflect the current limitations of GPU hardware and the underlying programming model.

### 2.4.1 Restrictions on Types

Not all types can be supported on current GPU hardware. The *amp* restriction modifier restricts functions from using unsupported types, in their function signature or in their function bodies.

We refer to the set of supported types as being *amp-compatible*. Any type referenced within an *amp* restriction function shall be *amp-compatible*. Some uses require further restrictions.

#### 2.4.1.1 Type Qualifiers

The *volatile* type qualifier is not supported within an *amp*-restricted function. A variable or member qualified with *volatile* may not be declared or accessed in *amp* restricted code.

<sup>3</sup> Compilers are always free to optimize this if they can determine the target statically.

### 2.4.1.2 Fundamental Types

Of the set of C++ fundamental types only the following are supported within an amp-restricted function as *amp-compatible* types.

- *bool*
- *int, unsigned int*
- *long, unsigned long*
- *float, double*
- *void*

The representation of these types on a device running an *amp* function is identical to that of its host.

### 2.4.1.3 Compound Types

Pointers shall only point to *amp-compatible* types or *concurrency::array* or *concurrency::graphics::texture*. Pointers to pointers are not supported. *std::nullptr\_t* type is supported and treated as a pointer type. No pointer type is considered *amp-compatible*. Pointers are only supported as local variables and/or function parameters and/or function return types.

References (lvalue and rvalue) shall refer only to *amp-compatible* types and/or *concurrency::array* and/or *concurrency::graphics::texture*. Additionally, references to pointers are supported as long as the pointer type is itself supported. Reference to *std::nullptr\_t* is not allowed. No reference type is considered *amp-compatible*. References are only supported as local variables and/or function parameters and/or function return types.

*array\_view* and *writeonly\_texture\_view* are *amp-compatible* types.

A user defined type (class, struct, union) is *amp-compatible* if

- it contains only data members whose types are *amp-compatible*, except for references to instances of classes *array* and *texture*, and
- the offset of its data members and base classes are at least four bytes aligned, and
- its data members shall not be bitfields, and
- it shall not have *virtual* base classes, and *virtual* member function, and
- all of its base classes are *amp-compatible*.

The element type of an array shall be *amp-compatible* and four byte aligned.

Pointers to members (C++11 8.3.3) shall only refer to non-static data members.

Enumeration types shall have underlying types consisting of *int, unsigned int, long, or unsigned long*.

The representation of an *amp-compatible* compound type (with the exception of pointer & reference) on a device is identical to that of its host.

### 2.4.2 Restrictions on Function Declarators

The function declarator (C++11 8.3.5) of an amp-restricted function:

- shall not have a trailing ellipsis (...) in its parameter list
- shall have no parameters, or shall have parameters whose types are *amp-compatible*
- shall have a return type that is *void* or is *amp-compatible*
- shall not be *virtual*
- shall not have a throw specification
- shall not have *extern "C"* linkage when multiple restriction specifiers are present

### 2.4.3 Restrictions on Function Scopes

The function scope of an amp-restricted function may contain any valid C++ declaration, statement, or expression except for those which are specified here.

#### 2.4.3.1 Literals

A C++ AMP program is ill-formed if the value of an integer constant or floating point constant exceeds the allowable range of any of the above types.

#### 2.4.3.2 Primary Expressions (C++11 5.1)

An identifier or qualified identifier that refers to an object shall refer only to:

- a parameter to the function, or
- a local variable declared at a block scope within the function, or
- a non-static member of the class of which this function is a member, or
- a *static const* type that can be reduced to a integer literal and is only used as an rvalue, or
- a global *const* type that can be reduced to a integer literal and is only used as an rvalue, or
- a captured variable in a lambda expression.

#### 2.4.3.3 Lambda Expressions

If a lambda expression appears within the body of an amp-restricted function, the *amp* modifier may be elided and the lambda is still considered an amp lambda.

A lambda expression shall not capture any context variable by reference, except for context variables of type *concurrency::array* and *concurrency::graphics::texture*.

#### 2.4.3.4 Function Calls (C++11 5.2.2)

The target of a function call operator:

- shall not be a virtual function
- shall not be a pointer to a function
- shall not recursively invoke itself or any other function that is directly or indirectly recursive.

These restrictions apply to all function-like invocations including:

- object constructors & destructors
- overloaded operators, including *new* and *delete*.

#### 2.4.3.5 Local Declarations

Local declarations shall not specify any storage class other than *register*, or *tile\_static*. Variables shall have types that are *amp-compatible*, pointers to *amp-compatible* types, or references to *amp-compatible* types.

##### 2.4.3.5.1 *tile\_static* Variables

A variable declared with the *tile\_static* storage class can be accessed by all threads within a tile (group of threads). (The *tile\_static* storage class is valid only within a *restrict(amp)* context.) The storage lifetime of a *tile\_static* variable begins when the execution of a thread in a tile reaches the point of declaration, and ends when the kernel function is exited by the last thread in the tile. Each thread tile accessing the variable shall perceive to access a separate, per-tile, instance of the variable.

A *tile\_static* variable declaration does not constitute a barrier (see 8.1.1). *tile\_static* variables are not initialized by the compiler and assume no default initial values.

The *tile\_static* storage class shall only be used to declare local (function or block scope) variables. The type of a *tile\_static* variable shall not be a pointer or reference type.

The type of a *tile\_static* variable or array must be *amp-compatible* and shall not directly or recursively contain any concurrency containers (e.g. *concurrency::array\_view*) or reference to concurrency containers.

A *tile\_static* variable shall not have an initializer and no constructors or destructors will be called for it; its initial contents are undefined.

#### 2.4.3.6 Type-Casting Restrictions

A type-cast shall not be used to convert a pointer to an integral type, nor an integral type to a pointer. This restriction applies to *reinterpret\_cast* (C++11 5.2.10) as well as to C-style casts (C++11 5.4).

Casting away *const*-ness may result in a compiler warning and/or undefined behavior.

#### 2.4.3.7 Miscellaneous Restrictions

The pointer-to-member operators *.\** and *->\** shall only be used to access pointer-to-data member objects.

Pointer arithmetic shall not be performed on pointers to *bool* values.

Furthermore, an amp-restricted function shall not contain any of the following:

- *dynamic\_cast* or *typeid* operators
- *goto* statements or labeled statements
- *asm* declarations
- Function *try* block, *try* blocks, *catch* blocks, or *throw*.

## 3 Device Modeling

### 3.1 The concept of a compute accelerator

A compute accelerator is a hardware capability that is optimized for data-parallel computing. An accelerator may be a device attached to a PCIe bus (such as a GPU), or it might be an extended instruction set on the main CPU (such as SSE or AVX).

*Informative: Future architectures might bridge these two extremes, such as AMD's Fusion or Intel's Knight's Ferry.*

In the C++ AMP model, an accelerator may have private memory which is not generally accessible by the host. C++ AMP allows data to be allocated in the accelerator memory and references to this data may be manipulated on the host. It is assumed that all data accessed within a kernel must be stored in accelerator memory although some C++ AMP scenarios will implicitly make copies of data logically stored on the host.

C++ AMP has functionality for copying data between host and accelerator memories. A copy from accelerator-to-host is always a synchronization point, unless an explicit asynchronous copy is specified. In general, for optimal performance, memory content should stay on an accelerator as long as possible.

In some cases, accelerator memory and CPU memory are one and the same. And depending upon the architecture, there may never be any need to copy between the two physical locations of memory. C++ AMP provides for coding patterns that allow the C++ AMP runtime to avoid or perform copies as required.

### 3.2 accelerator

An *accelerator* is an abstraction of a physical data-parallel-optimized compute node. An accelerator is often a discrete GPU, but can also be a virtual host-side entity such as the Microsoft DirectX *REF* device, or *WARP* (a CPU-side device accelerated using SSE instructions), or can refer to the CPU itself.

### 3.2.1 Default Accelerator

C++ AMP supports the notion of a default accelerator, an accelerator which is chosen automatically when the program does not explicitly do so.

A user may explicitly create a default accelerator object in one of two ways:

1. Invoke the default constructor:

```
accelerator def;
```

2. Use the `default_accelerator` device path:

```
accelerator def(accelerator::default_accelerator);
```

The user may also influence which accelerator is chosen as the default by calling `accelerator::set_default` prior to invoking any operation which would otherwise choose the default. Such operations include the above two calls, as well as invoking `parallel_for_each` without an explicit `accelerator_view` argument, creating an `array` not bound to an explicit `accelerator_view`, etc.

If the user does not call `accelerator::set_default`, the default is chosen in an implementation specific manner.

#### Microsoft-specific:

*The Microsoft implementation of C++ AMP uses the the following heuristic to select a default accelerator when one is not specified by a call to `accelerator::set_default`:*

1. *If using the debug runtime, prefer an accelerator that supports debugging.*
2. *If the process environment variable `CPPAMP_DEFAULT_ACCELERATOR` is set, interpret its value as a device path and prefer the device that corresponds to it.*
3. *Otherwise, the following criteria are used to determine the 'best' accelerator:*
  - a. *Prefer non-emulated devices*
  - b. *Prefer the device with the most available memory.*
  - c. *Prefer the device which is not attached to the display.*

*Note that the `cpu_accelerator` is never considered among the candidates in the above heuristic.*

### 3.2.2 Synopsis

```
class accelerator
{
public:
    static const wchar_t default_accelerator[]; // = L"default"
```

```
// Microsoft-specific:
```

```
static const wchar_t direct3d_warp[]; // = L"direct3d\\warp"
```

```
static const wchar_t direct3d_ref[]; // = L"direct3d\\ref"
```

```
static const wchar_t cpu_accelerator[]; // = L"cpu"
```

```
accelerator();
```

```
explicit accelerator(const wstring& path);
```

```
accelerator(const accelerator& other);
```

```
static vector<accelerator> get_all();
```

```
static bool set_default(const wstring& path);
```

```

819
820     accelerator& operator=(const accelerator& other);
821
822     __declspec(property(get)) wstring device_path;
823     __declspec(property(get)) unsigned int version; // hiword=major, loword=minor
824     __declspec(property(get)) wstring description;
825     __declspec(property(get)) bool is_debug;
826     __declspec(property(get)) bool is_emulated;
827     __declspec(property(get)) bool has_display;
828     __declspec(property(get)) bool supports_double_precision;
829     __declspec(property(get)) bool supports_limited_double_precision;
830     __declspec(property(get)) size_t dedicated_memory;
831     __declspec(property(get)) accelerator_view default_view;
832
833     accelerator_view create_view();
834     accelerator_view create_view(queuing_mode qmode);
835
836     bool operator==(const accelerator& other) const;
837     bool operator!=(const accelerator& other) const;
838 };
839
840

```

#### class accelerator

Represents a physical accelerated computing device. An object of this type can be created by enumerating the available devices, or getting the default device, the reference device, or the WARP device.

*Microsoft-specific:*

*The WARP device may not be available on all platforms, not even all Microsoft platforms.*

### 3.2.3 Static Members

#### static vector<accelerator> accelerator::get\_all()

Returns a std::vector of accelerator objects (in no specific order) representing all accelerators that are available, including reference accelerators and WARP accelerators if available.

##### Return Value:

A vector of accelerators.

#### static bool set\_default(const wstring& path);

Sets the default accelerator to the device path named by the "path" argument. See the constructor "accelerator(const wstring& path)" for a description of the allowable path strings.

This establishes a process-wide default accelerator and influences all subsequent operations that might create a default accelerator.

##### Parameters

<i>path</i>	The device path of the default accelerator.
-------------	---

##### Return Value:

A Boolean flag indicating whether the default was set. If the default has already been set for this process, this value will be *false*, and the function will have no effect.

845

### 3.2.4 Constructors

847

#### accelerator()



Constructs a new accelerator object that represents the default accelerator. This is equivalent to calling the constructor "accelerator(accelerator::default\_accelerator)".

The actual accelerator chosen as the default can be affected by calling "accelerator::set\_default" prior to calling this constructor.

#### Parameters:

None.

848

### accelerator(const wstring& path)

Constructs a new accelerator object that represents the physical device named by the "path" argument. If the path represents an unknown or unsupported device, an exception will be thrown.

The path can be one of the following:

1. accelerator::default\_accelerator (or L"default"), which represents the path of the fastest accelerator available, as chosen by the runtime.
2. accelerator::cpu\_accelerator (or L"cpu"), which represents the CPU. Note that parallel\_for\_each shall not be invoked over this accelerator.
3. A valid device path that uniquely identifies a hardware accelerator available on the host system.

*Microsoft-specific:*

4. accelerator::direct3d\_warp (or L"direct3d\\warp"), which represents the WARP accelerator
5. accelerator::direct3d\_ref (or L"direct3d\\ref"), which represents the REF accelerator.

#### Parameters:

path

The device path of this accelerator.

849

### accelerator(const accelerator& other);

Copy constructs an accelerator object. This function does a shallow copy with the newly created accelerator object pointing to the same underlying device as the passed accelerator parameter.

#### Parameters:

other

The accelerator object to be copied.

850

## 3.2.5 Members

851

852

```
static const wchar_t default_accelerator[]
static const wchar_t direct3d_warp[]
static const wchar_t direct3d_ref[]
static const wchar_t cpu_accelerator[]
```

These are static constant string literals that represent device paths for known accelerators, or in the case of "default\_accelerator", direct the runtime to choose an accelerator automatically.

**default\_accelerator:** The string L"default" represents the default accelerator, which directs the runtime to choose the fastest accelerator available. The selection criteria are discussed in section 3.2.1 Default Accelerator.

**cpu\_accelerator:** The string L"cpu" represents the host system. This accelerator is used to provide a location for system-allocated memory such as arrays and staging arrays. It is not a valid target for accelerated computations.

*Microsoft-specific:*

**direct3d\_warp:** The string L"direct3d\\warp" represents the device path of the CPU-accelerated Warp device. On other non-direct3d platforms, this member may not exist.

**direct3d\_ref:** The string L"direct3d\\ref" represents the software rasterizer, or Reference, device. This particular device is useful for debugging. On other non-direct3d platforms, this member may not exist.

853

**accelerator& operator=(const accelerator& other)**

Assigns an accelerator object to "this" accelerator object and returns a reference to "this" object. This function does a shallow assignment with the newly created accelerator object pointing to the same underlying device as the passed accelerator parameter.

**Parameters:**

<i>other</i>	The accelerator object to be assigned from.
--------------	---

**Return Value:**

A reference to "this" accelerator object.

854

**\_\_declspec(property(get)) accelerator\_view default\_view**

Returns the default accelerator view associated with the accelerator. The queuing\_mode of the default accelerator\_view is queuing\_mode\_automatic.

**Return Value:**

The default **accelerator\_view** object associated with the accelerator.

855

**accelerator\_view create\_view(queuing\_mode qmode)**

Creates and returns a new accelerator view on the accelerator with the supplied queuing mode.

**Return Value:**

The new **accelerator\_view** object created on the compute device.

**Parameters:**

<i>qmode</i>	The queuing mode of the accelerator_view to be created. See "- Queuing Mode".
--------------	---

856

**accelerator\_view create\_view()**

Creates and returns a new resource view on the accelerator. Equivalent to "create\_view(queuing\_mode\_automatic)".

**Return Value:**

The new **accelerator\_view** object created on the compute device.

857

858

**bool operator==(const accelerator& other) const**

Compares "this" accelerator with the passed accelerator object to determine if they represent the same underlying device.

**Parameters:**

<i>other</i>	The accelerator object to be compared against.
--------------	--

**Return Value:**

A boolean value indicating whether the passed accelerator object is same as "this" accelerator.

859

860

**bool operator!=(const accelerator& other) const**

Compares "this" accelerator with the passed accelerator object to determine if they represent different devices.

**Parameters:**

<i>other</i>	The accelerator object to be compared against.
--------------	--

**Return Value:**

A boolean value indicating whether the passed accelerator object is different from "this" accelerator.

### 3.2.6 Properties

The following read-only properties are part of the public interface of the class ***accelerator***, to enable querying the accelerator characteristics:

**`__declspec(property(get)) wstring device_path`**

Returns a system-wide unique device instance path that matches the “Device Instance Path” property for the device in Device Manager, or one of the predefined path constants **`direct3d_warp`** or **`direct3d_ref`**.

**`__declspec(property(get)) wstring description`**

Returns a short textual description of the accelerator device.

**`__declspec(property(get)) unsigned int version`**

Returns a 32-bit unsigned integer representing the version number of this accelerator. The format of the integer is major.minor, where the major version number is in the high-order 16 bits, and the minor version number is in the low-order bits.

**`__declspec(property(get)) bool has_display`**

Returns a boolean value indicating whether the accelerator is attached to a display.

**`__declspec(property(get)) bool dedicated_memory`**

Returns the amount of dedicated memory (in KB) on an accelerator device. There is no guarantee that this amount of memory is actually available to use.

**`__declspec(property(get)) bool supports_double_precision`**

Returns a Boolean value indicating whether this accelerator supports double-precision (**`double`**) computations.

**`__declspec(property(get)) bool supports_limited_double_precision`**

Returns a boolean value indicating whether the accelerator has limited double precision support (excludes double division, precise\_math functions, int to double, double to int conversions) for a parallel\_for\_each kernel.

**`__declspec(property(get)) bool is_debug`**

Returns a boolean value indicating whether the accelerator supports debugging.

**`__declspec(property(get)) bool is_emulated`**

Returns a boolean value indicating whether the accelerator is emulated. This is true, for example, with the reference accelerator.

## 3.3 accelerator\_view

An ***accelerator\_view*** represents a logical view of an accelerator. A single physical compute device may have many logical (isolated) accelerator views. Each accelerator has a default accelerator view and additional accelerator views may be optionally created by the user. Physical devices must potentially be shared amongst many client threads. Client threads may choose to use the same ***accelerator\_view*** of an accelerator or each client may communicate with a compute device via an independent ***accelerator\_view*** object for isolation from other client threads. Work submitted to an ***accelerator\_view*** is guaranteed to be executed in the order that it was submitted; there are no such ordering guarantees for work submitted on different ***accelerator\_views***.

An ***accelerator\_view*** can be created with a queuing mode of “immediate” or “automatic”. (See “Queuing Mode”).

### 3.3.1 Synopsis

**`class accelerator_view`**

```

890 {
891 public:
892     accelerator_view() = delete;
893     accelerator_view(const accelerator_view& other);
894
895     accelerator_view& operator=(const accelerator_view& other);
896
897     __declspec(property(get)) Concurrency::accelerator accelerator;
898     __declspec(property(get)) bool is_debug;
899     __declspec(property(get)) unsigned int version;
900     __declspec(property(get)) queuing_mode queuing_mode;
901
902     void flush();
903     void wait();
904     completion_future create_marker();
905
906     bool operator==(const accelerator_view& other) const;
907     bool operator!=(const accelerator_view& other) const;
908 };
909

```

#### class accelerator\_view

Represents a logical (isolated) accelerator view of a compute accelerator. An object of this type can be obtained by calling the [default\\_view](#) property or [create\\_view](#) member functions on an accelerator object.

910

### 911 3.3.2 Queuing Mode

912

913 An [accelerator\\_view](#) can be created with a queuing mode in one of two states:

```

914     enum queuing_mode {
915         queuing_mode_immediate,
916         queuing_mode_automatic
917     };
918

```

919

920 If the queuing mode is [queuing\\_mode\\_immediate](#), then any commands (such as copy or [parallel\\_for\\_each](#)) are sent to the

921 corresponding accelerator before control is returned to the caller.

922

923 If the queuing mode is [queuing\\_mode\\_automatic](#), then such commands are queued up on a command queue

924 corresponding to this [accelerator\\_view](#). Commands are not actually sent to the device until [flush\(\)](#) is called.

925

### 926 3.3.3 Constructors

927

928 An [accelerator\\_view](#) object may only be constructed using a copy or move constructor. There is no default constructor.

929

#### accelerator\_view(const accelerator\_view& other)

Copy-constructs an [accelerator\\_view](#) object. This function does a shallow copy with the newly created [accelerator\\_view](#) object pointing to the same underlying view as the "other" parameter.

##### Parameters:

<i>other</i>	The <a href="#">accelerator_view</a> object to be copied.
--------------	---

930

### 931 3.3.4 Members

932

#### accelerator\_view& operator=(const accelerator\_view& other)

Assigns an `accelerator_view` object to “this” `accelerator_view` object and returns a reference to “this” object. This function does a shallow assignment with the newly created `accelerator_view` object pointing to the same underlying view as the passed `accelerator_view` parameter.

**Parameters:**

<i>other</i>	The <code>accelerator_view</code> object to be assigned from.
--------------	---

**Return Value:**

A reference to “this” `accelerator_view` object.

933

`__declspec(property(get)) queuing_mode queuing_mode`

Returns the queuing mode that this `accelerator_view` was created with. See “Queuing Mode”.

**Return Value:**

The queuing mode.

934

`__declspec(property(get)) unsigned int version`

Returns a 32-bit unsigned integer representing the version number of this `accelerator_view`. The format of the integer is major.minor, where the major version number is in the high-order 16 bits, and the minor version number is in the low-order bits.

The version of the `accelerator_view` is usually the same as that of the parent `accelerator`.

*Microsoft-specific: The version may differ from the `accelerator` only when the `accelerator_view` is created from a `direct3d` device using the interop API.*

935

`__declspec(property(get)) Concurrency::accelerator accelerator`

Returns the `accelerator` that this `accelerator_view` has been created on.

936

`__declspec(property(get)) bool is_debug`

Returns a boolean value indicating whether the `accelerator_view` supports debugging through extensive error reporting.

The `is_debug` property of the `accelerator_view` is usually same as that of the parent `accelerator`. The value may differ from the `accelerator` only when the `accelerator_view` is created from a `direct3d` device using the interop API.

937

`void wait()`

Performs a blocking wait for completion of all commands submitted to the `accelerator_view` prior to calling `wait`.

**Return Value:**

None

938

`void flush()`

Sends the queued up commands in the `accelerator_view` to the device for execution.

An `accelerator_view` internally maintains a buffer of commands such as data transfers between the host memory and device buffers, and kernel invocations (parallel\_for\_each calls). This member function sends the commands to the device for processing. Normally, these commands are sent to the GPU automatically whenever the runtime determines that they need to be, such as when the command buffer is full or when waiting for transfer of data from the device buffers to host memory. The `flush` member function will send the commands manually to the device.

Calling this member function incurs an overhead and must be used with discretion. A typical use of this member function would be when the CPU waits for an arbitrary amount of time and would like to force the execution of queued device commands in the meantime. It can also be used to ensure that resources on the `accelerator` are reclaimed after all references to them have been removed.

Because *flush* operates asynchronously, it can return either before or after the device finishes executing the buffered commands. However, the commands will eventually always complete.

If the *queuing\_mode* is *queuing\_mode\_immediate*, this function does nothing.

**Return Value:**

None

### completion\_future create\_marker()

This command inserts a marker event into the accelerator\_view's command queue. This marker is returned as a completion\_future object. When all commands that were submitted prior to the marker event creation have completed, the future is ready.

**Return Value:**

A future which can be waited on, and will block until the current batch of commands has completed.

### bool operator==(const accelerator\_view& other) const

Compares "this" accelerator\_view with the passed accelerator\_view object to determine if they represent the same underlying object.

**Parameters:**

*other*

The accelerator\_view object to be compared against.

**Return Value:**

A boolean value indicating whether the passed accelerator\_view object is same as "this" accelerator\_view.

### bool operator!=(const accelerator\_view& other) const

Compares "this" accelerator\_view with the passed accelerator\_view object to determine if they represent different underlying objects.

**Parameters:**

*other*

The accelerator\_view object to be compared against.

**Return Value:**

A boolean value indicating whether the passed accelerator\_view object is different from "this" accelerator\_view.

## 3.4 Device enumeration and selection API

The physical compute devices can be enumerated or selected by calling the following static member function of the class accelerator.

### 3.4.1 Synopsis

```
vector<accelerator> accelerator::get_all();
```

As an example, if one wants to find an accelerator that is not emulated and is not attached to a display, one could do the following:

```
vector<accelerator> gpus = accelerator::get_all();
auto headlessIter = std::find_if(gpus.begin(), gpus.end(), [] (accelerator& acc) {
    return !acc.has_display && !acc.is_emulated;
});
```

## 4 Basic Data Elements

C++ AMP enables programmers to express solutions to data-parallel problems in terms of N-dimensional data aggregates and operations over them.

Fundamental to C++ AMP is the concept of an array. An array associates values in an index space with an element type. For example an array could be the set of pixels on a screen where each pixel is represented by four 32-bit values: [Red](#), [Green](#), [Blue](#) and [Alpha](#). The index space would then be the screen resolution, for example all points:

$$\{ \{y, x\} \mid 0 \leq y < 1200, 0 \leq x < 1600, x \text{ and } y \text{ are integers} \}.$$

### 4.1 index<N>

Defines an N-dimensional index point; which may also be viewed as a vector based at the origin in N-space.

The index<N> type represents an N-dimensional vector of [int](#) which specifies a unique position in an N-dimensional space. The dimensions in the coordinate vector are ordered from most-significant to least-significant. Thus, in Cartesian 3-dimensional space, where a common convention exists that the Z dimension (plane) is most significant, the Y dimension (row) is second in significance and the X dimension (column) is the least significant, the index vector (2,0,4) represents the position at (Z=2, Y=0, X=4).

The position is relative to the origin in the N-dimensional space, and can contain negative component values.

*Informative: As a scoping decision, it was decided to limit specializations of index, extent, etc. to 1, 2, and 3 dimensions. This also applies to arrays and array\_views. General N-dimensional support is still provided with slightly reduced convenience.*

#### 4.1.1 Synopsis

```
template <int N>
class index {
public:
    static const int rank = N;
    typedef int value_type;

    index() restrict(amp,cpu);
    index(const index& other) restrict(amp,cpu);
    explicit index(int i0) restrict(amp,cpu); // N==1
    index(int i0, int i1) restrict(amp,cpu); // N==2
    index(int i0, int i1, int i2) restrict(amp,cpu); // N==3
    explicit index(const int components[]) restrict(amp,cpu);

    index& operator=(const index& other) restrict(amp,cpu);

    int operator[](unsigned int c) const restrict(amp,cpu);
    int& operator[](unsigned int c) restrict(amp,cpu);

    template <int N>
        friend bool operator==(const index<N>& lhs, const index<N>& rhs) restrict(amp,cpu);
    template <int N>
        friend bool operator!=(const index<N>& lhs, const index<N>& rhs) restrict(amp,cpu);
    template <int N>
        friend index<N> operator+(const index<N>& lhs,
```

```

1012                                     const index<N>& rhs) restrict(amp,cpu);
1013     template <int N>
1014         friend index<N> operator-(const index<N>& lhs,
1015                                   const index<N>& rhs) restrict(amp,cpu);
1016
1017     index& operator+=(const index& rhs) restrict(amp,cpu);
1018     index& operator-=(const index& rhs) restrict(amp,cpu);
1019
1020     template <int N>
1021         friend index<N> operator+(const index<N>& lhs, int rhs) restrict(amp,cpu);
1022     template <int N>
1023         friend index<N> operator+(int lhs, const index<N>& rhs) restrict(amp,cpu);
1024     template <int N>
1025         friend index<N> operator-(const index<N>& lhs, int rhs) restrict(amp,cpu);
1026     template <int N>
1027         friend index<N> operator-(int lhs, const index<N>& rhs) restrict(amp,cpu);
1028     template <int N>
1029         friend index<N> operator*(const index<N>& lhs, int rhs) restrict(amp,cpu);
1030     template <int N>
1031         friend index<N> operator*(int lhs, const index<N>& rhs) restrict(amp,cpu);
1032     template <int N>
1033         friend index<N> operator/(const index<N>& lhs, int rhs) restrict(amp,cpu);
1034     template <int N>
1035         friend index<N> operator/(int lhs, const index<N>& rhs) restrict(amp,cpu);
1036     template <int N>
1037         friend index<N> operator%(const index<N>& lhs, int rhs) restrict(amp,cpu);
1038     template <int N>
1039         friend index<N> operator%(int lhs, const index<N>& rhs) restrict(amp,cpu);
1040
1041     index& operator+=(int rhs) restrict(amp,cpu);
1042     index& operator-=(int rhs) restrict(amp,cpu);
1043     index& operator*=(int rhs) restrict(amp,cpu);
1044     index& operator/=(int rhs) restrict(amp,cpu);
1045     index& operator%=(int rhs) restrict(amp,cpu);
1046
1047     index& operator++() restrict(amp,cpu);
1048     index& operator++(int) restrict(amp,cpu);
1049     index& operator--() restrict(amp,cpu);
1050     index& operator--(int) restrict(amp,cpu);
1051 };
1052
1053
1054

```

```
template <int N> class index
```

Represents a unique position in N-dimensional space.

#### Template Arguments

*N*

The dimensionality space into which this index applies. Special constructors are supplied for the cases where  $N \in \{1, 2, 3\}$ , but *N* can be any integer greater than 0.

```
static const int rank = N
```

A static member of `index<N>` that contains the rank of this index.

```
typedef int value_type;
```

The element type of `index<N>`.



## 4.1.2 Constructors

**index()** *restrict(amp,cpu)*

Default constructor. The value at each dimension is initialized to zero. Thus, "**index<3> ix;**" initializes the variable to the position (0,0,0).

**index(const index& other)** *restrict(amp,cpu)*

Copy constructor. Constructs a new **index<N>** from the supplied argument "other".

**Parameters:**

*other*

An object of type **index<N>** from which to initialize this new index.

```
explicit index(int i0) restrict(amp,cpu) // N==1
index(int i0, int i1) restrict(amp,cpu) // N==2
index(int i0, int i1, int i2) restrict(amp,cpu) // N==3
```

Constructs an **index<N>** with the coordinate values provided by  $i_0 \dots i_2$ . These are specialized constructors that are only valid when the rank of the index  $N \in \{1,2,3\}$ . Invoking a specialized constructor whose argument count  $\neq N$  will result in a compilation error.

**Parameters:**

$i_0$  [,  $i_1$  [,  $i_2$  ] ]

The component values of the index vector.

**explicit index(const int components[])** *restrict(amp,cpu)*

Constructs an **index<N>** with the coordinate values provided the array of **int** component values. If the coordinate array length  $\neq N$ , the behavior is undefined. If the array value is NULL or not a valid pointer, the behavior is undefined.

**Parameters:**

*components*

An array of **N int** values.

## 4.1.3 Members

**index& operator=(const index& other)** *restrict(amp,cpu)*

Assigns the component values of "other" to this **index<N>** object.

**Parameters:**

*other*

An object of type **index<N>** from which to copy into this index.

**Return Value:**

Returns **\*this**.

```
int operator[](unsigned int c) const restrict(amp,cpu)
int& operator[](unsigned int c) restrict(amp,cpu)
```

Returns the index component value at position **c**.

**Parameters:**

**c**

The dimension axis whose coordinate is to be accessed.

**Return Value:**

A the component value at position **c**.

## 4.1.4 Operators

```
template <int N>
    friend bool operator==(const index<N>& lhs, const index<N>& rhs) restrict(amp,cpu)
template <int N>
    friend bool operator!=(const index<N>& lhs, const index<N>& rhs) restrict(amp,cpu)
```

Compares two objects of **index<N>**.

The expression

$\text{leftIdx} \oplus \text{rightIdx}$

is true if  $\text{leftIdx}[i] \oplus \text{rightIdx}[i]$  for every  $i$  from 0 to  $N-1$ .

**Parameters:**

<i>lhs</i>	The left-hand <b>index&lt;N&gt;</b> to be compared.
<i>rhs</i>	The right-hand <b>index&lt;N&gt;</b> to be compared.

```
template <int N>
    friend index<N> operator+(const index<N>& lhs, const index<N>& rhs) restrict(amp,cpu)
template <int N>
    friend index<N> operator-(const index<N>& lhs, const index<N>& rhs) restrict(amp,cpu)
```

Binary arithmetic operations that produce a new **index<N>** that is the result of performing the corresponding pair-wise binary arithmetic operation on the elements of the operands. The *result* **index<N>** is such that for a given operator  $\oplus$ ,  
 $\text{result}[i] = \text{leftIdx}[i] \oplus \text{rightIdx}[i]$   
 for every  $i$  from 0 to  $N-1$ .

**Parameters:**

<i>lhs</i>	The left-hand <b>index&lt;N&gt;</b> of the arithmetic operation.
<i>rhs</i>	The right-hand <b>index&lt;N&gt;</b> of the arithmetic operation.

```
index& operator+=(const index& rhs) restrict(amp,cpu)
index& operator-=(const index& rhs) restrict(amp,cpu)
```

For a given operator  $\oplus$ , produces the same effect as  
 $(*this) = (*this) \oplus rhs;$

The return value is `"*this"`.

**Parameters:**

<i>rhs</i>	The right-hand <b>index&lt;N&gt;</b> of the arithmetic operation.
------------	---

```
template <int N>
    friend index<N> operator+(const index<N>& idx, int value) restrict(amp,cpu)
template <int N>
    friend index<N> operator+(int value, const index<N>& idx) restrict(amp,cpu)
template <int N>
    friend index<N> operator-(const index<N>& idx, int value) restrict(amp,cpu)
template <int N>
    friend index<N> operator-(int value, const index<N>& idx) restrict(amp,cpu)
template <int N>
    friend index<N> operator*(const index<N>& idx, int value) restrict(amp,cpu)
template <int N>
    friend index<N> operator*(int value, const index<N>& idx) restrict(amp,cpu)
template <int N>
    friend index<N> operator/(const index<N>& idx, int value) restrict(amp,cpu)
template <int N>
    friend index<N> operator/(int value, const index<N>& idx) restrict(amp,cpu)
template <int N>
    friend index<N> operator%(const index<N>& idx, int value) restrict(amp,cpu)
template <int N>
    friend index<N> operator%(int value, const index<N>& idx) restrict(amp,cpu)
```

Binary arithmetic operations that produce a new **index<N>** that is the result of performing the corresponding binary arithmetic operation on the elements of the index operands. The *result* **index<N>** is such that for a given operator  $\oplus$ ,  
 $\text{result}[i] = \text{idx}[i] \oplus \text{value}$   
 or  
 $\text{result}[i] = \text{value} \oplus \text{idx}[i]$

for every  $i$  from 0 to  $N-1$ .

**Parameters:**

<i>idx</i>	The <b>index&lt;N&gt;</b> operand
<i>value</i>	The integer operand

```

index& operator+=(int value) restrict(amp,cpu)
index& operator-=(int value) restrict(amp,cpu)
index& operator*=(int value) restrict(amp,cpu)
index& operator/=(int value) restrict(amp,cpu)
index& operator%=(int value) restrict(amp,cpu)

```

For a given operator  $\oplus$ , produces the same effect as

$(*this) = (*this) \oplus value;$

The return value is `"*this"`.

**Parameters:**

<code>value</code>	The right-hand <code>int</code> of the arithmetic operation.
--------------------	--

```

index& operator++() restrict(amp,cpu)
index operator++(int) restrict(amp,cpu)
index& operator--() restrict(amp,cpu)
index operator--(int) restrict(amp,cpu)

```

For a given operator  $\oplus$ , produces the same effect as

$(*this) = (*this) \oplus 1;$

For prefix increment and decrement, the return value is `"*this"`. Otherwise a new `index<N>` is returned.

## 4.2 extent<N>

The `extent<N>` type represents an N-dimensional vector of `int` which specifies the bounds of an N-dimensional space with an origin of 0. The values in the coordinate vector are ordered from most-significant to least-significant. Thus, in Cartesian 3-dimensional space, where a common convention exists that the Z dimension (plane) is most significant, the Y dimension (row) is second in significance and the X dimension (column) is the least significant, the extent vector (7,5,3) represents a space where the Z coordinate ranges from 0 to 6, the Y coordinate ranges from 0 to 4, and the X coordinate ranges from 0 to 2.

### 4.2.1 Synopsis

```

template <int N>
class extent {
public:
    static const int rank = N;
    typedef int value_type;

    extent() restrict(amp,cpu);
    extent(const extent& other) restrict(amp,cpu);
    explicit extent(int e0) restrict(amp,cpu); // N==1
    extent(int e0, int e1) restrict(amp,cpu); // N==2
    extent(int e0, int e1, int e2) restrict(amp,cpu); // N==3
    explicit extent(const int components[]) restrict(amp,cpu);

    extent& operator=(const extent& other) restrict(amp,cpu);

    int operator[](unsigned int c) const restrict(amp,cpu);
    int& operator[](unsigned int c) restrict(amp,cpu);

    int size() const restrict(amp,cpu);

    bool contains(const index<N>& idx) const restrict(amp,cpu);

    template <int D0>
        tiled_extent<D0> tile() const;

```

```

1111     template <int D0, int D1>         tiled_extent<D0,D1> tile() const;
1112     template <int D0, int D1, int D2> tiled_extent<D0,D1,D2> tile() const;
1113
1114     extent operator+(const index<N>& idx) restrict(amp,cpu);
1115     extent operator-(const index<N>& idx) restrict(amp,cpu);
1116
1117     template <int N>
1118         friend bool operator==(const extent<N>& lhs, const extent<N>& rhs) restrict(amp,cpu);
1119     template <int N>
1120         friend bool operator!=(const extent<N>& lhs, const extent<N>& rhs) restrict(amp,cpu);
1121
1122     template <int N>
1123         friend extent<N> operator+(const extent<N>& lhs, int rhs) restrict(amp,cpu);
1124     template <int N>
1125         friend extent<N> operator+(int lhs, const extent<N>& rhs) restrict(amp,cpu);
1126     template <int N>
1127         friend extent<N> operator-(const extent<N>& lhs, int rhs) restrict(amp,cpu);
1128     template <int N>
1129         friend extent<N> operator-(int lhs, const extent<N>& rhs) restrict(amp,cpu);
1130     template <int N>
1131         friend extent<N> operator*(const extent<N>& lhs, int rhs) restrict(amp,cpu);
1132     template <int N>
1133         friend extent<N> operator*(int lhs, const extent<N>& rhs) restrict(amp,cpu);
1134     template <int N>
1135         friend extent<N> operator/(const extent<N>& lhs, int rhs) restrict(amp,cpu);
1136     template <int N>
1137         friend extent<N> operator/(int lhs, const extent<N>& rhs) restrict(amp,cpu);
1138     template <int N>
1139         friend extent<N> operator%(const extent<N>& lhs, int rhs) restrict(amp,cpu);
1140     template <int N>
1141         friend extent<N> operator%(int lhs, const extent<N>& rhs) restrict(amp,cpu);
1142
1143     extent& operator+=(int rhs) restrict(amp,cpu);
1144     extent& operator-=(int rhs) restrict(amp,cpu);
1145     extent& operator*=(int rhs) restrict(amp,cpu);
1146     extent& operator/=(int rhs) restrict(amp,cpu);
1147     extent& operator%=(int rhs) restrict(amp,cpu);
1148
1149     extent& operator++() restrict(amp,cpu);
1150     extent operator++(int) restrict(amp,cpu);
1151     extent& operator--() restrict(amp,cpu);
1152     extent operator--(int) restrict(amp,cpu);
1153 };
1154
1155

```

```
template <int N> class extent
```

Represents a unique position in N-dimensional space.

#### Template Arguments

<i>N</i>	The dimension to this extent applies. Special constructors are supplied for the cases where $N \in \{1, 2, 3\}$ , but $N$ can be any integer greater than or equal to 1.
----------	--

```
static const int rank = N
```

A static member of `extent<N>` that contains the rank of this extent.

```
typedef int value_type;
```

The element type of `extent<N>`.

## 4.2.2 Constructors

```
extent() restrict(amp,cpu);
```

Default constructor. The value at each dimension is initialized to zero. Thus, "`extent<3> ix;`" initializes the variable to the position (0,0,0).

**Parameters:**

None.

```
extent(const extent& other) restrict(amp,cpu)
```

Copy constructor. Constructs a new `extent<N>` from the supplied argument `ix`.

**Parameters:**

*other*

An object of type `extent<N>` from which to initialize this new extent.

```
explicit extent(int e0) restrict(amp,cpu) // N==1
extent(int e0, int e1) restrict(amp,cpu) // N==2
extent(int e0, int e1, int e2) restrict(amp,cpu) // N==3
```

Constructs an `extent<N>` with the coordinate values provided by `e0..2`. These are specialized constructors that are only valid when the rank of the extent  $N \in \{1,2,3\}$ . Invoking a specialized constructor whose argument count  $\neq N$  will result in a compilation error.

**Parameters:**

*e0 [, e1 [, e2 ]]*

The component values of the extent vector.

```
explicit extent(const int components[]) restrict(amp,cpu);
```

Constructs an `extent<N>` with the coordinate values provided the array of `int` component values. If the coordinate array length  $\neq N$ , the behavior is undefined. If the array value is NULL or not a valid pointer, the behavior is undefined.

**Parameters:**

*components*

An array of  $N$  `int` values.

## 4.2.3 Members

```
extent& operator=(const extent& other) restrict(amp,cpu)
```

Assigns the component values of "other" to this `extent<N>` object.

**Parameters:**

*other*

An object of type `extent<N>` from which to copy into this extent.

**Return Value:**

Returns `*this`.

```
int operator[](unsigned int c) const restrict(amp,cpu)
int& operator[](unsigned int c) restrict(amp,cpu)
```

Returns the extent component value at position `c`.

**Parameters:**

*c*

The dimension axis whose coordinate is to be accessed.

**Return Value:**

A the component value at position `c`.

```
bool contains(const index<N>& idx) const restrict(amp,cpu)
```

Tests whether the index "idx" is properly contained within this extent (with an assumed origin of zero).

**Parameters:**

*idx*

An object of type `index<N>`

**Return Value:**

Returns `true` if the "idx" is contained within the space defined by this extent (with an assumed origin of zero).

```
int size() const restrict(amp,cpu)
```

This member function returns the total linear size of this `extent<N>` (in units of elements), which is computed as:

```
extent[0] * extent[1] ... * extent[N-1]
```

```
template <int D0>          tiled_extent<D0> tile() const restrict(amp,cpu)
template <int D0, int D1>  tiled_extent<D0,D1> tile() const restrict(amp,cpu)
template <int D0, int D1, int D2> tiled_extent<D0,D1,D2> tile() const restrict(amp,cpu)
```

Produces a `tiled_extent` object with the tile extents given by `D0`, `D1`, and `D2`.

`tile<D0,D1,D2>()` is only supported on `extent<3>`. It will produce a compile-time error if used on an `extent` where  $N \neq 3$ .

`tile<D0,D1>()` is only supported on `extent<2>`. It will produce a compile-time error if used on an `extent` where  $N \neq 2$ .

`tile<D0>()` is only supported on `extent<1>`. It will produce a compile-time error if used on an `extent` where  $N \neq 1$ .

#### 4.2.4 Operators

```
template <int N>
    friend bool operator==(const extent<N>& lhs, const extent<N>& rhs) restrict(amp,cpu)
template <int N>
    friend bool operator!=(const extent<N>& lhs, const extent<N>& rhs) restrict(amp,cpu)
```

Compares two objects of `extent<N>`.

The expression

`leftExt  $\oplus$  rightExt`

is true if `leftExt[i]  $\oplus$  rightExt[i]` for every  $i$  from 0 to  $N-1$ .

##### Parameters:

<i>lhs</i>	The left-hand <code>extent&lt;N&gt;</code> to be compared.
<i>rhs</i>	The right-hand <code>extent&lt;N&gt;</code> to be compared.

```
extent<N> operator+(const index<N>& idx) restrict(amp,cpu)
extent<N> operator-(const index<N>& idx) restrict(amp,cpu)
```

Adds (or subtracts) an object of type `index<N>` from this extent to form a new extent. The *result* `extent<N>` is such that for a given operator  $\oplus$ ,

`result[i] = this[i]  $\oplus$  idx[i]`

##### Parameters:

<i>idx</i>	The right-hand <code>index&lt;N&gt;</code> to be added or subtracted.
------------	---

```
template <int N>
    friend extent<N> operator+(const extent<N>& ext, int value) restrict(amp,cpu)
template <int N>
    friend extent<N> operator+(int value, const extent<N>& ext) restrict(amp,cpu)
template <int N>
    friend extent<N> operator-(const extent<N>& ext, int value) restrict(amp,cpu)
template <int N>
    friend extent<N> operator-(int value, const extent<N>& ext) restrict(amp,cpu)
template <int N>
    friend extent<N> operator*(const extent<N>& ext, int value) restrict(amp,cpu)
template <int N>
    friend extent<N> operator*(int value, const extent<N>& ext) restrict(amp,cpu)
template <int N>
    friend extent<N> operator/(const extent<N>& ext, int value) restrict(amp,cpu)
template <int N>
    friend extent<N> operator/(int value, const extent<N>& ext) restrict(amp,cpu)
template <int N>
    friend extent<N> operator%(const extent<N>& ext, int value) restrict(amp,cpu)
template <int N>
    friend extent<N> operator%(int value, const extent<N>& ext) restrict(amp,cpu)
```

Binary arithmetic operations that produce a new **extent<N>** that is the result of performing the corresponding binary arithmetic operation on the elements of the extent operands. The *result extent<N>* is such that for a given operator  $\oplus$ ,  
 $result[i] = ext[i] \oplus value$   
or  
 $result[i] = value \oplus ext[i]$   
for every  $i$  from 0 to  $N-1$ .

**Parameters:**

<i>ext</i>	The <b>extent&lt;N&gt;</b> operand
<i>value</i>	The integer operand

```
extent& operator+=(int value) restrict(amp,cpu)
extent& operator-=(int value) restrict(amp,cpu)
extent& operator*=(int value) restrict(amp,cpu)
extent& operator/=(int value) restrict(amp,cpu)
extent& operator%=(int value) restrict(amp,cpu)
```

For a given operator  $\oplus$ , produces the same effect as  
 $(*this) = (*this) \oplus value$

The return value is `"*this"`.

**Parameters:**

<i>Value</i>	The right-hand <b>int</b> of the arithmetic operation.
--------------	--

```
extent& operator++() restrict(amp,cpu)
extent operator++(int) restrict(amp,cpu)
extent& operator--() restrict(amp,cpu)
extent operator--(int) restrict(amp,cpu)
```

For a given operator  $\oplus$ , produces the same effect as  
 $(*this) = (*this) \oplus 1$

For prefix increment and decrement, the return value is `"*this"`. Otherwise a new **extent<N>** is returned.

### 4.3 tiled\_extent<D0,D1,D2>

A *tiled\_extent* is an extent of 1 to 3 dimensions which also subdivides the index space into 1-, 2-, or 3-dimensional tiles. It has three specialized forms: *tiled\_extent<D0>*, *tiled\_extent<D0,D1>*, and *tiled\_extent<D0,D1,D2>*, where  $D_{0-2}$  specify the positive length of the tile along each dimension, with *D0* being the most-significant dimension and *D2* being the least-significant. Partial template specializations are provided to represent 2-D and 1-D tiled extents.

A *tiled\_extent* can be formed from an extent by calling *extent<N>::tile<D0,D1,D2>()* or one of the other two specializations of *extent<N>::tile()*.

A *tiled\_extent* inherits from *extent*, thus all public members of *extent* are available on *tiled\_extent*.

#### 4.3.1 Synopsis

```
template <int D0, int D1=0, int D2=0>
class tiled_extent : public extent<3>
{
public:
    static const int rank = 3;
```

```

1203 tiled_extent() restrict(amp,cpu);
1204 tiled_extent(const tiled_extent& other) restrict(amp,cpu);
1205 tiled_extent(const extent<3>& extent) restrict(amp,cpu);
1206
1207 tiled_extent& operator=(const tiled_extent& other) restrict(amp,cpu);
1208
1209 tiled_extent pad() const restrict(amp,cpu);
1210 tiled_extent truncate() const restrict(amp,cpu);
1211
1212 __declspec(property(get)) extent<3> tile_extent;
1213
1214 static const int tile_dim0 = D0;
1215 static const int tile_dim1 = D1;
1216 static const int tile_dim2 = D2;
1217
1218 friend bool operator==(const tiled_extent& lhs,
1219                        const tiled_extent& rhs) restrict(amp,cpu);
1220 friend bool operator!=(const tiled_extent& lhs,
1221                        const tiled_extent& rhs) restrict(amp,cpu);
1222 };
1223
1224
1225 template <int D0, int D1>
1226 class tiled_extent<D0,D1,0> : public extent<2>
1227 {
1228 public:
1229     static const int rank = 2;
1230
1231     tiled_extent() restrict(amp,cpu);
1232     tiled_extent(const tiled_extent& other) restrict(amp,cpu);
1233     tiled_extent(const extent<2>& extent) restrict(amp,cpu);
1234
1235     tiled_extent& operator=(const tiled_extent& other) restrict(amp,cpu);
1236
1237     tiled_extent pad() const restrict(amp,cpu);
1238     tiled_extent truncate() const restrict(amp,cpu);
1239
1240     __declspec(property(get)) extent<2> tile_extent;
1241
1242     static const int tile_dim0 = D0;
1243     static const int tile_dim1 = D1;
1244
1245     friend bool operator==(const tiled_extent& lhs,
1246                          const tiled_extent& rhs) restrict(amp,cpu);
1247     friend bool operator!=(const tiled_extent& lhs,
1248                          const tiled_extent& rhs) restrict(amp,cpu);
1249 };
1250
1251 template <int D0>
1252 class tiled_extent<D0,0,0> : public extent<1>
1253 {
1254 public:
1255     static const int rank = 1;
1256
1257     tiled_extent() restrict(amp,cpu);
1258     tiled_extent(const tiled_extent& other) restrict(amp,cpu);
1259     tiled_extent(const extent<1>& extent) restrict(amp,cpu);
1260

```



```

1261 tiled_extent& operator=(const tiled_extent& other) restrict(amp,cpu);
1262
1263 tiled_extent pad() const restrict(amp,cpu);
1264 tiled_extent truncate() const restrict(amp,cpu);
1265
1266 __declspec(property(get)) extent<1> tile_extent;
1267
1268 static const int tile_dim0 = D0;
1269
1270 friend bool operator==(const tiled_extent& lhs,
1271                        const tiled_extent& rhs) restrict(amp,cpu);
1272 friend bool operator!=(const tiled_extent& lhs,
1273                        const tiled_extent& rhs) restrict(amp,cpu);
1274 };
1275
1276
1277

```

```

template <int D0, int D1=0, int D2=0> class tiled_extent
template <int D0, int D1>           class tiled_extent<D0,D1,0>
template <int D0>                   class tiled_extent<D0,0,0>

```

Represents an extent subdivided into 1-, 2-, or 3-dimensional tiles.

#### Template Arguments

*D0, D1, D2*

The length of the tile in each specified dimension, where D0 is the most-significant dimension and D2 is the least-significant.

1278

```
static const int rank = N
```

A static member of `tiled_extent` that contains the rank of this tiled extent, and is either 1, 2, or 3 depending on the specialization used.

1279

### 1280 4.3.2 Constructors

1281

```
tiled_extent() restrict(amp,cpu)
```

Default constructor. The origin and extent is default-constructed and thus zero.

#### Parameters:

None.

1282

```
tiled_extent(const tiled_extent& other) restrict(amp,cpu)
```

Copy constructor. Constructs a new `tiled_extent` from the supplied argument "other".

#### Parameters:

*other*

An object of type `tiled_extent` from which to initialize this new extent.

1283

```
tiled_extent(const extent<N>& extent) restrict(amp,cpu)
```

Constructs a `tiled_extent<N>` with the extent "extent". The origin is default-constructed and thus zero. Notice that this constructor allows implicit conversions from `extent<N>` to `tiled_extent<N>`.

#### Parameters:

*extent*

The extent of this `tiled_extent`

1284

### 1285 4.3.3 Members

1286

```
tiled_extent& operator=(const tiled_extent& other) restrict(amp,cpu)
```

Assigns the component values of "other" to this `tiled_extent<N>` object.

#### Parameters:

*Other*

An object of type `tiled_extent<N>` from which to copy into this.

#### Return Value:

Returns `*this`.

1287

```
 tiled_extent pad() const restrict(amp,cpu)
```

Returns a new tiled\_extent with the extents adjusted up to be evenly divisible by the tile dimensions. The origin of the new tiled\_extent is the same as the origin of this one.

1288

```
 tiled_extent truncate() const restrict(amp,cpu)
```

Returns a new tiled\_extent with the extents adjusted down to be evenly divisible by the tile dimensions. The origin of the new tiled\_extent is the same as the origin of this one.

1289

```
 __declspec(property(get)) extent<N> tile_extent
```

Returns an instance of an extent<N> that captures the values of the tiled\_extent template arguments D0, D1, and D2. For example:

```
 tiled_extent<64,16,4> tg;
 extent<3> myTileExtent = tg.tile_extent;
 assert(myTileExtent[0] == 64);
 assert(myTileExtent[1] == 16);
 assert(myTileExtent[2] == 4);
```

1290

```
 static const int tile_dim0
 static const int tile_dim1
 static const int tile_dim2
```

These constants allow access to the template arguments of tiled\_extent.

1291

#### 1292 4.3.4 Operators

1293

```
 friend bool operator==(const tiled_extent& lhs,
                        const tiled_extent& rhs) restrict(amp,cpu)
 friend bool operator!=(const tiled_extent& lhs,
                        const tiled_extent& rhs) restrict(amp,cpu)
```

Compares two objects of tiled\_extent<N>.

The expression

lhs  $\oplus$  rhs

is true if lhs.extent  $\oplus$  rhs.extent and lhs.origin  $\oplus$  rhs.origin.

##### Parameters:

lhs	The left-hand tiled_extent to be compared.
rhs	The right-hand tiled_extent to be compared.

1294

1295

#### 1296 4.4 tiled\_index<D0,D1,D2>

1297

1298 A tiled\_index is a set of indices of 1 to 3 dimensions which have been subdivided into 1-, 2-, or 3-dimensional tiles in a  
 1299 tiled\_extent. It has three specialized forms: tiled\_index<D0>, tiled\_index<D0,D1>, and tiled\_index<D0,D1,D2>, where  $D_{0-2}$   
 1300 specify the length of the tile along each dimension, with  $D0$  being the most-significant dimension and  $D2$  being the least-  
 1301 significant. Partial template specializations are provided to represent 2-D and 1-D tiled indices.

1302

1303 A tiled\_index is implicitly convertible to an index<N>, where the implicit index represents the global index.

1304

1305 A tiled\_index contains 4 member indices which are related to each other mathematically and help the user to pinpoint a  
 1306 global index to an index within a tiled space.

1307

1308 A tiled\_index contains a global index into an extent space. The other indices obey the following relations:

1309

```

1310 .local ≡ .global % (D0,D1,D2)
1311 .tile ≡ .global / (D0,D1,D2)
1312 .tile_origin ≡ .global - .local
1313


```

This is shown visually in the following example:

```

1315 parallel_for_each(extent<2>(20,24).tile<5,4>(),
1316                   [&](tiled_index<5,4> ti) { /* ... */ });
1317

```



1. Each cell in the diagram represents one thread which is scheduled by the `parallel_for_each` call. We see that, as with the non-tiled `parallel_for_each`, the number of threads scheduled is given by the extent parameter to the `parallel_for_each` call.
2. Using vector notation, we see that the total number of tiles scheduled is  $\langle 20, 24 \rangle / \langle 5, 4 \rangle = \langle 4, 6 \rangle$ , which we see in the above diagram as 4 tiles along the vertical axis, and 6 tiles along the horizontal axis.
3. The tile in red is tile number  $\langle 0, 0 \rangle$ . The tile in yellow is tile number  $\langle 1, 2 \rangle$ .
4. The thread in blue:
  - a. has a global id of  $\langle 5, 8 \rangle$
  - b. Has a local id  $\langle 0, 0 \rangle$  within its tile. i.e., it lies on the origin of the tile.
5. The thread in green:
  - a. has a global id of  $\langle 6, 9 \rangle$
  - b. has a local id of  $\langle 1, 1 \rangle$  within its tile
  - c. The blue thread (number  $\langle 5, 8 \rangle$ ) is the green thread's tile origin.

#### 4.4.1 Synopsis

```

1334 template <int D0, int D1=0, int D2=0>
1335 class tiled_index
1336 {
1337

```

```

1339 public:
1340     static const int rank = 3;
1341
1342     const index<3> global;
1343     const index<3> local;
1344     const index<3> tile;
1345     const index<3> tile_origin;
1346     const tile_barrier barrier;
1347
1348     tiled_index(const index<3>& global,
1349                const index<3> local,
1350                const index<3> tile,
1351                const index<3> tile_origin,
1352                const tile_barrier& barrier) restrict(amp,cpu);
1353     tiled_index(const tiled_index& other) restrict(amp,cpu);
1354
1355     operator const index<3>() const restrict(amp,cpu);
1356
1357     __declspec(property(get)) extent<3> tile_extent;
1358
1359     static const int tile_dim0 = D0;
1360     static const int tile_dim1 = D1;
1361     static const int tile_dim2 = D2;
1362 };
1363
1364 template <int D0, int D1>
1365 class tiled_index<D0,D1,0>
1366 {
1367 public:
1368     static const int rank = 2;
1369
1370     const index<2> global;
1371     const index<2> local;
1372     const index<2> tile;
1373     const index<2> tile_origin;
1374     const tile_barrier barrier;
1375
1376     tiled_index(const index<2>& global,
1377                const index<2> local,
1378                const index<2> tile,
1379                const index<2> tile_origin,
1380                const tile_barrier& barrier) restrict(amp,cpu);
1381     tiled_index(const tiled_index& other) restrict(amp,cpu);
1382
1383     operator const index<2>() const restrict(amp,cpu);
1384
1385     __declspec(property(get)) extent<2> tile_extent;
1386
1387     static const int tile_dim0 = D0;
1388     static const int tile_dim1 = D1;
1389 };
1390
1391 template <int D0>
1392 class tiled_index<D0,0,0>
1393 {
1394 public:
1395     static const int rank = 1;

```

```

1397
1398     const index<1> global;
1399     const index<1> local;
1400     const index<1> tile;
1401     const index<1> tile_origin;
1402     const tile_barrier barrier;
1403
1404     tiled_index(const index<1>& global,
1405               const index<1> local,
1406               const index<1> tile,
1407               const index<1> tile_origin,
1408               const tile_barrier& barrier) restrict(amp,cpu);
1409     tiled_index(const tiled_index& other) restrict(amp,cpu);
1410
1411     operator const index<1>() const restrict(amp,cpu);
1412
1413     __declspec(property(get)) extent<1> tile_extent;
1414
1415     static const int tile_dim0 = D0;
1416 };
1417
1418
1419

```

```

template <int D0, int D1=0, int D2=0> class tiled_index
template <int D0, int D1>           class tiled_index<D0,D1,0>
template <int D0 >                 class tiled_index<D0,0,0>

```

Represents a set of related indices subdivided into 1-, 2-, or 3-dimensional tiles.

#### Template Arguments

<i>D0, D1, D2</i>	The length of the tile in each specified dimension, where D0 is the most-significant dimension and D2 is the least-significant.
-------------------	---

1420

```
static const int rank = N
```

A static member of `tiled_index` that contains the rank of this tiled extent, and is either 1, 2, or 3 depending on the specialization used.

1421

## 1422 4.4.2 Constructors

1423

1424 The `tiled_index` class has no default constructor.

1425

```

tiled_index(const index<N>& global,
           const index<N>& local,
           const index<N>& tile,
           const index<N>& tile_origin,
           const tile_barrier& barrier) restrict(amp,cpu)

```

Construct a new `tiled_index` out of the constituent indices.

Note that it is permissible to create a `tiled_index` instance for which the geometric identities which are guaranteed for system-created tiled indices, which are passed as a kernel parameter to the tiled overloads of `parallel_for_each`, do not hold. In such cases, it is up to the application to assign application-specific meaning to the member indices of the instance.

#### Parameters:

<i>global</i>	An object of type <code>index&lt;N&gt;</code> which is taken to be the global index of this tile.
<i>local</i>	An object of type <code>index&lt;N&gt;</code> which is taken to be the local index within this tile.
<i>tile</i>	An object of type <code>index&lt;N&gt;</code> which is taken to be the coordinates of the current tile.
<i>tile_origin</i>	An object of type <code>index&lt;N&gt;</code> which is taken to be the global index of the

	top-left corner of the tile.
<i>barrier</i>	An object of type <code>tile_barrier</code> .

1426

<b>tiled_index(const tiled_index&amp; other) restrict(amp,cpu)</b>	
Copy constructor. Constructs a new <code>tiled_index</code> from the supplied argument "other".	
<b>Parameters:</b>	
<i>other</i>	An object of type <code>tiled_index</code> from which to initialize this.

1427

#### 1428 4.4.3 Members

1429

<b>const index&lt;N&gt; global</b>	
An index of rank 1, 2, or 3 that represents the global index within an extent.	

1430

<b>const index&lt;N&gt; local</b>	
An index of rank 1, 2, or 3 that represents the relative index within the current tile of a tiled extent.	

1431

<b>const index&lt;N&gt; tile</b>	
An index of rank 1, 2, or 3 that represents the coordinates of the current tile of a tiled extent.	

1432

<b>const index&lt;N&gt; tile_origin</b>	
An index of rank 1, 2, or 3 that represents the global coordinates of the origin of the current tile within a tiled extent.	

1433

<b>const tile_barrier barrier</b>	
An object which represents a barrier within the current tile of threads.	

1434

<b>operator const index&lt;N&gt;() const restrict(amp,cpu)</b>	
Implicit conversion operator that converts a <code>tiled_index&lt;D0,D1,D2&gt;</code> into an <code>index&lt;N&gt;</code> . The implicit conversion converts to the <code>.global</code> index member.	

1435

<b>__declspec(property(get)) extent&lt;N&gt; tile_extent</b>	
Returns an instance of an <code>extent&lt;N&gt;</code> that captures the values of the <code>tiled_index</code> template arguments D0, D1, and D2. For example:	
<pre> index&lt;3&gt; zero; tiled_index&lt;64,16,4&gt; ti(index&lt;3&gt;(256,256,256), zero, zero, zero, mybarrier); extent&lt;3&gt; myTileExtent = ti.tile_extent; assert(myTileExtent.tile_dim0 == 64); assert(myTileExtent.tile_dim1 == 16); assert(myTileExtent.tile_dim2 == 4); </pre>	

1436

<b>static const int tile_dim0</b> <b>static const int tile_dim1</b> <b>static const int tile_dim2</b>	
These constants allow access to the template arguments of <code>tiled_index</code> .	

1437

#### 1438 4.5 tile\_barrier

1439

1440 The `tile_barrier` class is a capability class that is only creatable by the system, and passed to a tiled `parallel_for_each`  
1441 function object as part of the `tiled_index` parameter. It provides member functions, such as `wait`, whose purpose is to  
1442 synchronize execution of threads running within the thread tile.

1443

1444 A call to `wait` shall not occur in non-uniform code within a thread tile. Section 8 defines uniformity and lack thereof  
1445 formally.

#### 1446 4.5.1 Synopsis

```

1447
1448 class tile_barrier
1449 {
1450 public:
1451     tile_barrier(const tile_barrier& other) restrict(amp,cpu);
1452
1453     void wait() restrict(amp);
1454     void wait_with_all_memory_fence() restrict(amp);
1455     void wait_with_global_memory_fence() restrict(amp);
1456     void wait_with_tile_static_memory_fence() restrict(amp);
1457 };
1458

```

#### 1459 4.5.2 Constructors

1460 The `tile_barrier` class does not have a public default constructor, only a copy-constructor.

1461 `tile_barrier(const tile_barrier& other) restrict(amp,cpu)`

1462 Copy constructor. Constructs a new `tile_barrier` from the supplied argument "other".

##### Parameters:

<i>other</i>	An object of type <code>tile_barrier</code> from which to initialize this.
--------------	--

#### 1464 4.5.3 Members

1465 The `tile_barrier` class does not have an assignment operator. Section 8 provides a complete description of the C++ AMP memory model, of which class `tile_barrier` is an important part.

1466 `void wait() restrict(amp)`

1467 Blocks execution of all threads in the thread tile until all threads in the tile have reached this call. Establishes a memory fence on all `tile_static` and global memory operations executed by the threads in the tile such that all memory operations issued prior to hitting the barrier are visible to all other threads after the barrier has completed and none of the memory operations occurring after the barrier are executed before hitting the barrier. This is identical to [wait\\_with\\_all\\_memory\\_fence](#).

1468 `void wait_with_all_memory_fence() restrict(amp)`

1469 Blocks execution of all threads in the thread tile until all threads in the tile have reached this call. Establishes a memory fence on all `tile_static` and global memory operations executed by the threads in the tile such that all memory operations issued prior to hitting the barrier are visible to all other threads after the barrier has completed and none of the memory operations occurring after the barrier are executed before hitting the barrier. This is identical to [wait](#).

1470 `void wait_with_global_memory_fence() restrict(amp)`

1471 Blocks execution of all threads in the thread tile until all threads in the tile have reached this call. Establishes a memory fence on global memory operations (but not `tile-static` memory operations) executed by the threads in the tile such that all global memory operations issued prior to hitting the barrier are visible to all other threads after the barrier has completed and none of the global memory operations occurring after the barrier are executed before hitting the barrier.

1472 `void wait_with_tile_static_memory_fence() restrict(amp)`

1473 Blocks execution of all threads in the thread tile until all threads in the tile have reached this call. Establishes a memory fence on `tile-static` memory operations (but not global memory operations) executed by the threads in the tile such that all global memory operations issued prior to hitting the barrier are visible to all other threads after the barrier has completed and none of the `tile-static` memory operations occurring after the barrier are executed before hitting the barrier.

#### 1474 4.5.4 Other Memory Fences and Barriers

1475 C++ AMP provides functions that serve as memory fences, which establish a happens-before relationship between memory  
 1476 operations performed by threads within the same thread tile. These functions are available in the concurrency namespace.  
 1477 Section 8 provides a complete description of the C++ AMP memory model.  
 1478

```
void all_memory_fence(const tile_barrier&) restrict(amp)
```

Establishes a thread-tile scoped memory fence for both global and tile-static memory operations. This function does not imply a barrier and is therefore permitted in divergent code.

```
void global_memory_fence(const tile_barrier&) restrict(amp)
```

Establishes a thread-tile scoped memory fence for global (but not tile-static) memory operations. This function does not imply a barrier and is therefore permitted in divergent code.

```
void tile_static_memory_fence(const tile_barrier&) restrict(amp)
```

Establishes a thread-tile scoped memory fence for tile-static (but not global) memory operations. This function does not imply a barrier and is therefore permitted in divergent code.

## 1483 4.6 completion\_future

1484 This class is the return type of all C++ AMP asynchronous APIs and has an interface analogous to `std::shared_future<void>`.  
 1485 Similar to `std::shared_future`, this type provides member methods such as ***wait*** and ***get*** to wait for C++ AMP asynchronous  
 1486 operations to finish, and the type additionally provides a member method ***then***, to specify a completion callback *functor* to  
 1487 be executed upon completion of a C++ AMP asynchronous operation. Further this type also contains a member method  
 1488 ***to\_task*** (Microsoft specific extension) which returns a `concurrency::task` object which can be used to avail the capabilities of  
 1489 PPL tasks with C++ AMP asynchronous operations; viz. chaining continuations, cancellation etc. This essentially enables  
 1490 “wait-free” composition of C++ AMP asynchronous tasks on accelerators with CPU tasks.

### 1491 4.6.1 Synopsis

```
1492 class completion_future
1493 {
1494 public:
1495     completion_future();
1496     completion_future(const completion_future& _Other);
1497     completion_future(completion_future&& _Other);
1498     ~completion_future();
1499     completion_future& operator=(const completion_future& _Other);
1500     completion_future& operator=(completion_future&& _Other);
1501
1502     void get() const;
1503
1504     bool valid() const;
1505
1506     void wait() const;
1507     template <class _Rep, class _Period>
1508     std::future_status::future_status wait_for(const std::chrono::duration<_Rep, _Period>&
1509     _Rel_time) const;
1510     template <class _Clock, class _Duration>
1511     std::future_status::future_status wait_until(const std::chrono::time_point<_Clock,
1512     _Duration>& _Abs_time) const;
1513
1514     operator std::shared_future<void>() const;
```



```

1518     void then(const _Functor &_Func) const;
1519
1520     concurrency::task<void> to_task() const;
1521 };

```

## 1522 4.6.2 Constructors

1523

### completion\_future()

Default constructor. Constructs an empty uninitialized completion\_future object which does not refer to any asynchronous operation. Default constructed completion\_future objects have `valid() == false`

1524

### completion\_future (const completion\_future& other)

Copy constructor. Constructs a new completion\_future object that refers to the same asynchronous operation as the other completion\_future object.

#### Parameters:

other	An object of type completion_future from which to initialize this.
-------	--

1525

1526

1527

### completion\_future (completion\_future&& other)

Move constructor. Move constructs a new completion\_future object that refers to the same asynchronous operation as originally referred by the other completion\_future object. After this constructor returns, `other.valid() == false`

#### Parameters:

other	An object of type completion_future which the new completion_future object is to be move constructed from.
-------	--

1528

### completion\_future& operator=(const completion\_future& other)

Copy assignment. Copy assigns the contents of other to this. This method causes this to stop referring its current asynchronous operation and start referring the same asynchronous operation as other.

#### Parameters:

other	An object of type completion_future which is copy assigned to this.
-------	---

1529

### completion\_future& operator=(completion\_future&& other)

Move assignment. Move assigns the contents of other to this. This method causes this to stop referring its current asynchronous operation and start referring the same asynchronous operation as other. After this method returns, `other.valid() == false`

#### Parameters:

other	An object of type completion_future which is move assigned to this.
-------	---

1530

## 1531 4.6.3 Members

1532

1533

### void get() const

This method is functionally identical to `std::shared_future<void>::get`. This method waits for the associated asynchronous operation to finish and returns only upon the completion of the asynchronous operation. If an exception was encountered during the execution of the asynchronous operation, this method throws that stored exception.

1534

### bool valid() const

This method is functionally identical to `std::shared_future<void>::valid`. This returns true if this completion\_future is associated with an asynchronous operation.

1535

### void wait() const

```

template <class Rep, class Period>
std::future_status::future_status wait_for(const std::chrono::duration<Rep, Period>&
rel_time) const

```

```
template <class Clock, class Duration>
std::future_status::future_status wait_until(const std::chrono::time_point<Clock, Duration>&
abs_time) const
```

These methods are functionally identical to the corresponding `std::shared_future<void>` methods.

The `wait` method waits for the associated asynchronous operation to finish and returns only upon completion of the associated asynchronous operation or if an exception was encountered when executing the asynchronous operation.

The other variants are functionally identical to the `std::shared_future<void>` member methods with same names.

```
operator shared_future<void>() const
```

Conversion operator to `std::shared_future<void>`. This method returns a `shared_future<void>` object corresponding to `this` `completion_future` object and refers to the same asynchronous operation.

```
template <typename Functor>
void then(const Functor &func) const
```

This method enables specification of a completion callback `func` which is executed upon completion of the asynchronous operation associated with `this` `completion_future` object. The completion callback `func` should have an `operator()` that is valid when invoked with non arguments, i.e., "`func()`".

#### Parameters:

<code>func</code>	A function object or lambda whose <code>operator()</code> is invoked upon completion of <code>this</code> 's associated asynchronous operation.
-------------------	---

```
concurrency::task<void> to_task() const
```

This method returns a `concurrency::task<void>` object corresponding to `this` `completion_future` object and refers to the same asynchronous operation. This method is a Microsoft specific extension.

## 5 Data Containers

### 5.1 array<T,N>

The type `array<T,N>` represent a dense and regular (not jagged) N-dimensional array which resides on a specific location such as an accelerator or the CPU. The element type of the array is `T`, which is necessarily of a type compatible with the target accelerator. While the rank of the array is determined statically and is part of the type, the extent of the array is runtime-determined, and is expressed using class `extent<N>`. A specific element of an array is selected using an instance of `index<N>`. If "`idx`" is a valid index for an array with extent "`e`", then  $0 \leq \text{idx}[k] < e[k]$  for  $0 \leq k < N$ . Here each "`k`" is referred to as a dimension and higher-numbered dimensions are referred to as less significant.

The array element type `T` shall be an *amp-compatible* whose size is a multiple of 4 bytes and shall not directly or recursively contain any concurrency containers or reference to concurrency containers.

Array data is laid out contiguously in memory. Elements which differ by one in the least significant dimension are adjacent in memory. This storage layout is typically referred to as *row major* and is motivated by achieving efficient memory access given the standard mapping rules that GPUs use for assigning compute domain values to warps.

Arrays are logically considered to be value types in that when an array is copied to another array, a deep copy is performed. Two arrays never point to the same data.

The `array<T,N>` type is used in several distinct scenarios:

- 1562 • As a data container to be used in computations on an accelerator
- 1563 • As a data container to hold memory on the host CPU (to be used to copy to and from other arrays)
- 1564 • As a staging object to act as a fast intermediary in host-to-accelerator copies

1565 An array can have any number of dimensions, although some functionality is specialized for `array<T,1>`, `array<T,2>`, and  
 1566 `array<T,3>`. The dimension defaults to 1 if the template argument is elided.

### 1568 5.1.1 Synopsis

```

1569
1570 template <typename T, int N=1>
1571 class array
1572 {
1573 public:
1574     static const int rank = N;
1575     typedef T value_type;
1576
1577     array() = delete;
1578
1579     explicit array(const extent<N>& extent);
1580     array(const extent<N>& extent, accelerator_view av, accelerator_view associated_av); //
1581     staging
1582
1583     template <typename InputIterator>
1584         array(const extent<N>& extent, InputIterator srcBegin);
1585     template <typename InputIterator>
1586         array(const extent<N>& extent, InputIterator srcBegin, InputIterator srcEnd);
1587     template <typename InputIterator>
1588         array(const extent<N>& extent, InputIterator srcBegin,
1589             accelerator_view av, accelerator_view associated_av); // staging
1590     template <typename InputIterator>
1591         array(const extent<N>& extent, InputIterator srcBegin, InputIterator srcEnd,
1592             accelerator_view av, accelerator_view associated_av); // staging
1593     template <typename InputIterator>
1594         array(const extent<N>& extent, InputIterator srcBegin, accelerator_view av);
1595     template <typename InputIterator>
1596         array(const extent<N>& extent, InputIterator srcBegin, InputIterator srcEnd,
1597             accelerator_view av);
1598
1599     explicit array(const array_view<const T,N>& src);
1600     array(const array_view<const T,N>& src,
1601         accelerator_view av, accelerator_view associated_av); // staging
1602     array(const array_view<const T,N>& src, accelerator_view av);
1603
1604     array(const array& other);
1605     array(array&& other);
1606
1607     array& operator=(const array& other);
1608     array& operator=(array&& other);
1609
1610     array& operator=(const array_view<const T,N>& src);
1611
1612     void copy_to(array& dest) const;
1613     void copy_to(const array_view<T,N>& dest) const;
1614
1615     __declspec(property(get)) extent<N> extent;
```

```

1616
1617 __declspec(property(get)) accelerator_view accelerator_view;
1618 __declspec(property(get)) accelerator_view associated_accelerator_view;
1619
1620 T& operator[](const index<N>& idx) restrict(amp,cpu);
1621 const T& operator[](const index<N>& idx) const restrict(amp,cpu);
1622 array_view<T,N-1> operator[](int i) restrict(amp,cpu);
1623 array_view<const T,N-1> operator[](int i) const restrict(amp,cpu);
1624
1625 const T& operator()(const index<N>& idx) const restrict(amp,cpu);
1626 T& operator()(const index<N>& idx) restrict(amp,cpu);
1627 array_view<T,N-1> operator()(int i) restrict(amp,cpu);
1628 array_view<const T,N-1> operator()(int i) const restrict(amp,cpu);
1629
1630 array_view<T,N> section(const index<N>& idx, const extent<N>& ext) restrict(amp,cpu);
1631 array_view<const T,N> section(const index<N>& idx, const extent<N>& ext) const
1632 restrict(amp,cpu);
1633 array_view<T,N> section(const index<N>& idx) restrict(amp,cpu);
1634 array_view<const T,N> section(const index<N>& idx) const restrict(amp,cpu);
1635
1636 template <typename ElementType>
1637 array_view<ElementType,1> reinterpret_as() restrict(amp,cpu);
1638 template <typename ElementType>
1639 array_view<const ElementType,1> reinterpret_as() const restrict(amp,cpu);
1640
1641 template <int K>
1642 array_view<T,K> view_as(const extent<K>& viewExtent) restrict(amp,cpu);
1643 template <int K>
1644 array_view<const T,K> view_as(const extent<K>& viewExtent) const restrict(amp,cpu);
1645
1646 operator std::vector<T>() const;
1647
1648 T* data() restrict(amp,cpu);
1649 const T* data() const restrict(amp,cpu);
1650 };
1651
1652 template<typename T>
1653 class array<T,1>
1654 {
1655 public:
1656     static const int rank = 1;
1657     typedef T value_type;
1658
1659     array() = delete;
1660
1661     explicit array(const extent<1>& extent);
1662     explicit array(int e0);
1663     array(const extent<1>& extent,
1664           accelerator_view av, accelerator_view associated_av); // staging
1665     array(int e0, accelerator_view av, accelerator_view associated_av); // staging
1666     array(const extent<1>& extent, accelerator_view av);
1667     array(int e0, accelerator_view av);
1668
1669     template <typename InputIterator>
1670     array(const extent<1>& extent, InputIterator srcBegin);
1671     template <typename InputIterator>
1672     array(const extent<1>& extent, InputIterator srcBegin, InputIterator srcEnd);
1673     template <typename InputIterator>

```

```

1674     array(int e0, InputIterator srcBegin);
1675 template <typename InputIterator>
1676     array(int e0, InputIterator srcBegin, InputIterator srcEnd);
1677 template <typename InputIterator>
1678     array(const extent<1>& extent, InputIterator srcBegin,
1679           accelerator_view av, accelerator_view associated_av); // staging
1680 template <typename InputIterator>
1681     array(const extent<1>& extent, InputIterator srcBegin, InputIterator srcEnd,
1682           accelerator_view av, accelerator_view associated_av); // staging
1683 template <typename InputIterator>
1684     array(int e0, InputIterator srcBegin,
1685           accelerator_view av, accelerator_view associated_av); // staging
1686 template <typename InputIterator>
1687     array(int e0, InputIterator srcBegin, InputIterator srcEnd,
1688           accelerator_view av, accelerator_view associated_av); // staging
1689 template <typename InputIterator>
1690     array(const extent<1>& extent, InputIterator srcBegin, accelerator_view av);
1691 template <typename InputIterator>
1692     array(const extent<1>& extent, InputIterator srcBegin, InputIterator srcEnd,
1693           accelerator_view av);
1694 template <typename InputIterator>
1695     array(int e0, InputIterator srcBegin, InputIterator srcEnd, accelerator_view av);
1696
1697 array(const array_view<const T,1>& src);
1698 array(const array_view<const T,1>& src,
1699       accelerator_view av, accelerator_view associated_av); // staging
1700 array(const array_view<const T,1>& src, accelerator_view av);
1701
1702 array(const array& other);
1703 array(array&& other);
1704
1705 array& operator=(const array& other);
1706 array& operator=(array&& other);
1707
1708 array& operator=(const array_view<const T,1>& src);
1709
1710 void copy_to(array& dest) const;
1711 void copy_to(const array_view<T,1>& dest) const;
1712
1713 __declspec(property(get)) extent<1> extent;
1714
1715 __declspec(property(get)) accelerator_view accelerator_view;
1716
1717 T& operator[](const index<1>& idx) restrict(amp,cpu);
1718 const T& operator[](const index<1>& idx) const restrict(amp,cpu);
1719 T& operator[](int i0) restrict(amp,cpu);
1720 const T& operator[](int i0) const restrict(amp,cpu);
1721
1722 T& operator()(const index<1>& idx) restrict(amp,cpu);
1723 const T& operator()(const index<1>& idx) const restrict(amp,cpu);
1724 T& operator()(int i0) restrict(amp,cpu);
1725 const T& operator()(int i0) const restrict(amp,cpu);
1726
1727 array_view<T,1> section(const index<1>& idx, const extent<1>& ext) restrict(amp,cpu);
1728 array_view<const T,1> section(const index<1>& idx, const extent<1>& ext) const
1729 restrict(amp,cpu);
1730 array_view<T,1> section(const index<1>& idx) restrict(amp,cpu);
1731 array_view<const T,1> section(const index<1>& idx) const restrict(amp,cpu);

```

```

1732 array_view<T,1> section(int i0, int e0) restrict(amp,cpu);
1733 array_view<const T,1> section(int i0, int e0) const restrict(amp,cpu);
1734
1735 template <typename ElementType>
1736     array_view<ElementType,1> reinterpret_as() restrict(amp,cpu);
1737 template <typename ElementType>
1738     array_view<const ElementType,1> reinterpret_as() const restrict(amp,cpu);
1739
1740 template <int K>
1741     array_view<T,K> view_as(const extent<K>& viewExtent) restrict(amp,cpu);
1742 template <int K>
1743     array_view<const T,K> view_as(const extent<K>& viewExtent) const restrict(amp,cpu);
1744
1745 operator std::vector<T>() const;
1746
1747 T* data() restrict(amp,cpu);
1748 const T* data() const restrict(amp,cpu);
1749 };
1750
1751
1752 template<typename T>
1753 class array<T,2>
1754 {
1755 public:
1756     static const int rank = 2;
1757     typedef T value_type;
1758
1759     array() = delete;
1760     explicit array(const extent<2>& extent);
1761     array(int e0, int e1);
1762     array(const extent<2>& extent,
1763           accelerator_view av, accelerator_view associated_av); // staging
1764     array(int e0, int e1, accelerator_view av, accelerator_view associated_av); // staging
1765     array(const extent<2>& extent, accelerator_view av);
1766     array(int e0, int e1, accelerator_view av);
1767
1768     template <typename InputIterator>
1769         array(const extent<2>& extent, InputIterator srcBegin);
1770     template <typename InputIterator>
1771         array(const extent<2>& extent, InputIterator srcBegin, InputIterator srcEnd);
1772     template <typename InputIterator>
1773         array(int e0, int e1, InputIterator srcBegin);
1774     template <typename InputIterator>
1775         array(int e0, int e1, InputIterator srcBegin, InputIterator srcEnd);
1776     template <typename InputIterator>
1777         array(const extent<2>& extent, InputIterator srcBegin,
1778               accelerator_view av, accelerator_view associated_av); // staging
1779     template <typename InputIterator>
1780         array(const extent<2>& extent, InputIterator srcBegin, InputIterator srcEnd,
1781               accelerator_view av, accelerator_view associated_av); // staging
1782     template <typename InputIterator>
1783         array(int e0, int e2, InputIterator srcBegin,
1784               accelerator_view av, accelerator_view associated_av); // staging
1785     template <typename InputIterator>
1786         array(int e0, int e2, InputIterator srcBegin, InputIterator srcEnd,
1787               accelerator_view av, accelerator_view associated_av); // staging
1788     template <typename InputIterator>
1789         array(const extent<2>& extent, InputIterator srcBegin, accelerator_view av);

```

```

1790     template <typename InputIterator>
1791         array(const extent<2>& extent, InputIterator srcBegin, InputIterator srcEnd,
1792             accelerator_view av);
1793     template <typename InputIterator>
1794         array(int e0, int e1, InputIterator srcBegin, accelerator_view av);
1795     template <typename InputIterator>
1796         array(int e0, int e1, InputIterator srcBegin, InputIterator srcEnd, accelerator_view av);
1797
1798     array(const array_view<const T,2>& src);
1799     array(const array_view<const T,2>& src,
1800         accelerator_view av, accelerator_view associated_av); // staging
1801     array(const array_view<const T,2>& src, accelerator_view av);
1802
1803     array(const array& other);
1804     array(array&& other);
1805
1806     array& operator=(const array& other);
1807     array& operator=(array&& other);
1808
1809     array& operator=(const array_view<const T,2>& src);
1810
1811     void copy_to(array& dest) const;
1812     void copy_to(const array_view<T,2>& dest) const;
1813
1814     __declspec(property(get)) extent<2> extent;
1815
1816     __declspec(property(get)) accelerator_view accelerator_view;
1817
1818     T& operator[] (const index<2>& idx) restrict(amp,cpu);
1819     const T& operator[] (const index<2>& idx) const restrict(amp,cpu);
1820     array_view<T,1> operator[] (int i0) restrict(amp,cpu);
1821     array_view<const T,1> operator[] (int i0) const restrict(amp,cpu);
1822
1823     T& operator() (const index<2>& idx) restrict(amp,cpu);
1824     const T& operator() (const index<2>& idx) const restrict(amp,cpu);
1825     T& operator() (int i0, int i1) restrict(amp,cpu);
1826     const T& operator() (int i0, int i1) const restrict(amp,cpu);
1827
1828     array_view<T,2> section(const index<2>& idx, const extent<2>& ext) restrict(amp,cpu);
1829     array_view<const T,2> section(const index<2>& idx, const extent<2>& ext) const
1830     restrict(amp,cpu);
1831     array_view<T,2> section(const index<2>& idx) restrict(amp,cpu);
1832     array_view<const T,2> section(const index<2>& idx) const restrict(amp,cpu);
1833     array_view<T,2> section(int i0, int i1, int e0, int e1) restrict(amp,cpu);
1834     array_view<const T,2> section(int i0, int i1, int e0, int e1) const restrict(amp,cpu);
1835
1836     template <typename ElementType>
1837         array_view<ElementType,1> reinterpret_as() restrict(amp,cpu);
1838     template <typename ElementType>
1839         array_view<const ElementType,1> reinterpret_as() const restrict(amp,cpu);
1840
1841     template <int K>
1842         array_view<T,K> view_as(const extent<K>& viewExtent) restrict(amp,cpu);
1843     template <int K>
1844         array_view<const T,K> view_as(const extent<K>& viewExtent) const restrict(amp,cpu);
1845
1846     operator std::vector<T>() const;
1847

```



```

1848     T* data() restrict(amp,cpu);
1849     const T* data() const restrict(amp,cpu);
1850 };
1851
1852
1853 template<typename T>
1854 class array<T,3>
1855 {
1856 public:
1857     static const int rank = 3;
1858     typedef T value_type;
1859
1860     array() = delete;
1861
1862     explicit array(const extent<3>& extent);
1863     array(int e0, int e1, int e2);
1864     array(const extent<3>& extent,
1865           accelerator_view av, accelerator_view associated_av); // staging
1866     array(int e0, int e1, int e2,
1867           accelerator_view av, accelerator_view associated_av); // staging
1868     array(const extent<3>& extent, accelerator_view av);
1869     array(int e0, int e1, int e2, accelerator_view av);
1870
1871     template <typename InputIterator>
1872         array(const extent<3>& extent, InputIterator srcBegin);
1873     template <typename InputIterator>
1874         array(const extent<3>& extent, InputIterator srcBegin, InputIterator srcEnd);
1875     template <typename InputIterator>
1876         array(int e0, int e1, int e2, InputIterator srcBegin);
1877     template <typename InputIterator>
1878         array(int e0, int e1, int e2, InputIterator srcBegin, InputIterator srcEnd);
1879     template <typename InputIterator>
1880         array(const extent<3>& extent, InputIterator srcBegin,
1881               accelerator_view av, accelerator_view associated_av); // staging
1882     template <typename InputIterator>
1883         array(const extent<3>& extent, InputIterator srcBegin, InputIterator srcEnd,
1884               accelerator_view av, accelerator_view associated_av); // staging
1885     template <typename InputIterator>
1886         array(int e0, int e2, int e2, InputIterator srcBegin,
1887               accelerator_view av, accelerator_view associated_av); // staging
1888     template <typename InputIterator>
1889         array(int e0, int e2, int e2, InputIterator srcBegin, InputIterator srcEnd,
1890               accelerator_view av, accelerator_view associated_av); // staging
1891     template <typename InputIterator>
1892         array(const extent<3>& extent, InputIterator srcBegin, accelerator_view av);
1893     template <typename InputIterator>
1894         array(const extent<3>& extent, InputIterator srcBegin, InputIterator srcEnd,
1895               accelerator_view av);
1896     template <typename InputIterator>
1897         array(int e0, int e1, int e2, InputIterator srcBegin, accelerator_view av);
1898     template <typename InputIterator>
1899         array(int e0, int e1, int e2, InputIterator srcBegin, InputIterator srcEnd,
1900               accelerator_view av);
1901
1902     array(const array_view<const T,3>& src);
1903     array(const array_view<const T,3>& src,
1904           accelerator_view av, accelerator_view associated_av); // staging
1905     array(const array_view<const T,3>& src, accelerator_view av);

```



```

1906
1907     array(const array& other);
1908     array(array&& other);
1909
1910     array& operator=(const array& other);
1911     array& operator=(array&& other);
1912
1913     array& operator=(const array_view<const T,3>& src);
1914
1915     void copy_to(array& dest) const;
1916     void copy_to(const array_view<T,3>& dest) const;
1917
1918     __declspec(property(get)) extent<3> extent;
1919
1920     __declspec(property(get)) accelerator_view accelerator_view;
1921
1922     T& operator[](const index<3>& idx) restrict(amp,cpu);
1923     const T& operator[](const index<3>& idx) const restrict(amp,cpu);
1924     array_view<T,2> operator[](int i0) restrict(amp,cpu);
1925     array_view<const T,2> operator[](int i0) const restrict(amp,cpu);
1926
1927     T& operator()(const index<3>& idx) restrict(amp,cpu);
1928     const T& operator()(const index<3>& idx) const restrict(amp,cpu);
1929     T& operator()(int i0, int i1, int i2) restrict(amp,cpu);
1930     const T& operator()(int i0, int i1, int i2) const restrict(amp,cpu);
1931
1932     array_view<T,3> section(const index<3>& idx, const extent<3>& ext) restrict(amp,cpu);
1933     array_view<const T,3> section(const index<3>& idx, const extent<3>& ext) const
1934     restrict(amp,cpu);
1935     array_view<T,3> section(const index<3>& idx) restrict(amp,cpu);
1936     array_view<const T,3> section(const index<3>& idx) const restrict(amp,cpu);
1937     array_view<T,3> section(int i0, int i1, int i2,
1938                             int e0, int e1, int e2) restrict(amp,cpu);
1939     array_view<const T,3> section(int i0, int i1, int i2,
1940                                 int e0, int e1, int e2) const restrict(amp,cpu);
1941
1942     template <typename ElementType>
1943         array_view<ElementType,1> reinterpret_as() restrict(amp,cpu);
1944     template <typename ElementType>
1945         array_view<const ElementType,1> reinterpret_as() const restrict(amp,cpu);
1946
1947     template <int K>
1948         array_view<T,K> view_as(const extent<K>& viewExtent) restrict(amp,cpu);
1949     template <int K>
1950         array_view<const T,K> view_as(const extent<K>& viewExtent) const restrict(amp,cpu);
1951
1952     operator std::vector<T>() const;
1953
1954     T* data() restrict(amp,cpu);
1955     const T* data() const restrict(amp,cpu);
1956 };
1957
1958

```

template <typename T, int N=1> class array	
Represents an N-dimensional region of memory (with type T) located on an accelerator.	
Template Arguments	
<i>T</i>	The element type of this array
<i>N</i>	The dimensionality of the array, defaults to 1 if elided.

1959

```
static const int rank = N
```

The rank of this array.

1960

```
typedef T value_type;
```

The element type of this array.

1961

### 1962 5.1.2 Constructors

1963 There is no default constructor for `array<T,N>`. All constructors are restricted to run on the CPU only (can't be executed on  
1964 an amp target).

1965

```
array(const array& other)
```

Copy constructor. Constructs a new `array<T,N>` from the supplied argument `other`. The new array is located on the same `accelerator_view` as the source array. A deep copy is performed.

**Parameters:**

*Other*

An object of type `array<T,N>` from which to initialize this new array.

1966

```
array(array&& other)
```

Move constructor. Constructs a new `array<T,N>` by moving from the supplied argument `other`.

**Parameters:**

*Other*

An object of type `array<T,N>` from which to initialize this new array.

1967

```
explicit array(const extent<N>& extent)
```

Constructs a new array with the supplied extent, located on the default view of the default accelerator. If any components of the extent are non-positive, an exception will be thrown.

**Parameters:**

*Extent*

The extent in each dimension of this array.

1968

```
explicit array<T,1>::array(int e0)
array<T,2>::array(int e0, int e1)
array<T,3>::array(int e0, int e1, int e2)
```

Equivalent to construction using `"array(extent<N>(e0 [, e1 [, e2 ]]))"`.

**Parameters:**

*e0 [, e1 [, e2 ]]*

The component values that will form the extent of this array.

1969

```
template <typename InputIterator>
array(const extent<N>& extent, InputIterator srcBegin [, InputIterator srcEnd])
```

Constructs a new array with the supplied extent, located on the default accelerator, initialized with the contents of a source container specified by a beginning and optional ending iterator. The source data is copied by value into this array as if by calling `"copy ()"`.

If the number of available container elements is less than `this->extent.size()`, undefined behavior results.

**Parameters:**

*extent*

The extent in each dimension of this array.

*srcBegin*

A beginning iterator into the source container.

*srcEnd*

An ending iterator into the source container.

1970

<pre>template &lt;typename InputIterator&gt;     array&lt;T,1&gt;::array(int e0, InputIterator srcBegin [, InputIterator srcEnd]) template &lt;typename InputIterator&gt;     array&lt;T,2&gt;::array(int e0, int e1, InputIterator srcBegin [, InputIterator srcEnd]) template &lt;typename InputIterator&gt;     array&lt;T,3&gt;::array(int e0, int e1, int e2, InputIterator srcBegin [, InputIterator srcEnd])</pre>	
Equivalent to construction using "array(extent<N>(e0 [, e1 [, e2 ]]), src)".	
<b>Parameters:</b>	
<i>e0</i> [, <i>e1</i> [, <i>e2</i> ]]	The component values that will form the extent of this array.
<i>srcBegin</i>	A beginning iterator into the source container.
<i>srcEnd</i>	An ending iterator into the source container.

1971

<b>explicit</b> array(const array_view<const T,N>& src)	
Constructs a new array, located on the default view of the default accelerator, initialized with the contents of the array_view "src". The extent of this array is taken from the extent of the source array_view. The "src" is copied by value into this array as if by calling "copy(src, *this)" (see 5.3.2).	
<b>Parameters:</b>	
<i>src</i>	An array_view object from which to copy the data into this array (and also to determine the extent of this array).

1972

<b>explicit</b> array(const extent<N>& extent, accelerator_view av)	
Constructs a new array with the supplied extent, located on the accelerator bound to the accelerator_view "av".	
<b>Parameters:</b>	
<i>extent</i>	The extent in each dimension of this array.
<i>av</i>	An accelerator_view object which specifies the location of this array.

1973

<pre>array&lt;T,1&gt;::array(int e0, accelerator_view av) array&lt;T,2&gt;::array(int e0, int e1, accelerator_view av) array&lt;T,3&gt;::array(int e0, int e1, int e2, accelerator_view av)</pre>	
Equivalent to construction using "array(extent<N>(e0 [, e1 [, e2 ]]), av)".	
<b>Parameters:</b>	
<i>e0</i> [, <i>e1</i> [, <i>e2</i> ]]	The component values that will form the extent of this array.
<i>av</i>	An accelerator_view object which specifies the location of this array.

1974

<pre>template &lt;typename InputIterator&gt;     array(const extent&lt;N&gt;&amp; extent, InputIterator srcBegin [, InputIterator srcEnd],           accelerator_view av)</pre>	
Constructs a new array with the supplied extent, located on the accelerator bound to the accelerator_view "av", initialized with the contents of the source container specified by a beginning and optional ending iterator. The data is copied by value into this array as if by calling "copy()".	
<b>Parameters:</b>	
<i>extent</i>	The extent in each dimension of this array.
<i>srcBegin</i>	A beginning iterator into the source container.

<i>srcEnd</i>	An ending iterator into the source container.
<i>av</i>	An <code>accelerator_view</code> object which specifies the location of this array.

1975

<b>array</b> ( <code>const array_view&lt;const T,N&gt;&amp; src, accelerator_view av</code> )	
Constructs a new array initialized with the contents of the <code>array_view</code> "src". The extent of this array is taken from the extent of the source <code>array_view</code> . The "src" is copied by value into this array as if by calling " <code>copy(src, *this)</code> " (see 5.3.2). The new array is located on the accelerator bound to the <code>accelerator_view</code> "av".	
<b>Parameters:</b>	
<i>src</i>	An <code>array_view</code> object from which to copy the data into this array (and also to determine the extent of this array).
<i>av</i>	An <code>accelerator_view</code> object which specifies the location of this array

1976

<pre> template &lt;typename InputIterator&gt;     array&lt;T,1&gt;::array(int e0, InputIterator srcBegin [, InputIterator srcEnd],                     accelerator_view av) template &lt;typename InputIterator&gt;     array&lt;T,2&gt;::array(int e0, int e1, InputIterator srcBegin [, InputIterator srcEnd],                     accelerator_view av) template &lt;typename InputIterator&gt;     array&lt;T,3&gt;::array(int e0, int e1, int e2, InputIterator srcBegin [, InputIterator srcEnd],                     accelerator_view av) </pre>	
Equivalent to construction using " <code>array(extent&lt;N&gt;(e0 [, e1 [, e2 ]]), srcBegin [, srcEnd], av)</code> ".	
<b>Parameters:</b>	
<i>e0 [, e1 [, e2 ]]</i>	The component values that will form the extent of this array.
<i>srcBegin</i>	A beginning iterator into the source container.
<i>srcEnd</i>	An ending iterator into the source container.
<i>av</i>	An <code>accelerator_view</code> object which specifies the location of this array.

1977

1978 

### 5.1.2.1 Staging Array Constructors

1979 Staging arrays are used as a hint to optimize repeated copies between two accelerators (in V1 practically this is between  
1980 the CPU and an accelerator). Staging arrays are optimized for data transfers, and do not have stable user-space memory.

1981 *Microsoft-specific: On Windows, staging arrays are backed by DirectX staging buffers which have the correct hardware  
1982 alignment to ensure efficient DMA transfer between the CPU and a device.*

1983 Staging arrays are differentiated from normal arrays by their construction with a second accelerator. Note that the  
1984 `accelerator_view` property of a staging array returns the value of the first accelerator argument it was constructed with  
1985 (`acclSrc`, below).

1986

1987 It is illegal to change or examine the contents of a staging array while it is involved in a transfer operation (i.e., between  
1988 lines 17 and 22 in the following example):

1989

```

1990 1. class SimulationServer
1991 2. {
1992 3.     array<float,2> acceleratorArray;
1993 4.     array<float,2> stagingArray;
1994 5. public:
1995 6.     SimulationServer(const accelerator_view& av)
1996 7.         :acceleratorArray(extent<2>(1000,1000), av),

```

```

8.         stagingArray(extent<2>(1000,1000), accelerator("cpu").default_view,
9.         accelerator("gpu").default_view)
10.    };
11. };
12.
13. void OnCompute()
14. {
15.     array<float,2> &a = acceleratorArray;
16.     ApplyNetworkChanges(stagingArray.data());
17.     a = stagingArray;
18.     parallel_for_each(a.extents, [&](index<2> idx)
19.     {
20.         // update a[idx] according to simulation
21.     })
22.     stagingArray = a;
23.     SendToClient(stagingArray.data());
24. }
25. };

```

array(const extent<N>& extent, accelerator_view av, accelerator_view associated_av)	
Constructs a staging array with the given extent, which acts as a staging area between accelerators "acclSrc" and "acclDest". If "acclSrc" is a cpu accelerator, this will construct a staging array which is optimized for data transfers between the CPU and "acclDest".	
Parameters:	
<i>extent</i>	The extent in each dimension of this array.
<i>acclSrc</i>	An <b>accelerator</b> object which specifies the home location of this array.
<i>acclDest</i>	An <b>accelerator</b> object which specifies a target device accelerator.

array<T,1>::array(int e0, accelerator_view av, accelerator_view associated_av) array<T,2>::array(int e0, int e1, accelerator_view av, accelerator_view associated_av) array<T,3>::array(int e0, int e1, int e2, accelerator_view av, accelerator_view associated_av)	
Equivalent to construction using "array(extent<N>(e0 [, e1 [, e2 ]]), acclSrc, acclDest)".	
Parameters:	
<i>e0 [, e1 [, e2 ]]</i>	The component values that will form the extent of this array.
<i>acclSrc</i>	An <b>accelerator</b> object which specifies the home location of this array.
<i>acclDest</i>	An <b>accelerator</b> object which specifies a target device accelerator.

template <typename InputIterator> array(const extent<N>& extent, InputIterator srcBegin [, InputIterator srcEnd], accelerator_view av, accelerator_view associated_av)	
Constructs a staging array with the given extent, which acts as a staging area between accelerators "acclSrc" (which must be the CPU accelerator) and "acclDest". The staging array will be initialized with the data specified by "src" as if by calling "copy(src, *this)" (see 5.3.2).	
Parameters:	
<i>extent</i>	The extent in each dimension of this array.
<i>srcBegin</i>	A beginning iterator into the source container.
<i>srcEnd</i>	An ending iterator into the source container.

<i>acclSrc</i>	An <b>accelerator</b> object which specifies the home location of this array.
<i>acclDest</i>	An <b>accelerator</b> object which specifies a target device accelerator.

2019  
2020

<b>array</b> ( <b>const</b> array_view< <b>const</b> T,N>& src, accelerator_view av, accelerator_view associated_av)	
Constructs a staging array initialized with the array_view given by "src", which acts as a staging area between accelerators "acclSrc" (which must be the CPU accelerator) and "acclDest". The extent of this array is taken from the extent of the source array_view. The staging array will be initialized from "src" as if by calling "copy(src, *this)" (see 5.3.2).	
<b>Parameters:</b>	
<i>src</i>	An <b>array_view</b> object from which to copy the data into this array (and also to determine the extent of this array).
<i>acclSrc</i>	An <b>accelerator</b> object which specifies the home location of this array.
<i>acclDest</i>	An <b>accelerator</b> object which specifies a target device accelerator.

2021

<pre>template &lt;typename InputIterator&gt;     array&lt;T,1&gt;::array(int e0, InputIterator srcBegin [, InputIterator srcEnd], accelerator_view     av, accelerator_view associated_av) template &lt;typename InputIterator&gt;     array&lt;T,2&gt;::array(int e0, int e1, InputIterator srcBegin [, InputIterator srcEnd],                       accelerator_view av, accelerator_view associated_av) template &lt;typename InputIterator&gt;     array&lt;T,3&gt;::array(int e0, int e1, int e2, InputIterator srcBegin [, InputIterator srcEnd],                       accelerator_view av, accelerator_view associated_av)</pre>	
Equivalent to construction using " <b>array</b> ( <b>extent</b> <N>(e0 [, e1 [, e2 ]]), src, acclSrc, acclDest)".	
<b>Parameters:</b>	
<i>e0 [, e1 [, e2 ]]</i>	The component values that will form the extent of this array.
<i>srcBegin</i>	A beginning iterator into the source container.
<i>srcEnd</i>	An ending iterator into the source container.
<i>acclSrc</i>	An <b>accelerator</b> object which specifies the home location of this array.
<i>acclDest</i>	An <b>accelerator</b> object which specifies a target device accelerator.

2022  
2023

### 2024 5.1.3 Members

2025

<b>__declspec(property(get))</b> extent<N> extent
extent<N> get_extent() <b>const restrict</b> (cpu,amp)
Access the extent that defines the shape of this array.

2026

<b>__declspec(property(get))</b> accelerator_view accelerator_view
This property returns the accelerator_view representing the location where this array has been allocated. This property is only accessible on the CPU.

2027

<b>array</b> & <b>operator</b> =( <b>const</b> array& other)
--

Assigns the contents of the array "other" to this array, using a deep copy. This function can only be called on the CPU.

**Parameters:**

*other* An object of type `array<T,N>` from which to copy into this array.

**Return Value:**

Returns `*this`.

2028

`array& operator=(array&& other)`

Moves the contents of the array "other" to this array. This function can only be called on the CPU.

**Parameters:**

*other* An object of type `array<T,N>` from which to move into this array.

**Return Value:**

Returns `*this`.

2029

`array& operator=(const array_view<const T,N>& src)`

Assigns the contents of the array\_view "src", as if by calling "copy(src, \*this)" (see 5.3.2).

**Parameters:**

*src* An object of type `array_view<T,N>` from which to copy into this array.

**Return Value:**

Returns `*this`.

2030

`void copy_to(array<T,N>& dest)`

Copies the contents of this array to the array given by "dest", as if by calling "copy(\*this, dest)" (see 5.3.2).

**Parameters:**

*dest* An object of type `array <T,N>` to which to copy data from this array.

2031

`void copy_to(const array_view<T,N>& dest)`

Copies the contents of this array to the array\_view given by "dest", as if by calling "copy(\*this, dest)" (see 5.3.2).

**Parameters:**

*dest* An object of type `array_view<T,N>` to which to copy data from this array.

2032

`T* data() restrict(amp,cpu)`

`const T* data() const restrict(amp,cpu)`

Returns a pointer to the raw data underlying this array.

**Return Value:**

A (const) pointer to the first element in the linearized array.

2033

`operator std::vector<T>() const`

Implicitly converts an array to a `std::vector`, as if by "copy(\*this, vector)" (see 5.3.2).

**Return Value:**

An object of type `vector<T>` which contains a copy of the data contained on the array.

2034

## 5.1.4 Indexing

2035

2036

`T& operator[](const index<N>& idx) restrict(amp,cpu)`

`T& operator()(const index<N>& idx) restrict(amp,cpu)`

Returns a reference to the element of this array that is at the location in N-dimensional space specified by "idx".

**Parameters:**

*idx* An object of type `index<N>` from that specifies the location of the element.

2037

`const T& operator[](const index<N>& idx) const restrict(amp,cpu)`

`const T& operator()(const index<N>& idx) const restrict(amp,cpu)`

Returns a const reference to the element of this array that is at the location in N-dimensional space specified by "idx".

**Parameters:**

<i>idx</i>	An object of type <code>index&lt;N&gt;</code> from that specifies the location of the element.
------------	--

2038

<pre>T&amp; array&lt;T,1&gt;::operator()(int i0) restrict(amp,cpu) T&amp; array&lt;T,2&gt;::operator()(int i0, int i1) restrict(amp,cpu) T&amp; array&lt;T,3&gt;::operator()(int i0, int i1, int i2) restrict(amp,cpu)</pre>	
Equivalent to " <code>array&lt;T,N&gt;::operator()(index&lt;N&gt;(i0 [, i1 [, i2 ]])</code> ".	
<b>Parameters:</b>	
<i>i0</i> [, <i>i1</i> [, <i>i2</i> ] ]	The component values that will form the index into this array.

2039

<pre>const T&amp; array&lt;T,1&gt;::operator()(int i0) const restrict(amp,cpu) const T&amp; array&lt;T,2&gt;::operator()(int i0, int i1) const restrict(amp,cpu) const T&amp; array&lt;T,3&gt;::operator()(int i0, int i1, int i2) const restrict(amp,cpu)</pre>	
Equivalent to " <code>array&lt;T,N&gt;::operator()(index&lt;N&gt;(i0 [, i1 [, i2 ]]) const</code> ".	
<b>Parameters:</b>	
<i>i0</i> [, <i>i1</i> [, <i>i2</i> ] ]	The component values that will form the index into this array.

2040

<pre>array_view&lt;T,N-1&gt; operator[](int i0) restrict(amp,cpu) array_view&lt;const T,N-1&gt; operator[](int i0) const restrict(amp,cpu)</pre>	
This overload is defined for <code>array&lt;T,N&gt;</code> where $N \geq 2$ .	
This mode of indexing is equivalent to projecting on the most-significant dimension. It allows C-style indexing. For example:	
<pre>array&lt;float,4&gt; myArray(myExtents, ...);  myArray[index&lt;4&gt;(5,4,3,2)] = 7; assert(myArray[5][4][3][2] == 7);</pre>	
<b>Parameters:</b>	
<i>i0</i>	An integer that is the index into the most-significant dimension of this array.
<b>Return Value:</b>	
Returns an <code>array_view</code> whose dimension is one lower than that of this array.	

2041

2042 

## 5.1.5 View Operations

2043

<pre>array_view&lt;T,N&gt; section(const index&lt;N&gt;&amp; offset, const extent&lt;N&gt;&amp; ext) restrict(amp,cpu) array_view&lt;const T,N&gt; section(const index&lt;N&gt;&amp; offset, const extent&lt;N&gt;&amp; ext) const restrict(amp,cpu)</pre>	
See " <code>array_view&lt;T,N&gt;::section(const index&lt;N&gt;&amp;, const extent&lt;N&gt;&amp;)</code> " in section 5.2.2 for a description of this function.	

2044

<pre>array_view&lt;T,N&gt; section(const index&lt;N&gt;&amp; idx) restrict(amp,cpu) array_view&lt;const T,N&gt; section(const index&lt;N&gt;&amp; idx) const restrict(amp,cpu)</pre>	
Equivalent to " <code>section(idx, this-&gt;extent - idx)</code> ".	

2045

<pre>array_view&lt;T,1&gt; array&lt;T,1&gt;::section(int i0, int e0) restrict(amp,cpu) array_view&lt;const T,1&gt; array&lt;T,1&gt;::section(int i0, int e0) const restrict(amp,cpu) array_view&lt;T,2&gt; array&lt;T,2&gt;::section(int i0, int i1, int e0, int e1) restrict(amp,cpu) array_view&lt;const T,2&gt; array&lt;T,2&gt;::section(int i0, int i1,  int e0, int e1) const restrict(amp,cpu) array_view&lt;T,3&gt; array&lt;T,3&gt;::section(int i0, int i1, int i2,                                      int e0, int e1, int e2) restrict(amp,cpu) array_view&lt;const T,3&gt; array&lt;T,3&gt;::section(int i0, int i1, int i2,</pre>	
--	--



<code>int e0, int e1, int e2) const restrict(amp,cpu)</code>	
Equivalent to " <code>array&lt;T,N&gt;::section(index&lt;N&gt;(i0 [, i1 [, i2 ]]), extent&lt;N&gt;(e0 [, e1 [, e2 ]])) const</code> ".	
<b>Parameters:</b>	
<code>i0 [, i1 [, i2 ]]</code>	The component values that will form the origin of the section
<code>e0 [, e1 [, e2 ]]</code>	The component values that will form the extent of the section

2046

<pre>template&lt;typename ElementType&gt;     array_view&lt;ElementType,1&gt; reinterpret_as() restrict(amp,cpu) template&lt;typename ElementType&gt;     array_view&lt;const ElementType,1&gt; reinterpret_as() const restrict(amp,cpu)</pre>	
<p>Sometimes it is desirable to view the data of an N-dimensional array as a linear array, possibly with a (unsafe) reinterpretation of the element type. This can be achieved through the <code>reinterpret_as</code> member function. Example:</p> <pre>struct RGB { float r; float g; float b; };  array&lt;RGB,3&gt; a = ...; array_view&lt;float,1&gt; v = a.reinterpret_as&lt;float&gt;();  assert(v.extent == 3*a.extent);</pre> <p>The size of the reinterpreted <code>ElementType</code> must evenly divide into the total size of this array.</p>	
<b>Return Value:</b>	
Returns an <code>array_view</code> from this <code>array&lt;T,N&gt;</code> with the element type reinterpreted from <code>T</code> to <code>ElementType</code> , and the rank reduced from <code>N</code> to 1.	

2047

<pre>template &lt;int K&gt;     array_view&lt;T,K&gt; view_as(extent&lt;K&gt; viewExtent) restrict(amp,cpu) template &lt;int K&gt;     array_view&lt;const T,K&gt; view_as(extent&lt;K&gt; viewExtent) const restrict(amp,cpu)</pre>	
<p>An array of higher rank can be reshaped into an array of lower rank, or vice versa, using the <code>view_as</code> member function. Example:</p> <pre>array&lt;float,1&gt; a(100);  array_view&lt;float,2&gt; av = a.view_as(extent&lt;2&gt;(2,50));</pre>	
<b>Return Value:</b>	
Returns an <code>array_view</code> from this <code>array&lt;T,N&gt;</code> with the rank changed to <code>K</code> from <code>N</code> .	

2048

## 5.2 array\_view<T,N>

2049

2050

2051

2052

2053

2054

2055

2056

2057

2058

2059

2060

2061

2062

2063

The `array_view<T,N>` type represents a possibly cached view into the data held in an `array<T,N>`, or a section thereof. It also provides such views over native CPU data. It exposes an indexing interface congruent to that of `array<T,N>`.

Like an `array`, an `array_view` is an N-dimensional object, where N defaults to 1 if it is omitted.

The array element type `T` shall be an *amp-compatible* whose size is a multiple of 4 bytes and shall not directly or recursively contain any concurrency containers or reference to concurrency containers.

`array_views` may be accessed locally, where their source data lives, or remotely on a different accelerator\_view or coherence domain. When they are accessed remotely, views are copied and cached as necessary. Except for the effects of automatic caching, `array_views` have a performance profile similar to that of arrays (small to negligible access penalty when accessing the data through views).

2064 There are three remote usage scenarios:

- 2065 1. A view to a system memory pointer is passed through a *parallel\_for\_each* call to an accelerator and accessed on  
2066 the accelerator.
- 2067 2. A view to an accelerator-residing array is passed using a *parallel\_for\_each* to another accelerator\_view and is  
2068 accessed there.
- 2069 3. A view to an accelerator-residing array is accessed on the CPU.

2070 When any of these scenarios occur, the referenced views are implicitly copied by the system to the remote location and, if  
2071 modified through the *array\_view*, copied back to the home location. The Implementation is free to optimize copying  
2072 changes back; may only copy changed elements, or may copy unchanged portions as well. Overlapping *array\_views* to the  
2073 same data source are *not guaranteed to maintain aliasing between arrays/array\_views* on a remote location.

2074

2075 Multi-threaded access to the same data source, either directly or through views, must be synchronized by the user.

2076

2077 The runtime makes the following guarantees regarding caching of data inside array views.

- 2078 1. Let A be an array and V a view to the array. Then, all well-synchronized accesses to A and V in program order obey  
2079 a serial happens-before relationship.
- 2080 2. Let A be an array and V1 and V2 be overlapping views to the array.
  - 2081 • When executing on the accelerator where A has been allocated, all well-synchronized accesses through A,  
2082 V1 and V2 are aliased through A and induce a total happens-before relationship which obeys program  
2083 order. (No caching.)
  - 2084 • Otherwise, if they are executing on different accelerators, then the behaviour of writes to V1 and V2 is  
2085 undefined (a race).

2086 When an *array\_view* is created over a pointer in system memory, the user commits to:

- 2087 1. only changing the data accessible through the view directly through the view class, **or**
- 2088 2. adhering to the following rules when accessing the data directly (not through the view):
  - 2089 a. Calling *synchronize()* before the data is accessed directly, **and**
  - 2090 b. If the underlying data is modified, calling *refresh()* prior to further accessing it through the view.

2091 Either action will notify the *array\_view* that the underlying native memory has changed and that any accelerator-residing  
2092 copies are now stale. If the user abides by these rules then the guarantees provided by the system for pointer-based views  
2093 are identical to those provided to views of data-parallel arrays.

## 2094 5.2.1 Synopsis

2095 The *array\_view<T,N>* has the following specializations:

- 2096 • *array\_view<T,1>*
- 2097 • *array\_view<T,2>*
- 2098 • *array\_view<T,3>*
- 2099 • *array\_view<const T,N>*
- 2100 • *array\_view<const T,1>*
- 2101 • *array\_view<const T,2>*
- 2102 • *array\_view<const T,3>*

### 2103 5.2.1.1 *array\_view<T,N>*

2104 The generic *array\_view<T,N>* represents a view over elements of type *T* with rank *N*. The elements are both readable and  
2105 writeable.

2106

2107 `template <typename T, int N = 1>`

```

2108 class array_view
2109 {
2110 public:
2111     static const int rank = N;
2112     typedef T value_type;
2113
2114     array_view() = delete;
2115     array_view(array<T,N>& src) restrict(amp,cpu);
2116     template <typename Container>
2117         array_view(const extent<N>& extent, Container& src);
2118     array_view(const extent<N>& extent, value_type* src) restrict(amp,cpu);
2119
2120     array_view(const array_view& other) restrict(amp,cpu);
2121
2122     array_view& operator=(const array_view& other) restrict(amp,cpu);
2123
2124     void copy_to(array<T,N>& dest) const;
2125     void copy_to(const array_view& dest) const;
2126
2127     __declspec(property(get)) extent<N> extent;
2128
2129     // These are restrict(amp,cpu)
2130     T& operator[] (const index<N>& idx) const restrict(amp,cpu);
2131     array_view<T,N-1> operator[] (int i) const restrict(amp,cpu);
2132
2133     T& operator() (const index<N>& idx) const restrict(amp,cpu);
2134     array_view<T,N-1> operator() (int i) const restrict(amp,cpu);
2135
2136     array_view<T,N> section(const index<N>& idx, const extent<N>& ext) restrict(amp,cpu);
2137     array_view<T,N> section(const index<N>& idx) const restrict(amp,cpu);
2138
2139     void synchronize() const;
2140     completion_future synchronize_async() const;
2141
2142     void refresh() const;
2143     void discard_data() const;
2144
2145 };
2146
2147 template <typename T>
2148 class array_view<T,1>
2149 {
2150 public:
2151     static const int rank = 1;
2152     typedef T value_type;
2153
2154     array_view() = delete;
2155     array_view(array<T,1>& src) restrict(amp,cpu);
2156     template <typename Container>
2157         array_view(const extent<1>& extent, Container& src);
2158     template <typename Container>
2159         array_view(int e0, Container& src);
2160     array_view(const extent<1>& extent, value_type* src) restrict(amp,cpu);
2161     array_view(int e0, value_type* src) restrict(amp,cpu);
2162
2163     array_view(const array_view& other) restrict(amp,cpu);
2164
2165     array_view& operator=(const array_view& other) restrict(amp,cpu);

```

```

2166
2167 void copy_to(array<T,1>& dest) const;
2168 void copy_to(const array_view& dest) const;
2169
2170 __declspec(property(get)) extent<1> extent;
2171
2172 T& operator[](const index<1>& idx) const restrict(amp,cpu);
2173 T& operator[](int i) const restrict(amp,cpu);
2174
2175 T& operator()(const index<1>& idx) const restrict(amp,cpu);
2176 T& operator()(int i) const restrict(amp,cpu);
2177
2178 array_view<T,1> section(const index<1>& idx, const extent<1>& ext) const restrict(amp,cpu);
2179 array_view<T,1> section(const index<1>& idx) const restrict(amp,cpu);
2180 array_view<T,1> section(const extent<1>& ext) const restrict(amp,cpu);
2181 array_view<T,1> section(int i0, int e0) restrict(amp,cpu);
2182
2183 template <typename ElementType>
2184     array_view<ElementType,1> reinterpret_as() const restrict(amp,cpu);
2185
2186 template <int K>
2187     array_view<T,K> view_as(extent<K> viewExtent) const restrict(amp,cpu);
2188
2189 T* data() const restrict(amp,cpu);
2190
2191 void synchronize() const;
2192 completion_future synchronize_async() const;
2193
2194 void refresh() const;
2195 void discard_data() const;
2196 };
2197
2198
2199 template <typename T>
2200 class array_view<T,2>
2201 {
2202 public:
2203     static const int rank = 2;
2204     typedef T value_type;
2205
2206     array_view() = delete;
2207     array_view(array<T,2>& src) restrict(amp,cpu);
2208     template <typename Container>
2209         array_view(const extent<2>& extent, Container& src);
2210     template <typename Container>
2211         array_view(int e0, int e1, Container& src);
2212     array_view(const extent<2>& extent, value_type* src) restrict(amp,cpu);
2213     array_view(int e0, int e1, value_type* src) restrict(amp,cpu);
2214
2215     array_view(const array_view& other) restrict(amp,cpu);
2216
2217     array_view& operator=(const array_view& other) restrict(amp,cpu);
2218
2219     void copy_to(array<T,2>& dest) const;
2220     void copy_to(const array_view& dest) const;
2221
2222     __declspec(property(get)) extent<2> extent;
2223

```

```

2224 T& operator[](const index<2>& idx) const restrict(amp,cpu);
2225 array_view<T,1> operator[](int i) const restrict(amp,cpu);
2226
2227 T& operator()(const index<2>& idx) const restrict(amp,cpu);
2228 T& operator()(int i0, int i1) const restrict(amp,cpu);
2229
2230 array_view<T,2> section(const index<2>& idx, const extent<2>& ext) const restrict(amp,cpu);
2231 array_view<T,2> section(const index<2>& idx) const restrict(amp,cpu);
2232 array_view<T,2> section(const extent<2>& ext) const restrict(amp,cpu);
2233 array_view<T,2> section(int i0, int i1, int e0, int e1) const restrict(amp,cpu);
2234
2235 void synchronize() const;
2236 completion_future synchronize_async() const;
2237
2238 void refresh() const;
2239 void discard_data() const;
2240 };
2241
2242 template <typename T>
2243 class array_view<T,3>
2244 {
2245 public:
2246     static const int rank = 3;
2247     typedef T value_type;
2248
2249     array_view() = delete;
2250     array_view(array<T,3>& src) restrict(amp,cpu);
2251     template <typename Container>
2252         array_view(const extent<3>& extent, Container& src);
2253     template <typename Container>
2254         array_view(int e0, int e1, int e2, Container& src);
2255     array_view(const extent<3>& extent, value_type* src) restrict(amp,cpu);
2256     array_view(int e0, int e1, int e2, value_type* src) restrict(amp,cpu);
2257
2258     array_view(const array_view& other) restrict(amp,cpu);
2259
2260     array_view& operator=(const array_view& other) restrict(amp,cpu);
2261
2262     void copy_to(array<T,3>& dest) const;
2263     void copy_to(const array_view& dest) const;
2264
2265     __declspec(property(get)) extent<3> extent;
2266
2267     T& operator[](const index<3>& idx) const restrict(amp,cpu);
2268     array_view<T,2> operator[](int i) const restrict(amp,cpu);
2269
2270     T& operator()(const index<3>& idx) const restrict(amp,cpu);
2271     T& operator()(int i0, int i1, int i2) const restrict(amp,cpu);
2272
2273     array_view<T,3> section(const index<3>& idx, const extent<3>& ext) const restrict(amp,cpu);
2274     array_view<T,3> section(const index<3>& idx) const restrict(amp,cpu);
2275     array_view<T,3> section(const extent<3>& ext) const restrict(amp,cpu);
2276     array_view<T,3> section(int i0, int i1, int i2, int e0, int e1, int e2) const
2277     restrict(amp,cpu);
2278
2279     void synchronize() const;
2280     completion_future synchronize_async() const;
2281

```

```

2282     void refresh() const;
2283     void discard_data() const;
2284 };
2285

```

#### 2286 5.2.1.2 array\_view<const T,N>

2287 The partial specialization `array_view<const T,N>` represents a view over elements of type `const T` with rank `N`. The  
 2288 elements are readonly. At the boundary of a call site (such as `parallel_for_each`), this form of `array_view` need only be  
 2289 copied to the target accelerator if it isn't already there. It will not be copied out.

```

2290
2291 template <typename T, int N=1>
2292 class array_view<const T,N>
2293 {
2294 public:
2295     static const int rank = N;
2296     typedef const T value_type;
2297
2298     array_view() = delete;
2299     array_view(const array<T,N>& src) restrict(amp,cpu);
2300     template <typename Container>
2301         array_view(const extent<N>& extent, const Container& src);
2302     array_view(const extent<N>& extent, const value_type* src) restrict(amp,cpu);
2303
2304     array_view(const array_view<T,N>& other) restrict(amp,cpu);
2305     array_view(const array_view<const T,N>& other) restrict(amp,cpu);
2306
2307     array_view& operator=(const array_view& other) restrict(amp,cpu);
2308
2309     void copy_to(array<T,N>& dest) const;
2310     void copy_to(const array_view<T,N>& dest) const;
2311
2312     __declspec(property(get)) extent<N> extent;
2313
2314     const T& operator[](const index<N>& idx) const restrict(amp,cpu);
2315     array_view<const T,N-1> operator[](int i) const restrict(amp,cpu);
2316
2317     const T& operator()(const index<N>& idx) const restrict(amp,cpu);
2318     array_view<const T,N-1> operator()(int i) const restrict(amp,cpu);
2319
2320     array_view<const T,N> section(const index<N>& idx, const extent<N>& ext) const
2321     restrict(amp,cpu);
2322     array_view<const T,N> section(const index<N>& idx) const restrict(amp,cpu);
2323
2324     void refresh() const;
2325 };
2326
2327 template <typename T>
2328 class array_view<const T,1>
2329 {
2330 public:
2331     static const int rank = 1;
2332     typedef const T value_type;
2333
2334     array_view() = delete;
2335     array_view(const array<T,1>& src) restrict(amp,cpu);
2336     template <typename Container>
2337         array_view(const extent<1>& extent, const Container& src);

```

```

2338     template <typename Container>
2339         array_view(int e0, const Container& src);
2340     array_view(const extent<1>& extent, const value_type* src) restrict(amp,cpu);
2341     array_view(int e0, const value_type* src) restrict(amp,cpu);
2342
2343     array_view(const array_view<T,1>& other) restrict(amp,cpu);
2344     array_view(const array_view<const T,1>& other) restrict(amp,cpu);
2345
2346     array_view& operator=(const array_view& other) restrict(amp,cpu);
2347
2348     void copy_to(array<T,1>& dest) const;
2349     void copy_to(const array_view<T,1>& dest) const;
2350
2351     __declspec(property(get)) extent<1> extent;
2352
2353     // These are restrict(amp,cpu)
2354     const T& operator[] (const index<1>& idx) const restrict(amp,cpu);
2355     const T& operator[] (int i) const restrict(amp,cpu);
2356
2357     const T& operator() (const index<1>& idx) const restrict(amp,cpu);
2358     const T& operator() (int i) const restrict(amp,cpu);
2359
2360     array_view<const T,1> section(const index<N>& idx, const extent<N>& ext) const
2361     restrict(amp,cpu);
2362     array_view<const T,1> section(const index<1>& idx) const restrict(amp,cpu);
2363     array_view<const T,1> section(const extent<1>& ext) const restrict(amp,cpu);
2364     array_view<const T,1> section(int i0, int e0) const restrict(amp,cpu);
2365
2366     template <typename ElementType>
2367         array_view<const ElementType,1> reinterpret_as() const restrict(amp,cpu);
2368
2369     template <int K>
2370         array_view<const T,K> view_as(extent<K> viewExtent) const restrict(amp,cpu);
2371
2372     const T* data() const restrict(amp,cpu);
2373
2374     void refresh() const;
2375 };
2376
2377 template <typename T>
2378 class array_view<const T,2>
2379 {
2380 public:
2381     static const int rank = 2;
2382     typedef const T value_type;
2383
2384     array_view() = delete;
2385     array_view(const array<T,2>& src) restrict(amp,cpu);
2386     template <typename Container>
2387         array_view(const extent<2>& extent, const Container& src);
2388     template <typename Container>
2389         array_view(int e0, int e1, const Container& src);
2390     array_view(const extent<2>& extent, const value_type* src) restrict(amp,cpu);
2391     array_view(int e0, int e1, const value_type* src) restrict(amp,cpu);
2392
2393     array_view(const array_view<T,2>& other) restrict(amp,cpu);
2394     array_view(const array_view<const T,2>& other) restrict(amp,cpu);
2395

```



```

2396     array_view& operator=(const array_view& other) restrict(amp,cpu);
2397
2398     void copy_to(array<T,2>& dest) const;
2399     void copy_to(const array_view<T,2>& dest) const;
2400
2401     __declspec(property(get)) extent<2> extent;
2402
2403     const T& operator[](const index<2>& idx) const restrict(amp,cpu);
2404     array_view<const T,1> operator[](int i) const restrict(amp,cpu);
2405
2406     const T& operator()(const index<2>& idx) const restrict(amp,cpu);
2407     const T& operator()(int i0, int i1) const restrict(amp,cpu);
2408
2409     array_view<const T,2> section(const index<2>& idx, const extent<2>& ext) const
2410 restrict(amp,cpu);
2411     array_view<const T,2> section(const index<2>& idx) const restrict(amp,cpu);
2412     array_view<const T,2> section(const extent<2>& ext) const restrict(amp,cpu);
2413     array_view<const T,2> section(int i0, int i1, int e0, int e1) const restrict(amp,cpu);
2414
2415     void refresh() const;
2416 };
2417
2418 template <typename T>
2419 class array_view<const T,3>
2420 {
2421 public:
2422     static const int rank = 3;
2423     typedef const T value_type;
2424
2425     array_view() = delete;
2426     array_view(const array<T,3>& src) restrict(amp,cpu);
2427     template <typename Container>
2428     array_view(const extent<3>& extent, const Container& src);
2429     template <typename Container>
2430     array_view(int e0, int e1, int e2, const Container& src);
2431     array_view(const extent<3>& extent, const value_type* src) restrict(amp,cpu);
2432     array_view(int e0, int e1, int e2, const value_type* src) restrict(amp,cpu);
2433
2434     array_view(const array_view<T,3>& other) restrict(amp,cpu);
2435     array_view(const array_view<const T,3>& other) restrict(amp,cpu);
2436
2437     array_view& operator=(const array_view& other) restrict(amp,cpu);
2438
2439     void copy_to(array<T,3>& dest) const;
2440     void copy_to(const array_view<T,3>& dest) const;
2441
2442     __declspec(property(get)) extent<3> extent;
2443
2444     // These are restrict(amp,cpu)
2445     const T& operator[](const index<3>& idx) const restrict(amp,cpu);
2446     array_view<const T,2> operator[](int i) const restrict(amp,cpu);
2447
2448     const T& operator()(const index<3>& idx) const restrict(amp,cpu);
2449     const T& operator()(int i0, int i1, int i2) const restrict(amp,cpu);
2450
2451     array_view<const T,3> section(const index<3>& idx, const extent<3>& ext) const
2452 restrict(amp,cpu);
2453     array_view<const T,3> section(const index<3>& idx) const restrict(amp,cpu);

```



```

2454     array_view<const T,3> section(const extent<3>& ext) const restrict(amp,cpu);
2455     array_view<const T,3> section(int i0, int i1, int i2, int e0, int e1, int e2) const
2456 restrict(amp,cpu);
2457
2458     void refresh() const;
2459 };

```

## 2460 5.2.2 Constructors

2461 The `array_view` type cannot be default-constructed. It must be bound at construction time to a memory location.

2463 No bounds-checking is performed when constructing `array_views`.

2465  
2466

```

array_view<T,N>::array_view(array<T,N>& src) restrict(amp,cpu)
array_view<const T,N>::array_view(const array<T,N>& src) restrict(amp,cpu)

```

Constructs an `array_view` which is bound to the data contained in the "src" array. The extent of the `array_view` is that of the src array, and the origin of the array view is at zero.

### Parameters:

<i>Src</i>	An array which contains the data that this <code>array_view</code> is bound to.
------------	---

2467

```

template <typename Container>
array_view<T,N>::array_view(const extent<N>& extent, Container& src)
template <typename Container>
array_view<const T,N>::array_view(const extent<N>& extent, const Container& src)

```

Constructs an `array_view` which is bound to the data contained in the "src" container. The extent of the `array_view` is that given by the "extent" argument, and the origin of the array view is at zero.

### Parameters:

<i>Src</i>	A template argument that must resolve to a linear container that supports <code>.data()</code> and <code>.size()</code> members (such as <code>std::vector</code> or <code>std::array</code> )
<i>Extent</i>	The extent of this <code>array_view</code> .

2468

```

array_view<T,N>::array_view(const extent<N>& extent, value_type* src) restrict(amp,cpu)
array_view<const T,N>::array_view(const extent<N>& extent,
                                const value_type* src) restrict(amp,cpu)

```

Constructs an `array_view` which is bound to the data contained in the "src" container. The extent of the `array_view` is that given by the "extent" argument, and the origin of the array view is at zero.

### Parameters:

<i>Src</i>	A pointer to the source data that will be copied into this array.
<i>Extent</i>	The extent of this <code>array_view</code> .

2469

```

template <typename Container>
array_view<T,1>::array_view(int e0, Container& src)
template <typename Container>
array_view<T,2>::array_view(int e0, int e1, Container& src)
template <typename Container>
array_view<T,3>::array_view(int e0, int e1, int e2, Container& src)

template <typename Container>
array_view<const T,1>::array_view(int e0, const Container& src)
template <typename Container>

```

<pre>array_view&lt;const T,2&gt;::array_view(int e0, int e1, const Container&amp; src) template &lt;typename Container&gt; array_view&lt;const T,3&gt;::array_view(int e0, int e1, int e2, const Container&amp; src)</pre>	
Equivalent to construction using "array_view(extent<N>(e0 [, e1 [, e2 ]]), src)".	
<b>Parameters:</b>	
<i>e0</i> [, <i>e1</i> [, <i>e2</i> ]]	The component values that will form the extent of this array_view.
<i>Src</i>	A template argument that must resolve to a contiguous container that supports .data() and .size() members (such as std::vector or std::array)

2470

<pre>array_view&lt;T,1&gt;::array_view(int e0, value_type* src) restrict(amp,cpu) array_view&lt;T,2&gt;::array_view(int e0, int e1, value_type* src) restrict(amp,cpu) array_view&lt;T,3&gt;::array_view(int e0, int e1, int e2, value_type* src) restrict(amp,cpu)  array_view&lt;const T,1&gt;::array_view(int e0, const value_type* src) restrict(amp,cpu) array_view&lt;const T,2&gt;::array_view(int e0, int e1, const value_type* src) restrict(amp,cpu) array_view&lt;const T,3&gt;::array_view(int e0, int e1, int e2,                                 const value_type* src) restrict(amp,cpu)</pre>	
Equivalent to construction using "array_view(extent<N>(e0 [, e1 [, e2 ]]), src)".	
<b>Parameters:</b>	
<i>e0</i> [, <i>e1</i> [, <i>e2</i> ]]	The component values that will form the extent of this array_view.
<i>Src</i>	A pointer to the source data that will be copied into this array.

2471

<pre>array_view(const array_view&lt;T,N&gt;&amp; other) restrict(amp,cpu) array_view(const array_view&lt;const T,N&gt;&amp; other) restrict(amp,cpu);</pre>	
Copy constructor. Constructs a new array_view<T,N> from the supplied argument other. A shallow copy is performed.	
<b>Parameters:</b>	
<i>Other</i>	An object of type array_view<T,N> or array_view<const T,N> from which to initialize this new array_view.

2472

### 2473 5.2.3 Members

2474

<pre>__declspec(property(get)) extent&lt;N&gt; extent extent&lt;N&gt; get_extent() const restrict(cpu,amp)</pre>	
Access the extent that defines the shape of this array_view.	

2475

<pre>array_view&amp; operator=(const array_view&amp; other) restrict(amp,cpu)</pre>	
Assigns the contents of the array_view "other" to this array_view, using a shallow copy. Both array_views will refer to the same data.	
<b>Parameters:</b>	
<i>other</i>	An object of type array_view<T,N> from which to copy into this array.
<b>Return Value:</b>	
Returns *this.	

2476

<pre>void copy_to(array&lt;T,N&gt;&amp; dest)</pre>	
Copies the data referred to by this array_view to the array given by "dest", as if by calling "copy(*this, dest)" (see 5.3.2).	
<b>Parameters:</b>	
<i>dest</i>	An object of type array <T,N> to which to copy data from this array.

2477

<pre>void copy_to(const array_view&amp; dest)</pre>	
Copies the contents of this array_view to the array_view given by "dest", as if by calling "copy(*this, dest)" (see 5.3.2).	
<b>Parameters:</b>	

<i>dest</i>	An object of type <code>array_view&lt;T,N&gt;</code> to which to copy data from this array.
-------------	---

2478

```
T* array_view<T,1>::data() const restrict(amp,cpu)
const T* array_view<const T,1>::data() const restrict(amp,cpu)
```

Returns a pointer to the raw data underlying this `array_view`. This is only available on `array_views` of rank 1.

**Return Value:**

A (const) pointer to the first element in the linearized array.

2479

```
void array_view<T, N>::refresh() const
void array_view<const T, N>::refresh() const
```

Calling this member function informs the `array_view` that its bound memory has been modified outside the `array_view` interface. This will render all cached information stale.

2480

```
void array_view<T, N>::synchronize() const
```

Calling this member function synchronizes any modifications made to “this” `array_view` to its underlying data container. For example, for an `array_view` on system memory, if the contents of the view are modified on a remote `accelerator_view` through a `parallel_for_each` invocation, calling `synchronize` ensures that the modifications are synchronized to the source data and will be visible through the system memory pointer which the `array_view` was created over.

2481

```
completion_future array_view<T, N>::synchronize_async() const
```

An asynchronous version of `synchronize`, which returns a completion future object. When the future is ready, the synchronization operation is complete.

2482

```
void array_view<T, N>::discard_data() const
```

Indicates to the runtime that it may discard the current logical contents of this `array_view`. This is an optimization hint to the runtime used to avoid copying the current contents of the view to a target `accelerator_view`, and its use is recommended if the existing content is not needed.

2483

2484 

## 5.2.4 Indexing

2485

2486 Accessing an `array_view` out of bounds yields undefined results.

2487

```
T& array_view<T,N>::operator[](const index<N>& idx) const restrict(amp,cpu)
T& array_view<T,N>::operator()(const index<N>& idx) const restrict(amp,cpu)
```

Returns a reference to the element of this `array_view` that is at the location in N-dimensional space specified by “idx”.

**Parameters:**

<i>Idx</i>	An object of type <code>index&lt;N&gt;</code> from that specifies the location of the element.
------------	--

2488

```
const T& array_view<const T,N>::operator[](const index<N>& idx) const restrict(amp,cpu)
const T& array_view<const T,N>::operator()(const index<N>& idx) const restrict(amp,cpu)
```

Returns a const reference to the element of this `array_view` that is at the location in N-dimensional space specified by “idx”.

**Parameters:**

<i>Idx</i>	An object of type <code>index&lt;N&gt;</code> from that specifies the location of the element.
------------	--

2489

```
T& array_view<T,1>::operator()(int i0) const restrict(amp,cpu)
T& array_view<T,1>::operator[](int i0) const restrict(amp,cpu)
T& array_view<T,2>::operator()(int i0, int i1) const restrict(amp,cpu)
T& array_view<T,3>::operator()(int i0, int i1, int i2) const restrict(amp,cpu)
```

Equivalent to “`array_view<T,N>::operator()(index<N>(i0 [, i1 [, i2 ]))`”.

**Parameters:**

<i>i0 [, i1 [, i2 ]]</i>	The component values that will form the index into this array.
--------------------------	--

2490

```
const T& array_view<const T,1>::operator()(int i0) const restrict(amp,cpu)
const T& array_view<const T,2>::operator()(int i0, int i1) const restrict(amp,cpu)
const T& array_view<const T,3>::operator()(int i0, int i1, int i2) const restrict(amp,cpu)
```

Equivalent to "array\_view<T,N>::operator()(index<N>(i0 [, i1 [, i2 ]])) const".

**Parameters:**

*i0 [, i1 [, i2 ]]*

The component values that will form the index into this array.

2491

```
array_view<T,N-1> array_view<T,N>::operator[](int i0) const restrict(amp,cpu)
array_view<const T,N-1> array_view<const T,N>::operator[](int i0) const restrict(amp,cpu)
```

This overload is defined for array\_view<T,N> where  $N \geq 2$ .

This mode of indexing is equivalent to projecting on the most-significant dimension. It allows C-style indexing. For example:

```
array<float,4> myArray(myExtents, ...);

myArray[index<4>(5,4,3,2)] = 7;
assert(myArray[5][4][3][2] == 7);
```

**Parameters:**

*i0*

An integer that is the index into the most-significant dimension of this array.

**Return Value:**

Returns an array\_view whose dimension is one lower than that of this array\_view.

2492

## 5.2.5 View Operations

2493

2494

```
array_view<T,N> array_view<T,N>::section(const index<N>& idx, const extent<N>& ext) const
restrict(amp,cpu)
array_view<const T,N> array_view<const T,N>::section(const index<N>& idx, const extent<N>& ext)
const restrict(amp,cpu)
```

Returns a subsection of the source array view at the origin specified by "idx" and with the extent specified by "ext"

Example:

```
array<float,2> a(extent<2>(200,100));
array_view<float,2> v1(a); // v1.extent = <200,100>
array_view<float,2> v2 = v1.section(index<2>(15,25), extent<2>(40,50));
assert(v2(0,0) == v1(15,25));
```

**Parameters:**

*idx*

Provides the offset/origin of the resulting section.

*ext*

Provides the extent of the resulting section.

**Return Value:**

Returns a subsection of the source array at specified origin, and with the specified extent.

2495

```
array_view<T,N> array_view<T,N>::section(const index<N>& idx) const restrict(amp,cpu)
array_view<const T,N> array_view<const T,N>::section(const index<N>& idx) const
restrict(amp,cpu)
```

Equivalent to "section(idx, this->extent - idx)".

2496

2497

```
array_view<T,N> array_view<T,N>::section(const extent<N>& ext) const restrict(amp,cpu)
array_view<const T,N> array_view<const T,N>::section(const extent<N>& ext) const
restrict(amp,cpu)
```

Equivalent to "section(index<N>(), extent)".

```
array_view<T,1> array_view<T,1>::section(int i0, int e0) const restrict(amp,cpu)
array_view<const T,1> array_view<const T,1>::section(int i0, int e0) const restrict(amp,cpu)

array_view<T,2> array_view<T,2>::section(int i0, int i1, int e0, int e1) const restrict(amp,cpu)
array_view<const T,2> array_view<const T,2>::section(int i0, int i1,
                                                    int e0, int e1) const restrict(amp,cpu)

array_view<T,3> array_view<T,3>::section(int i0, int i1, int i2,
                                        int e0, int e1, int e2) const restrict(amp,cpu)
array_view<const T,3> array_view<const T,3>::section(int i0, int i1, int i2,
                                                    int e0, int e1, int e2) const restrict(amp,cpu)
```

Equivalent to "section(index<N>(i0 [, i1 [, i2 ]]), extent<N>(e0 [, e1 [, e2 ]]))".

**Parameters:**

<i>i0</i> [, <i>i1</i> [, <i>i2</i> ]]	The component values that will form the origin of the section
<i>e0</i> [, <i>e1</i> [, <i>e2</i> ]]	The component values that will form the extent of the section

```
template<typename ElementType>
array_view<ElementType,1> array_view<T,1>::reinterpret_as() const restrict(amp,cpu)
template<typename ElementType>
array_view<const ElementType,1> array_view<const T,1>::reinterpret_as() const
restrict(amp,cpu)
```

This member function is similar to "array<T,N>::reinterpret\_as" (see 5.1.5), although it only supports array\_views of rank 1 (only those guarantee that all elements are laid out contiguously).

The size of the reinterpreted ElementType must evenly divide into the total size of this array\_view.

**Return Value:**

Returns an array\_view from this array\_view<T,1> with the element type reinterpreted from T to ElementType.

```
template <int K>
array_view<T,K> array_view<T,1>::view_as(extent<K> viewExtent) const restrict(amp,cpu)
template <int K>
array_view<const T,K> array_view<const T,1>::view_as(extent<K> viewExtent) const
restrict(amp,cpu)
```

This member function is similar to array<T,N>::view\_as" (see 5.1.5), although it only supports array\_views of rank 1 (only those guarantee that all elements are laid out contiguously).

**Return Value:**

Returns an array\_view from this array\_view<T,1> with the rank changed to K from 1.

### 5.3 Copying Data

C++ AMP offers a universal [copy](#) function which covers all synchronous data transfer requirements. In call cases, copying data is not supported while executing on an accelerator (in other words, the copy functions do not have a [restrict\(amp\)](#) clause). The general form of copy is:

```
copy(src, dest);
```

*Informative: Note that this more closely follows the STL convention (destination is the last argument, as in [std::copy](#)) and is opposite of the C-style convention (destination is the first argument, as in [memcpy](#)).*

Copying to [array](#) and [array\\_view](#) types is supported from the following sources:

- An [array](#) or [array\\_view](#) with the same rank and element type as the destination [array](#) or [array\\_view](#).
- A standard container whose element type is the same as the destination [array](#) or [array\\_view](#).

*Informative: Containers that expose `.size()` and `.data()` members (e.g., `std::vector`, and `std::array`) can be handled more efficiently.*

The copy operation always performs a deep copy.

Asynchronous copy has the same semantics as synchronous copy, except that they return a `completion_future` that can be waited on.

### 5.3.1 Synopsis

```

template <typename T, int N>
    void copy(const array<T,N>& src, array<T,N>& dest);
template <typename T, int N>
    void copy(const array<T,N>& src, const array_view<T,N>& dest);

template <typename T, int N>
    void copy(const array_view<const T,N>& src, array<T,N>& dest);
template <typename T, int N>
    void copy(const array_view<const T,N>& src, const array_view<T,N>& dest);

template <typename T, int N>
    void copy(const array_view<T,N>& src, array<T,N>& dest);
template <typename T, int N>
    void copy(const array_view<T,N>& src, const array_view<T,N>& dest);

template <typename InputIter, typename T, int N>
    void copy(InputIter srcBegin, InputIter srcEnd, array<T,N>& dest);
template <typename InputIter, typename T, int N>
    void copy(InputIter srcBegin, InputIter srcEnd, const array_view<T,N>& dest);

template <typename InputIter, typename T, int N>
    void copy(InputIter srcBegin, array<T,N>& dest);
template <typename InputIter, typename T, int N>
    void copy(InputIter srcBegin, const array_view<T,N>& dest);

template <typename OutputIter, typename T, int N>
    void copy(const array<T,N>& src, OutputIter destBegin);
template <typename OutputIter, typename T, int N>
    void copy(const array_view<T,N>& src, OutputIter destBegin);

template <typename T, int N>
    completion_future copy_async(const array<T,N>& src, array<T,N>& dest);
template <typename T, int N>
    completion_future copy_async(const array<T,N>& src, const array_view<T,N>& dest);

template <typename T, int N>
    completion_future copy_async(const array_view<const T,N>& src, array<T,N>& dest);
template <typename T, int N>
    completion_future copy_async(const array_view<const T,N>& src, const array_view<T,N>& dest);

template <typename T, int N>
    completion_future copy_async(const array_view<T,N>& src, array<T,N>& dest);

```

```

2569 template <typename T, int N>
2570     completion_future copy_async(const array_view<T,N>& src, const array_view<T,N>& dest);
2571
2572 template <typename InputIter, typename T, int N>
2573     completion_future copy_async(InputIter srcBegin, InputIter srcEnd, array<T,N>& dest);
2574 template <typename InputIter, typename T, int N>
2575     completion_future copy_async(InputIter srcBegin, InputIter srcEnd, const array_view<T,N>&
2576     dest);
2577
2578 template <typename InputIter, typename T, int N>
2579     completion_future copy_async(InputIter srcBegin, array<T,N>& dest);
2580 template <typename InputIter, typename T, int N>
2581     completion_future copy_async(InputIter srcBegin, const array_view<T,N>& dest);
2582
2583 template <typename OutputIter, typename T, int N>
2584     completion_future copy_async(const array<T,N>& src, OutputIter destBegin);
2585 template <typename OutputIter, typename T, int N>
2586     completion_future copy_async(const array_view<T,N>& src, OutputIter destBegin);
2587

```

### 5.3.2 Copying between array and array\_view

An `array<T,N>` can be copied to an object of type `array_view<T,N>`, and vice versa.

```

template <typename T, int N>
    void copy(const array<T,N>& src, array<T,N>& dest)

```

```

template <typename T, int N>
    completion_future copy_async(const array<T,N>& src, array<T,N>& dest)

```

The contents of "src" are copied into "dest". The source and destination may reside on different accelerators. If the extents of "src" and "dest" don't match, a runtime exception is thrown.

#### Parameters:

<i>Src</i>	An object of type <code>array&lt;T,N&gt;</code> to be copied from.
<i>Dest</i>	An object of type <code>array&lt;T,N&gt;</code> to be copied to.

```

template <typename T, int N>
    void copy(const array<T,N>& src, const array_view<T,N>& dest)

```

```

template <typename T, int N>
    completion_future copy_async(const array<T,N>& src, const array_view<T,N>& dest)

```

The contents of "src" are copied into "dest". If the extents of "src" and "dest" don't match, a runtime exception is thrown.

#### Parameters:

<i>src</i>	An object of type <code>array&lt;T,N&gt;</code> to be copied from.
<i>dest</i>	An object of type <code>array_view&lt;T,N&gt;</code> to be copied to.

```

template <typename T, int N>
    void copy(const array_view<const T,N>& src, array<T,N>& dest)

```

```

template <typename T, int N>
    void copy(const array_view<T,N>& src, array<T,N>& dest)

```

```
template <typename T, int N>
    completion_future copy_async(const array_view<const T,N>& src, array<T,N>& dest)
```

```
template <typename T, int N>
    completion_future copy_async(const array_view<T,N>& src, array<T,N>& dest)
```

The contents of "src" are copied into "dest". If the extents of "src" and "dest" don't match, a runtime exception is thrown.

**Parameters:**

<i>src</i>	An object of type <code>array_view&lt;T,N&gt;</code> (or <code>array_view&lt;const T,N&gt;</code> ) to be copied from.
<i>dest</i>	An object of type <code>array&lt;T,N&gt;</code> to be copied to.

2594

```
template <typename T, int N>
    void copy(const array_view<const T,N>& src, const array_view<T,N>& dest)
```

```
template <typename T, int N>
    completion_future copy_async(const array_view<const T,N>& src, const array_view<T,N>& dest)
```

The contents of "src" are copied into "dest". If the extents of "src" and "dest" don't match, a runtime exception is thrown.

**Parameters:**

<i>src</i>	An object of type <code>array_view&lt;T,N&gt;</code> (or <code>array_view&lt;const T,N&gt;</code> ) to be copied from.
<i>dest</i>	An object of type <code>array_view&lt;T,N&gt;</code> to be copied to.

2595

2596

### 5.3.3 Copying from standard containers to arrays or array\_views

A standard container can be copied into an [array](#) or [array\\_view](#) by specifying an iterator range.

*Informative: Standard containers that present a `.size()` and a `.data()` (such as `std::vector` and `std::array`) operation can be handled very efficiently.*

2597

2598

2599

2600

2601

2602

```
template <typename InputIter, typename T, int N>
    void copy(InputIter srcBegin, InputIter srcEnd, array<T,N>& dest)
```

```
template <typename InputIter, typename T, int N>
    void copy(InputIter srcBegin, array<T,N>& dest)
```

```
template <typename InputIter, typename T, int N>
    completion_future copy_async(InputIter srcBegin, InputIter srcEnd, array<T,N>& dest)
```

```
template <typename InputIter, typename T, int N>
    completion_future copy_async(InputIter srcBegin, array<T,N>& dest)
```

The contents of a source container from the iterator range [srcBegin,srcEnd) are copied into "dest". If the number of elements in the iterator range is not equal to "dest.extent.size()", an exception is thrown.

In the overloads which don't take an end-iterator it is assumed that the source iterator is able to provide at least `dest.extent.size()` elements, but no checking is performed (nor possible).

**Parameters:**

<i>srcBegin</i>	An iterator to the first element of a source container.
<i>srcEnd</i>	An iterator to the end of a source container.
<i>dest</i>	An object of type <code>array&lt;T,N&gt;</code> to be copied to.



2603

```

template <typename InputIter, typename T, int N>
    void copy(InputIter srcBegin, InputIter srcEnd, const array_view<T,N>& dest)

template <typename InputIter, typename T, int N>
    void copy(InputIter srcBegin, const array_view<T,N>& dest)

template <typename InputIter, typename T, int N>
    completion_future copy_async(InputIter srcBegin, InputIter srcEnd, const array_view<T,N>&
dest)

template <typename InputIter, typename T, int N>
    completion_future copy_async(InputIter srcBegin, const array_view<T,N>& dest)

```

The contents of a source container from the iterator range [srcBegin,srcEnd) are copied into "dest". If the number of elements in the iterator range is not equal to "dest.extent.size()", an exception is thrown.

**Parameters:**

<i>srcBegin</i>	An iterator to the first element of a source container.
<i>srcEnd</i>	An iterator to the end of a source container.
<i>Dest</i>	An object of type <code>array_view&lt;T,N&gt;</code> to be copied to.

2604

### 2605 5.3.4 Copying from arrays or array\_views to standard containers

2606

2607 An array or array\_view can be copied into a standard container by specifying the begin iterator. Standard containers that  
 2608 present a `.size()` and a `.data()` (such as `std::vector` and `std::array`) operation can be handled very  
 2609 efficiently.

2610

```

template <typename OutputIter, typename T, int N>
    void copy(const array<T,N>& src, OutputIter destBegin)

template <typename OutputIter, typename T, int N>
    completion_future copy_async(const array<T,N>& src, OutputIter destBegin)

```

The contents of a source array are copied into "dest" starting with iterator destBegin. If the number of elements in the range starting destBegin in the destination container is smaller than "src.extent.size()", an exception is thrown.

**Parameters:**

<i>src</i>	An object of type <code>array&lt;T,N&gt;</code> to be copied from.
<i>destBegin</i>	An output iterator addressing the position of the first element in the destination container.

2611

```

template <typename OutputIter, typename T, int N>
    void copy(const array_view<T,N>& src, OutputIter destBegin)

template <typename OutputIter, typename T, int N>
    completion_future copy_async(const array_view<T,N>& src, OutputIter destBegin)

```

The contents of a source array are copied into "dest" starting with iterator destBegin. If the number of elements in the range starting destBegin in the destination container is smaller than "src.extent.size()", an exception is thrown.

**Parameters:**

<i>src</i>	An object of type <code>array_view&lt;T,N&gt;</code> to be copied from.
------------	---

*destBegin*

An output iterator addressing the position of the first element in the destination container.

2612

## 2613 6 Atomic Operations

2614 C++ AMP provides a set of atomic operations in the [concurrency](#) namespace. These operations are applicable in  
 2615 [restrict\(amp\)](#) contexts and may be applied to memory locations within [concurrency::array](#) instances and to memory  
 2616 locations within [tile\\_static](#) variables. Section 8 provides a full description of the C++ AMP memory model and how atomic  
 2617 operations fit into it.

### 2618 6.1 Synopsis

```

2619
2620 int atomic_exchange(int * dest, int val) restrict(amp)
2621 unsigned int atomic_exchange(unsigned int * dest, unsigned int val) restrict(amp)
2622 float atomic_exchange(float * dest, float val) restrict(amp)
2623
2624 bool atomic_compare_exchange(int * dest, int * expected_value, int val) restrict(amp)
2625 bool atomic_compare_exchange(unsigned int * dest, unsigned int * expected_value, unsigned int
2626 val) restrict(amp)
2627
2628 int atomic_fetch_add(int * dest, int val) restrict(amp)
2629 unsigned int atomic_fetch_add(unsigned int * dest, unsigned int val) restrict(amp)
2630
2631 int atomic_fetch_sub(int * dest, int val) restrict(amp)
2632 unsigned int atomic_fetch_sub(unsigned int * dest, unsigned int val) restrict(amp)
2633
2634 int atomic_fetch_max(int * dest, int val) restrict(amp)
2635 unsigned int atomic_fetch_max(unsigned int * dest, unsigned int val)
2636
2637 int atomic_fetch_min(int * dest, int val) restrict(amp)
2638 unsigned int atomic_fetch_min(unsigned int * dest, unsigned int val)
2639
2640 int atomic_fetch_and(int * dest, int val) restrict(amp)
2641 unsigned int atomic_fetch_and(unsigned int * dest, unsigned int val)
2642
2643 int atomic_fetch_or(int * dest, int val) restrict(amp)
2644 unsigned int atomic_fetch_or(unsigned int * dest, unsigned int val)
2645
2646 int atomic_fetch_xor(int * dest, int val) restrict(amp)
2647 unsigned int atomic_fetch_xor(unsigned int * dest, unsigned int val) restrict(amp)
2648
2649 int atomic_fetch_inc(int * dest) restrict(amp)
2650 unsigned int atomic_fetch_inc(unsigned int * dest) restrict(amp)
2651
2652 int atomic_fetch_dec(int * dest) restrict(amp)
2653 unsigned int atomic_fetch_dec(unsigned int * dest) restrict(amp)
2654

```

### 2655 6.2 Atomically Exchanging Values

2656

```

int atomic_exchange(int * dest, int val) restrict(amp)
unsigned int atomic_exchange(unsigned int * dest, unsigned int val) restrict(amp)
float atomic_exchange(float * dest, float val) restrict(amp)

```

Atomically read the value stored in *dst*, replace it with the value given in *val* and return the old value to the caller. This function provides overloads for *int*, *unsigned int* and *float* parameters.

**Parameters:**

<i>dst</i>	An pointer to the location which needs to be atomically modified. The location may reside within a <i>concurrency::array</i> or within a <i>tile_static</i> variable.
<i>val</i>	The new value to be stored in the location pointed to be <i>dst</i> .

**Return value:**

These functions return the old value which was previously stored at *dst*, and that was atomically replaced. These functions always succeed.

2657

```
bool atomic_compare_exchange(int * dest, int * expected_val, int val) restrict(amp)
bool atomic_compare_exchange(unsigned int * dest, unsigned int * expected_val, unsigned int val)
restrict(amp)
```

These functions attempt to atomically perform these three steps atomically:

1. Read the value stored in the location pointed to by *dst*
2. Compare the value read in the previous step with the value contained in the location pointed by *expected\_val*
3. Carry the following operations depending on the result of the comparison of the previous step:
  - a. If the values are identical, then the function tries to atomically change the value pointed by *dst* to the value in *val*. The function indicates by its return value whether this transformation has been successful or not.
  - b. If the values are not identical, then the function stores the value read in step (1) into the location pointed to by *expected\_val*, and returns *false*.

In terms of sequential semantics, the function is equivalent to the following pseudo-code:

```
auto t = *dst;
bool eq = t == *expected_val;
if (eq)
    *dst = val;
*expected_val = t;
return eq;
```

The function may fail spuriously. It is guaranteed that the system as a whole will make progress when threads are contending to atomically modify a variable, but there is no upper bound on the number of failed attempts that any particular thread may experience.

**Parameters:**

<i>dst</i>	An pointer to the location which needs to be atomically modified. The location may reside within a <i>concurrency::array</i> or within a <i>tile_static</i> variable.
<i>expected_val</i>	A pointer to a local variable or function parameter. Upon calling the function, the location pointed by <i>expected_val</i> contains the value the caller expects <i>dst</i> to contain. Upon return from the function, <i>expected_val</i> will contain the most recent value read from <i>dst</i> .
<i>val</i>	The new value to be stored in the location pointed to be <i>dst</i> .

**Return value:**

The return value indicates whether the function has been successful in atomically reading, comparing and modifying the contents of the memory location.

### 2658 6.3 Atomically Applying an Integer Numerical Operation

2659

```
int atomic_fetch_add(int * dest, int val) restrict(amp)
unsigned int atomic_fetch_add(unsigned int * dest, unsigned int val) restrict(amp)
```

```

int atomic_fetch_sub(int * dest, int val) restrict(amp)
unsigned int atomic_fetch_sub(unsigned int * dest, unsigned int val) restrict(amp)

int atomic_fetch_max(int * dest, int val) restrict(amp)
unsigned int atomic_fetch_max(unsigned int * dest, unsigned int val)

int atomic_fetch_min(int * dest, int val) restrict(amp)
unsigned int atomic_fetch_min(unsigned int * dest, unsigned int val)

int atomic_fetch_and(int * dest, int val) restrict(amp)
unsigned int atomic_fetch_and(unsigned int * dest, unsigned int val)

int atomic_fetch_or(int * dest, int val) restrict(amp)
unsigned int atomic_fetch_or(unsigned int * dest, unsigned int val)

int atomic_fetch_xor(int * dest, int val) restrict(amp)
unsigned int atomic_fetch_xor(unsigned int * dest, unsigned int val) restrict(amp)

```

Atomically read the value stored in *dest*, apply the binary numerical operation specific to the function with the read value and *val* serving as input operands, and store the result back to the location pointed by *dest*.

In terms of sequential semantics, the operation performed by any of the above function is described by the following piece of pseudo-code:

```
*dest = *dest  $\otimes$  val;
```

Where the operation denoted by  $\otimes$  is one of: addition (`atomic_fetch_add`), subtraction (`atomic_fetch_sub`), find maximum (`atomic_fetch_max`), find minimum (`atomic_fetch_min`), bit-wise AND (`atomic_fetch_and`), bit-wise OR (`atomic_fetch_or`), bit-wise XOR (`atomic_fetch_xor`).

#### Parameters:

<i>Dst</i>	An pointer to the location which needs to be atomically modified. The location may reside within a <i>concurrency::array</i> or within a <i>tile_static</i> variable.
<i>val</i>	The second operand which participates in the calculation of the binary operation whose result is stored into the location pointed to be <i>dst</i> .

#### Return value:

These functions return the old value which was previously stored at *dst*, and that was atomically replaced. These functions always succeed.

2660

```

int atomic_fetch_inc(int * dest) restrict(amp)
unsigned int atomic_fetch_inc(unsigned int * dest) restrict(amp)

int atomic_fetch_dec(int * dest) restrict(amp)
unsigned int atomic_fetch_dec(unsigned int * dest) restrict(amp)

```

Atomically increment or decrement the value stored at the location point to by *dest*.

#### Parameters:

<i>Dst</i>	An pointer to the location which needs to be atomically modified. The location may reside within a <i>concurrency::array</i> or within a <i>tile_static</i> variable.
------------	---

#### Return value:

These functions return the old value which was previously stored at *dst*, and that was atomically replaced. These functions always succeed.

## 2661 7 Launching Computations: parallel\_for\_each

2662

Developers using C++ AMP will use a form of `parallel_for_each()` to launch data-parallel computations on accelerators. The behavior of `parallel_for_each` is similar to that of `std::for_each`: execute a function for each element in a container. The C++ AMP specialization over containers of type `extent` and `tiled_extent` allow execution of functions on accelerators.

The `parallel_for_each` function takes the following general forms:

1. Non-tiled:

```
template <int N, typename Kernel>
void parallel_for_each(extent<N> compute_domain, const Kernel& f);
```

2. Tiled:

```
template <int D0, int D1, int D2, typename Kernel>
void parallel_for_each(tiled_extent<D0,D1,D2> compute_domain, const Kernel& f);
```

```
template <int D0, int D1, typename Kernel>
void parallel_for_each(tiled_extent<D0,D1> compute_domain, const Kernel& f);
```

```
template <int D0, typename Kernel>
void parallel_for_each(tiled_extent<D0> compute_domain, const Kernel& f);
```

A `parallel_for_each` invocation may be explicitly requested on a specific accelerator view

1. Non-tiled:

```
template <int N, typename Kernel>
void parallel_for_each(const accelerator_view& accl_view,
    extent<N> compute_domain, const Kernel& f);
```

2. Tiled:

```
template <int D0, int D1, int D2, typename Kernel>
void parallel_for_each(const accelerator_view& accl_view,
    tiled_extent<D0,D1,D2> compute_domain, const Kernel& f);
```

```
template <int D0, int D1, typename Kernel>
void parallel_for_each(const accelerator_view& accl_view,
    tiled_extent<D0,D1> compute_domain, const Kernel& f);
```

```
template <int D0, typename Kernel>
void parallel_for_each(const accelerator_view& accl_view,
    tiled_extent<D0> compute_domain, const Kernel& f);
```

A `parallel_for_each` over an `extent` represents a dense loop nest of independent serial loops.

When `parallel_for_each` executes, a parallel activity is spawned for each index in the compute domain. Each parallel activity is associated with an index value. (This index is an `index<N>` in the case of a non-tiled `parallel_for_each`, or a `tiled_index<D0,D1,D2>` in the case of a tiled `parallel_for_each`.) A parallel activity typically uses its index to access the appropriate locations in the input/output arrays.

A call to `parallel_for_each` behaves as if it were synchronous. In practice, the call may be asynchronous because it executes on a separate device, but since data copy-out is a synchronizing event, the developer cannot tell the difference.

There are no guarantees on the order and concurrency of the parallel activities spawned by the non-tiled `parallel_for_each`. Thus it is not valid to assume that one activity can wait for another sibling activity to complete for itself to make progress. This is discussed in further detail in section 8.

The tiled version of `parallel_for_each` organizes the parallel activities into fixed-size tiles of 1, 2, or 3 dimensions, as given by the `tiled_extent<>` argument. The `tiled_extent` provided as the first parameter to `parallel_for_each` must be divisible, along

each of its dimensions, by the respective tile extent. Tiling beyond 3 dimensions is not supported. Threads (parallel activities) in the same tile have access to shared `tile_static` memory, and can use `tiled_index::barrier.wait` (4.5.3) to synchronize access to it.

When launching an `amp`-restricted kernel, the implementation of tiled `parallel_for_each` will provide the following minimum capabilities:

- The maximum number of tiles per dimension will be no less than 65535.
- The maximum number of threads in a tile will be no less than 1024.
  - In 3D tiling, the maximal value of D0 will be no less than 64.

**Microsoft-specific:**

When launching an `amp`-restricted kernel, the tiled `parallel_for_each` provides the above portable guarantees and no more. i.e.,

- The maximum number of tiles per dimension is 65535.
- The maximum number of threads in a tile is 1024
  - In 3D tiling, the maximum value supported for D0 is 64.

The execution behind the `parallel_for_each` occurs on a certain accelerator, in the context of a certain accelerator view. This accelerator view may be passed explicitly to `parallel_for_each` (as an optional first argument). Otherwise, the target accelerator and the view using which work is submitted to the accelerator, is chosen from the objects of type `array<T,N>` and `texture<T>` that were captured in the kernel lambda. An implementation may require that all arrays and textures captured in the lambda must be on the same accelerator view; if not, an implementation is free to throw an exception. An implementation may also arrange for the specified data to be accessible on the selected accelerator view, rather than reject the call.

*Microsoft-specific: the Microsoft implementation of C++ AMP requires that all array and texture objects are co-located on the same accelerator view which is used, implicitly or explicitly in a `parallel_for_each` call.*

If the `parallel_for_each` kernel functor does not capture an array/texture object and neither is the target `accelerator_view` for the kernel's execution is explicitly specified, the runtime is allowed to execute the kernel on any `accelerator_view` on the default accelerator.

*Microsoft-specific: In such a scenario, the Microsoft implementation of C++ AMP selects the target `accelerator_view` for executing the `parallel_for_each` kernel as follows:*

- a. Determine the set of `accelerator_views` where ALL `array_views` referenced in the `p_f_e` kernel have cached copies
- b. From the above set, filter out any `accelerator_views` that are not on the default accelerator. Additionally filter out `accelerator_views` that do not have the capabilities required by the `p_f_e` kernel (debug intrinsics, number of UAVs)
- c. The default `accelerator_view` of the default accelerator is selected as the target, if the resultant set from b. is empty, or contains, that `accelerator_view`

Otherwise, any `accelerator_view` from the resultant set from b., is arbitrarily selected as the target. The `tiled_index<>` argument passed to the kernel contains a collection of indices including those that are relative to the current tile.

The argument `f` of template-argument type `Kernel` to the `parallel_for_each` function must be a lambda or functor offering an appropriate function call operator which the implementation of `parallel_for_each` invokes with the instantiated index type. To execute on an accelerator, the function call operator must be marked `restrict(amp)` (but may have additional

restrictions), and it must be callable from a caller passing in the instantiated index type. Overload resolution is handled as if the caller contained this code:

```
template <typename IndexType, typename Kernel>
void parallel_for_each_stub(IndexType i, Kernel f) restrict(amp)
{
    f(i);
}
```

Where the *Kernel f* argument is the same one passed into *parallel\_for\_each* by the caller, and the index instance *i* is the thread identifier, where *IndexType* is the following type:

- Non-Tiled *parallel\_for\_each*: *index<N>*, where *N* must be the same rank as the *extent<N>* used in the *parallel\_for\_each*.
- Tiled *parallel\_for\_each*: *tiled\_index<D0 [, D1 [, D2]]>*, where the tile extents must match those of the *tiled\_extent* used in the *parallel\_for\_each*.

The value returned by the kernel function, if any, is ignored.

#### Microsoft-specific:

*In the Microsoft implementation of C++ AMP, every function that is referenced directly or indirectly by the kernel function, as well as the kernel function itself, must be inlineable<sup>4</sup>.*

## 7.1 Capturing Data in the Kernel Function Object

Since the kernel function object does not take any other arguments, all other data operated on by the kernel, other than the thread index, must be captured in the lambda or function object passed to *parallel\_for\_each*. The function object shall be any amp-compatible class, struct or union type, including those introduced by lambda expressions.

## 7.2 Exception Behaviour

If an error occurs trying to launch the *parallel\_for\_each*, an exception will be thrown. Exceptions can be thrown the following reasons:

1. Failure to create shader
2. Failure to create buffers
3. Invalid extent passed
4. Mismatched accelerators

## 8 Correctly Synchronized C++ AMP Programs

Correctly synchronized C++ AMP programs are correctly synchronized C++ programs which also adhere to a few additional C++ AMP rules, as follows:

1. Accelerator-side execution
  - a. Concurrency rules for arbitrary sibling threads launched by a *parallel\_for\_each* call.
  - b. Semantics and correctness of tile barriers.
  - c. Semantics of atomic and memory fence operations.
2. Host-side execution
  - a. Concurrency of accesses to C++ AMP containers between host-side operations: *copy*, *synchronize*, *parallel\_for\_each* and the application of the various subscript operators of arrays and array views on the host.

<sup>4</sup> An implementation can employ whole-program compilation (such as link-time code-gen) to achieve this.



- b. Accessing *arrays* or *array\_view* data on the host.

## 8.1 Concurrency of sibling threads launched by a *parallel\_for\_each* call

In this section we will consider the relationship between sibling threads in a single *parallel\_for\_each* call. Interaction between separate *parallel\_for\_each* calls, copy operations and other host-side operations will be considered in the following sub-sections.

A *parallel\_for\_each* call logically initiates the operation of multiple sibling threads, one for each coordinate in the *extent* or *tiled\_extent* passed to it.

All the threads launched by a *parallel\_for\_each* are potentially concurrent. Unless barriers are used, an implementation is free to schedule these threads in any order. In addition, the memory model for normal memory accesses is weak, that is operations could be arbitrarily reordered as long as each thread perceives to execute in its original program order. Thus any two memory operations from any two threads in a *parallel\_for\_each* are by default concurrent, unless the application has explicitly enforced an order between these two operations using atomic operations, fences or barriers.

Conversely, an implementation may also schedule only a single logical thread at a time, in a non-cooperative manner, i.e., without letting any other threads make any progress, with the exception of hitting a tile barrier or terminating. When a thread encounters a tile barrier, an implementation must wrest control from that thread and provide progress to some other thread in the tile until they all have reached the barrier. Similarly, when a thread finishes execution, the system is obligated to execute steps from some other thread. Thus an implementation is obligated to switch context between threads only when a thread has hit a barrier (barriers pertain just to the tiled *parallel\_for\_each*), or is finished. An implementation doesn't have to admit any concurrency at a finer level than that which is dictated by barriers and thread termination. All implementations, however, are obligated to ensure progress is continually made, until all threads launched by a *parallel\_for\_each* are completed.

An immediate corollary is that C++ AMP doesn't provide a mechanism using which a thread could, without using tile barriers, poll for a change which needs to be effected by another thread. In particular, C++ AMP doesn't support locks which are implemented using atomic operations and fences, since a thread could end up polling forever, waiting for a lock to become available. The usage of tile barriers allows for creating a limited form of locking scoped to a thread tile. For example:

```
void tile_lock_example()
{
    parallel_for_each(
        extent<1>(TILE_SIZE).tile<TILE_SIZE>(),
        [] (tiled_index<TILE_SIZE> tid) restrict(amp)
        {
            tile_static int lock;

            // Initialize lock:
            if (tid.local[0] == 0) lock = 0;
            tid.barrier.wait();

            bool performed_my_exclusive_work = false;
            for (;;) {
                // try to acquire the lock
                if (!performed_my_exclusive_work && atomic_compare_exchange(&lock, 0, 1)) {
                    // The lock has been acquired - mutual exclusion from the rest of the threads in the tile
                    // is provided here...
                    some_synchronized_op();

                    // Release the lock
                    atomic_exchange(&lock, 0);
                    performed_my_exclusive_work = true;
                }
                else {
                    // The lock wasn't acquired, or we are already finished. Perhaps we can do something
                    // else in the meanwhile.
                }
            }
        }
    );
}
```



```

2866         some_non_exclusive_op();
2867     }
2868
2869     // The tile barrier ensures progress, so threads can spin in the for loop until they
2870     // are successful in acquiring the lock.
2871     tidx.barrier.wait();
2872 }
2873 });
2874 }
2875

```

*Informative: More often than not, such non-deterministic locking within a tile is not really necessary, since a static schedule of the threads based on integer thread ID's is possible and results in more efficient and more maintainable code, but we bring this example here for completeness and to illustrate a valid form of polling.*

### 8.1.1 Correct usage of tile barriers

Correct C++ AMP programs require all threads in a tile to hit all tile barriers uniformly. That is, at a minimum, when a thread encounters a particular `tile_barrier::wait` call site (or any other barrier method of class `tile_barrier`), all other threads in the tile must encounter the same call site.

*Informative: This requirement, however, is typically not sufficient in order to allow for efficient implementations. For example, it allows for the call stack of threads to differ, when they hit a barrier. In order to be able to generate good quality code for vector targets, much stronger constraints should be placed on the usage of barriers, as explained below.*

C++ AMP requires all *active control flow expressions* leading to a tile barrier to be *tile-uniform*. Active control flow expressions are those guarding the scopes of all control flow constructs and logical expressions, which are actively being executed at a time a barrier is called. For example, the condition of an *if* statement is an active control flow expression as long as either the true or false hands of the *if* statement are still executing. If either of those hands contains a tile barrier, or leads to one through an arbitrary nesting of scopes and function calls, then the control flow expression controlling the *if* statement must be *tile-uniform*. What follows is an exhaustive list of control flow constructs which may lead to a barrier and their corresponding control expressions:

```

2895     if (<control-expression>) <statement> else <statement>
2896     switch (<control-expression> { <cases> }
2897     for (<init-expression>; <control-expression>; <iteration-expression>) <statement>
2898     while (<control-expression>) <statement>
2899     do <statement> while(<control-expression>);
2900     <control-expression> ? <expression> : <expression>
2901     <control-expression> && <expression>
2902     <control-expression> || <expression>
2903

```

All active control flow constructs are strictly nested in accordance to the program's text, starting from the scope of the lambda at the `parallel_for_each` all the way to the scope containing the barrier.

C++ AMP requires that, when a barrier is encountered by one thread:

1. That the same barrier will be encountered by all other threads in the tile.
2. That the sequence of active control flow statements and/or expressions be identical for all threads when they reach the barrier.
3. That each of the corresponding control expressions be *tile-uniform* (which is defined below).
4. That any active control flow statement or expression hasn't been departed (necessarily in a non-uniform fashion) by a *break*, *continue* or *return* statement. That is, any breaking statement which instructs the program to leave an active scope must in itself behave as if it was a barrier, i.e., adhere to these four preceding rules.

Informally, a *tile-uniform expression* is an expression only involving variables, literals and function calls which have a uniform value throughout the tile. Formally, C++ AMP specifies that:

5. *Tile-uniform* expressions may reference literals and template parameters

6. *Tile-uniform* expressions may reference *const* (or effectively *const*) data members of the function object parameter of *parallel\_for\_each*
7. *Tile-uniform* expressions may reference *tiled\_index<,,>::tile*
8. *Tile-uniform* expressions may reference values loaded from *tile\_static* variables as long as those values are loaded immediately and uniformly after a tile barrier. That is, if the barrier and the load of the value occur at the same function and the barrier dominates the load and no potential store into the same *tile\_static* variable intervenes between the barrier and the load, then the loaded value will be considered *tile-uniform*
9. Control expressions may reference *tile-uniform local variables and parameters*. Uniform local variables and parameters are variables and parameters which are always initialized and assigned-to under uniform control flow (that is, using the same rules which are defined here for barriers) and which are only assigned *tile-uniform* expressions
10. *Tile-uniform* expressions may reference the return values of functions which return *tile-uniform* expressions
11. *Tile-uniform* expressions may not reference any expression not explicitly listed by the previous rules

An implementation is not obligated to warn when a barrier does not meet the criteria set forth above. An implementation may disqualify the compilation of programs which contain incorrect barrier usage. Conversely, an implementation may accept programs containing incorrect barrier usage and may execute them with undefined behavior.

### 8.1.2 Establishing order between operations of concurrent *parallel\_for\_each* threads

Threads may employ atomic operations, barriers and fences to establish a happens-before relationship encompassing their cumulative execution. When considering the correctness of the synchronization of programs, the following three aspects of the programs are relevant:

1. The types of memory which are potentially accessed concurrently by different threads. The memory type can be:
  - a. Global memory
  - b. Tile-static memory
2. The relationship between the threads which could potentially access the same piece of memory. They could be:
  - a. Within the same thread tile
  - b. Within separate threads tiles or sibling threads in the basic (non-tiled) *parallel\_for\_each* model.
3. Memory operations which the program contains:
  - a. Normal memory reads and writes.
  - b. Atomic read-modify-write operations.
  - c. Memory fences and barriers

Informally, the C++ AMP memory model is a weak memory model consistent with the C++ memory model, with the following exceptions:

1. Atomic operations do not necessarily create a sequentially consistent subset of execution. Atomic operations are only coherent, not sequentially consistent. That is, there doesn't necessarily exist a global linear order containing all atomic operations affecting all memory locations which were subjects of such operations. Rather, a separate global order exists for each memory location, and these per-location memory orders are not necessarily combinable into a single global order. (Note: this means an atomic operation does not constitute a memory fence.)
2. Memory fence operations are limited in their effects to the thread tile they are performed within. When a thread from tile A executes a fence, the fence operation doesn't necessarily affect any other thread from any tile other than A.
3. As a result of (1) and (2), the only mechanism available for cross-tile communication is atomic operations, and even when atomic operations are concerned, a linear order is only guaranteed to exist on a per-location basis, but not necessarily globally.
4. Fences are bi-directional, meaning they have both acquire and release semantics.
5. Fences can also be further scoped to a particular memory type (global vs. tile-static).
6. Applying normal stores and atomic operations concurrently to the same memory location results in undefined behavior.

7. Applying a normal load and an atomic operation concurrently to the same memory location is allowed (i.e., results in defined behavior).

We will now provide a more formal characterization of the different categories of programs based on their adherence to synchronization rules. The three classes of adherence are

1. *barrier-incorrect* programs,
2. *racy* programs, and,
3. *correctly-synchronized* programs.

#### 8.1.2.1 Barrier-incorrect programs

A *barrier-incorrect* program is a program which doesn't adhere to the correct barrier usage rules specified in the previous section. Such programs always have undefined behavior. The remainder of this section discusses barrier-correct programs only.

#### 8.1.2.2 Compatible memory operations

The following definition is later used in the definition of racy programs.

Two memory operations applied to the same (or overlapping) memory location are *compatible* if they are both aligned and have the same data width, and either both operations are reads, or both operations are atomic, or one operation is a read and the other is atomic.

This is summarized by the following table in which  $T_1$  is a thread executing  $Op_1$  and  $T_2$  is a thread executing operation  $Op_2$ .

$Op_1$	$Op_2$	Compatible?
Atomic	Atomic	Yes
Read	Read	Yes
Read	Atomic	Yes
Write	Any	No

#### 8.1.2.3 Concurrent memory operations

The following definition is later used in the definition of racy programs.

Informally, two memory operations by different threads are considered *concurrent* if no order has been established between them. Order can be established between two memory operations only when they are executed by threads within the same tile. Thus any two memory operations by threads from different tiles are always concurrent, even if they are atomic. Within the same tile, order is established using fences and barriers. Barriers are a strong form of a fence.

Formally, Let  $\{T_1, \dots, T_N\}$  be the threads of a tile. Fix a sharable memory type (be it global or tile-static). Let  $M$  be the total set of memory operations of the given memory type performed by the collective of the threads in the tile.

Let  $F = \langle F_1, \dots, F_L \rangle$  be the set of memory fence operations of the given memory type, performed by the collective of threads in the tile, and organized arbitrarily into an ordered sequence.

Let  $P$  be a partitioning of  $M$  into a sequence of subsets  $P = \langle M_0, \dots, M_L \rangle$ , organized into an ordered sequence in an arbitrary fashion.

Let  $S$  be the interleaving of  $F$  and  $P$ ,  $S = \langle M_0, F_1, M_1, \dots, F_L, M_L \rangle$

$S$  is *conforming* if both of these conditions hold:

1. **Adherence to program order:** For each  $T_i$ ,  $S$  respects the fences performed<sup>5</sup> by  $T_i$ . That is any operation performed by  $T_i$  before  $T_i$  performed fence  $F_j$  appears strictly before  $F_j$  in  $S$ , and similarly any operations performed by  $T_i$  after  $F_j$  appears strictly after  $F_j$  in  $S$ .
2. **Self-consistency:** For  $i < j$ , let  $M_i$  be a subset containing at least one store (atomic or non-atomic) into location  $L$  and let  $M_j$  be a subset containing at least a single load of  $L$ , and no stores into  $L$ . Further assume that no subset in-between  $M_i$  and  $M_j$  stores into  $L$ . Then  $S$  provides that all loads in  $M_j$  shall:
  - a. Return values stored into  $L$  by operations in  $M_i$ , and
  - b. For each thread  $T_i$ , the subset of  $T_i$  operations in  $M_j$  reading  $L$  shall all return the same value (which is necessarily one stored by an operation in  $M_i$ , as specified by condition (a) above).
3. **Respecting initial values.** Let  $M_j$  be a subset containing a load of  $L$ , and no stores into  $L$ . Further assume that there is no  $M_i$  where  $i < j$  such that  $M_i$  contains a store into  $L$ . Then all loads of  $L$  in  $M_j$  will return the initial value of  $L$ .

In such a conforming sequence  $S$ , two operations are *concurrent* if they have been executed by different threads and they belong to some common subset  $M_i$ . Two operations are *concurrent in an execution history* of a tile, if there exists a conforming interleaving  $S$  as described herein in which the operations are concurrent. Two operations of a program are *concurrent* if there possibly exists an execution of the program in which they are concurrent.

A barrier behaves like a fence to establish order between operations, except it provides additional guarantees on the order of execution. Based on the above definition, a barrier is like a fence that only permits a certain kind of interleaving. Specifically, one in which the sequence of fences ( $F$  in the above formalization) has the fences, corresponding to the barrier execution by individual threads, appearing uninterrupted in  $S$ , without any memory operations interleaved between them. For example, consider the following program:

```
C1
Barrier
C2
```

Assume that  $C1$  and  $C2$  are arbitrary sequences of code. Assume this program is executed by two threads  $T1$  and  $T2$ , then the only possible conforming interleavings are given by the following pattern:

```
T1(C1) || T2(C1)
T1(Barrier) || T2(Barrier)
T1(C2) || T2(C2)
```

Where the  $||$  operator implies arbitrary interleaving of the two operand sequences.

#### 8.1.2.4 Racy programs

*Racy programs* are programs which have possible executions where at least two operations performed by two separate threads are both (a) incompatible AND (b) concurrent.

Racy programs do not have semantics assigned to them. They have undefined behavior.

#### 8.1.2.5 Race-free programs

Race-free programs are, simply, programs that are not racy. Race-free programs have the following semantics assigned to them:

1. If two memory operations are ordered (i.e., not concurrent) by fences and/or barriers, then the values loaded/stored will respect such an ordering.

---

<sup>5</sup> Here, performance of memory operations is assumed to strictly follow program order.

2. If two memory operations are concurrent then they must be atomic and/or reads performed by threads within the same tile. For each memory location X there exists an eventual total order including all such operations concurrent operations applied to X and obeying the semantics of loads and atomic read-modify-write transactions.

## 8.2 Cumulative effects of a `parallel_for_each` call

An invocation of `parallel_for_each` receives a function object, the contents of which are made available on the device. The function object may contain: `concurrency::array` reference data members, `concurrency::array_view` value data members, `concurrency::texture` reference data members, and `concurrency::writeonly_texture_view` value data members. (In addition, the function object may also contain additional, user defined data members.) Each of these members of the types `array`, `array_view`, `texture` and `write_only_texture_view`, could be constrained in the type of access it provides to kernel code. For example an `array<int,2>&` member provides both read and write access to the array, while a `const array<int,2>&` member provides just read access to the array. Similarly, an `array_view<int,2>` member provides read and write access, while an `array_view<const int,2>` member provides read access only.

The C++ AMP specification permits implementations in which the memory backing an `array`, `array_view` or `texture` could be shared between different accelerators, and possibly also the host, while also permitting implementations where data has to be copied, by the implementation, between different memory regions in order to support access by some hardware. Simulating coherence at a very granular level is too expensive in the case disjoint memory regions are required by the hardware. Therefore, in order to support both styles of implementation, this specification stipulates that `parallel_for_each` has the freedom to implement coherence over `array`, `array_view`, and `texture` using coarse copying. Specifically, while a `parallel_for_each` call is being evaluated, implementations may:

1. Load and/or store any location, in any order, any number of times, of each container which is passed into `parallel_for_each` in read/write mode.
2. Load from any location, in any order, any number of times, of each container which is passed into `parallel_for_each` in read-only mode.

A `parallel_for_each` always behaves synchronously. That is, any observable side effects caused by any thread executing within a `parallel_for_each` call, or any side effects further affected by the implementation, due to the freedom it has in moving memory around, as stipulated above, shall be visible by the time `parallel_for_each` return.

However, since the effects of `parallel_for_each` are constrained to changing values within `arrays`, `array_views` and `textures` and each of these objects can synchronize its contents lazily upon access, an asynchronous implementation of `parallel_for_each` is possible, and encouraged. Nonetheless, implementations should still honor calls to `accelerator_view::wait` by blocking until all lazily queued side-effects have been fully performed. Similarly, an implementation should ensure that all lazily queued side-effects preceding an `accelerator_view::create_marker` call have been fully performed before the `completion_future` object which is returned by `create_marker` is made ready.

*Informative: Future versions of `parallel_for_each` may be less constrained in the changes they may affect to shared memory, and at that point an asynchronous implementation will no longer be valid. At that point, an explicitly asynchronous `parallel_for_each_async` will be added to the specification.*

Even though an implementation could be coarse in the way it implements coherence, it still must provide true aliasing for `array_views` which refer to the same home location. For example, assuming that `a1` and `a2` are both `array_views` constructed on top of a 100-wide one dimensional `array`, with `a1` referring to elements [0...10] of the `array` and `a2` referring to elements [10...20] of the same `array`. If both `a1` and `a2` are accessible on a `parallel_for_each` call, then accessing `a1` at position 10 is identical to accessing the view `a2` at position 0, since they both refer to the same location of the `array` they are providing a view over, namely position 10 in the original `array`. This rule holds whenever and wherever `a1` and `a2` are accessible simultaneously, i.e., on the host and in `parallel_for_each` calls.

Thus, for example, an implementation could clone an `array_view` passed into a `parallel_for_each` in read-only mode, and pass the cloned data to the device. It can create the clone using any order of reads from the original. The implementation may read the original a multiple number of times, perhaps in order to implement load-balancing or reliability features.

Similarly, an implementation could copy back results from an internally cloned [array](#), [array\\_view](#) or [texture](#), onto the original data. It may overwrite any data in the original container, and it can do so multiple times in the realization of a single [parallel\\_for\\_each](#) call.

When two or more overlapping array views are passed to a [parallel\\_for\\_each](#), an implementation could create a temporary array corresponding to a section of the original container which contains at a minimum the union of the views necessary for the call. This temporary array will hold the clones of the overlapping [array\\_views](#) while maintaining their aliasing requirements.

The guarantee regarding aliasing of [array\\_views](#) is provided for views which share the same *home location*. The home location of an [array\\_view](#) is defined thus:

1. In the case of an [array\\_view](#) that is ultimately derived from an array, the home location is the array.
2. In the case of an [array\\_view](#) that is ultimately derived from a host pointer, the home location is the original array view created using the pointer.

This means that two different [array\\_views](#) which have both been created, independently, on top of the same memory region are not guaranteed to appear coherent. In fact, creating and using top-level [array\\_views](#) on the same host storage is not supported. In order for such [array\\_view](#) to appear coherent, they must have a common top-level [array\\_view](#) ancestor which they both ultimately were derived from, and that top-level [array\\_view](#) must be the only one which is constructed on top of the memory it refers to.

This is illustrated in the next example:

```
#include <assert.h>
#include <amp.h>

using namespace concurrency;

void coherence_buggy()
{
    int storage[10];
    array_view<int> av1(10, &storage[0]);
    array_view<int> av2(10, &storage[0]); // error: av2 is top-level and aliases av1
    array_view<int> av3(5, &storage[5]); // error: av3 is top-level and aliases av1, av2

    parallel_for_each( extent<1>(1), [=] (index<1>) restrict(amp) { av3[2] = 15; });
    parallel_for_each( extent<1>(1), [=] (index<1>) restrict(amp) { av2[7] = 16; });
    parallel_for_each( extent<1>(1), [=] (index<1>) restrict(amp) { av1[7] = 17; });

    assert(av1[7] == av2[7]); // undefined results
    assert(av1[7] == av3[2]); // undefined results
}

void coherence_ok()
{
    int storage[10];
    array_view<int> av1(10, &storage[0]);
    array_view<int> av2(av1);           // OK
    array_view<int> av3(av1.section(5,5)); // OK

    parallel_for_each( extent<1>(1), [=] (index<1>) restrict(amp) { av3[2] = 15; });
    parallel_for_each( extent<1>(1), [=] (index<1>) restrict(amp) { av2[7] = 16; });
    parallel_for_each( extent<1>(1), [=] (index<1>) restrict(amp) { av1[7] = 17; });
}
```



```

3159     assert(av1[7] == av2[7]); // OK, never fails, both equal 17
3160     assert(av1[7] == av3[2]); // OK, never fails, both equal 17
3161 }
3162

```

3163 An implementation is not obligated to report such programmer's errors.

### 3164 8.3 Effects of copy and copy\_async operations

3165 Copy operations are offered on [array](#), [array\\_view](#) and [texture](#).

3166 Copy operations copy a source host buffer, [array](#), [array\\_view](#) or a [texture](#) to a destination object which can also be one of  
3167 these four varieties (except host buffer to host buffer, which is handled by [std::copy](#)). A *copy* operation will read all  
3168 elements of its source. It may read each element multiple times and it may read elements in any order. It may employ  
3169 memory load instructions that are either coarser or more granular than the width of the primitive data types in the  
3170 container, but it is guaranteed to never read a memory location which is strictly outside of the source container.

3171 Similarly, *copy* will overwrite each and every element in its output range. It may do so multiple times and in any order and  
3172 may coarsen or break apart individual store operations, but it is guaranteed to never write a memory location which is  
3173 strictly outside of the target container.

3174 A synchronous copy operation extends from the time the function is called until it has returned. During this time, any  
3175 source location may be read and any destination location may be written. An asynchronous copy extends from the time  
3176 [copy\\_async](#) is called until the time the [std::future](#) returned is signaled.

3177 As always, it is the programmer's responsibility not to call functions which could result in a race. For example, this program  
3178 is racy because the two copy operations are concurrent and [b](#) is written to by the first parallel activity while it is being  
3179 updated by the second parallel activity.

```

3186     array<int> a(100), b(100), c(100);
3187     parallel_invoke(
3188         [&] { copy(a,b); },
3189         [&] { copy(b,c); });
3190

```

### 3192 8.4 Effects of array\_view::synchronize, synchronize\_async and refresh functions

3193 An [array\\_view](#) may be constructed to wrap over a host side pointer. For such [array\\_views](#), it is generally forbidden to  
3194 access the underlying [array\\_view](#) storage directly, as long as the [array\\_view](#) exists. Access to the storage area is generally  
3195 accomplished indirectly through the [array\\_view](#). However, [array\\_view](#) offers mechanisms to synchronize and refresh its  
3196 contents, which do allow accessing the underlying memory directly. These mechanisms are described below.

3197 Reading of the underlying storage is possible under the condition that the view has been first *synchronized* back to its home  
3198 storage. This is performed using the [synchronize](#) or [synchronize\\_async](#) member functions of [array\\_view](#).

3201 When a top-level view is initially created on top of a raw buffer, it is synchronized with it. After it has been constructed, a  
3202 top-level view, as well as derived views, may lose coherence with the underlying host-side raw memory buffer if the  
3203 [array\\_view](#) is passed to [parallel\\_for\\_each](#) as a mutable view, or if the view is a target of a copy operation. In order to  
3204 restore coherence with host-side underlying memory [synchronize](#) or [synchronize\\_async](#) must be called. Synchronization is  
3205 restored when [synchronize](#) returns, or when the completion\_future returned by [synchronize\\_async](#) is ready.

3206 For the sake of composition with [parallel\\_for\\_each](#), [copy](#), and all other host-side operations involving a view, [synchronize](#)  
3207 should be considered a read of the entire data section referred to by the view, as if it was both the source of a copy

operation, and thus it must not be executed concurrently with any other operation involving writing the view. Note that even though `synchronize` does potentially modify the underlying host memory, it is logically a no-op as it doesn't affect the logical contents of the array. As such, it is allowed to execute concurrently with other operations which read the array view. As with `copy`, `synchronize` works at the granularity of the view it is applied to, e.g., synchronizing a view representing a subsection of a parent view doesn't necessarily synchronize the entire parent view. It is just guaranteed to synchronize the overlapping portions of such related views.

`array_views` are also required to synchronize their home storage:

1. Before they are destructed if and only if it is the last view of the underlying data container.
2. When they are accessed using the subscript operator (on said home location)

As a result of (1), any errors in synchronization which may be encountered during destruction of arrays views will not be propagated through the destructor. Users are therefore encouraged to ensure that `array_views` which may contain unsynchronized data are explicitly synchronized before they are destructed.

As a result of (2), the implementation of the subscript operator may need to contain a coherence enforcing check, especially on platforms where the accelerator hardware and host memory are not shared, and therefore coherence is managed explicitly by the C++ AMP runtime. Such a check may be detrimental for code desiring to achieve high performance through vectorization of the array view accesses. Therefore it is recommended for such performance-sensitive code to obtain a pointer to the beginning of a "run" and perform the low-level accesses needed based off of the raw pointer into the `array_view`. `array_views` are guaranteed to be contiguous in the unit-stride dimension, which enables this style of coding. Furthermore, the code may explicitly synchronize the `array_view` and at that point read the home storage directly, without the mediation of the view.

Sometimes it is desirable to also allow refreshing of a view by directly from its underlying memory. The `refresh` member function is provided for this task. This function revokes any caches associated with the view and resynchronizes the view's contents with the underlying memory. As such it may not be invoked concurrently with any other operation that accesses the view's data. However, it is safe to assume that `refresh` doesn't modify the view's underlying data and therefore concurrent read access to the underlying data is allowed during `refresh`'s operation and after `refresh` has returned, till the point when coherence may have been lost again, as has been described above in the discussion on the `synchronize` member function.

## 9 Math Functions

C++ AMP contains a rich library of floating point math functions that can be used in an accelerated computation. The C++ AMP library comes in two flavors, each contained in a separate namespace. The functions contained in the `concurrency::fast_math` namespace support only single-precision (`float`) operands and are optimized for performance at the expense of accuracy. The functions contained in the `concurrency::precise_math` namespace support both single and double precision (`double`) operands and are optimized for accuracy at the expense of performance. The two namespaces cannot be used together without introducing ambiguities. The accuracy of the functions in the `concurrency::precise_math` namespace shall be at least as high as those in the `concurrency::fast_math` namespace.

All functions are available in the `<amp_math.h>` header file, and all are decorated `restrict(amp)`.

### 9.1 fast\_math

Functions in the `fast_math` namespace are designed for computations where accuracy is not a prime requirement, and therefore the minimum precision is implementation-defined.

Not all functions available in `precise_math` are available in `fast_math`.



3259

C++ API function	Description
float acosf(float x) float acos(float x)	Returns the arc cosine in radians and the value is mathematically defined to be between 0 and PI (inclusive).
float asinf(float x) float asin(float x)	Returns the arc sine in radians and the value is mathematically defined to be between -PI/2 and PI/2 (inclusive).
float atanf(float x) float atan(float x)	Returns the arc tangent in radians and the value is mathematically defined to be between -PI/2 and PI/2 (inclusive).
float atan2f(float y, float x) float atan2(float y, float x)	Calculates the arc tangent of the two variables x and y. It is similar to calculating the arc tangent of $y / x$ , except that the signs of both arguments are used to determine the quadrant of the result.). Returns the result in radians, which is between -PI and PI (inclusive).
float ceilf(float x) float ceil(float x)	Rounds x up to the nearest integer.
float cosf(float x) float cos(float x)	Returns the cosine of x.
float coshf(float x) float cosh(float x)	Returns the hyperbolic cosine of x.
float expf(float x) float exp(float x)	Returns the value of e (the base of natural logarithms) raised to the power of x.
float exp2f(float x) float exp2(float x)	Returns the value of 2 raised to the power of x.
float fabsf(float x) float fabs(float x)	Returns the absolute value of floating-point number
float floorf(float x) float floor(float x)	Rounds x down to the nearest integer.
float fmaxf(float x, float y) float fmax(float x, float y)	Selects the greater of x and y.
float fminf(float x, float y) float fmin(float x, float y)	Selects the lesser of x and y.
float fmodf(float x, float y) float fmod(float x, float y)	Computes the remainder of dividing x by y. The return value is $x - n * y$ , where n is the quotient of $x / y$ , rounded towards zero to an integer.
float frexpf(float x, int * exp) float frexp(float x, int * exp)	Splits the number x into a normalized fraction and an exponent which is stored in exp.
int isfinite(float x)	Determines if x is finite.
int isinf(float x)	Determines if x is infinite.
int isnan(float x)	Determines if x is NAN.
float ldexpf(float x, float exp) float ldexp(float x, float exp)	Returns the result of multiplying the floating-point number x by 2 raised to the power exp
float logf(float x) float log(float x)	Returns the natural logarithm of x.
float log10f(float x) float log10(float x)	Returns the base 10 logarithm of x.
float log2f(float x) float log2(float x)	Returns the base 2 logarithm of x.
float modff(float x, float * iptr) float modf(float x, float * iptr)	Breaks the argument x into an integral part and a fractional part, each of which has the same sign as x. The integral part is stored in iptr.
float powf(float x, float y) float pow(float x, float y)	Returns the value of x raised to the power of y.

float roundf(float x) float round(float x)	Rounds x to the nearest integer.
float rsqrtf(float x) float rsqrt(float x)	Returns the reciprocal of the square root of x.
int signbit(float x) int signbit(double x)	Returns a non-zero value if the value of X has its sign bit set.
float sinf(float x) float sin(float x)	Returns the sine of x.
void sincosf(float x, float* s, float* c) void sincos(float x, float* s, float* c)	Returns the sine and cosine of x.
float sinhf(float x) float sinh(float x)	Returns the hyperbolic sine of x.
float sqrtf(float x) float sqrt(float x)	Returns the non-negative square root of x
float tanf(float x) float tan(float x)	Returns the tangent of x.
float tanhf(float x) float tanh(float x)	Returns the hyperbolic tangent of x.
float truncf(float x) float trunc(float x)	Rounds x to the nearest integer not larger in absolute value.

The following list of standard math functions from the “std::” namespace shall be imported into the concurrency::fast\_math namespace:

```

using std::acosf;
using std::asinf;
using std::atanf;
using std::atan2f;
using std::ceilf;
using std::cosf;
using std::coshf;
using std::expf;
using std::fabsf;
using std::floorf;
using std::fmodf;
using std::frexpf;
using std::ldexpf;
using std::logf;
using std::log10f;
using std::modff;
using std::powf;
using std::sinf;
using std::sinhf;
using std::sqrtf;
using std::tanf;
using std::tanhf;

using std::acos;
using std::asin;
using std::atan;
using std::atan2;
using std::ceil;
using std::cos;
using std::cosh;
using std::exp;

```

```

3295     using std::fabs;
3296     using std::floor;
3297     using std::fmod;
3298     using std::frexp;
3299     using std::ldexp;
3300     using std::log;
3301     using std::log10;
3302     using std::modf;
3303     using std::pow;
3304     using std::sin;
3305     using std::sinh;
3306     using std::sqrt;
3307     using std::tan;
3308     using std::tanh;

```

Importing these names into the `fast_math` namespace enables each of them to be called in unqualified syntax from a function that has both “`restrict(cpu,amp)`” restrictions. E.g.,

```

3313 void compute() restrict(cpu,amp) {
3314     ...
3315     float x = cos(y); // resolves to std::cos in “cpu” context; else fast_math::cos in “amp” context
3316     ...
3317 }

```

## 9.2 precise\_math

Functions in the `precise_math` namespace are designed for computations where accuracy is required. In the table below, the precision of each function is stated in units of “ulps” (error in last position).

Functions in the `precise_math` namespace also support both single and double precision, and are therefore dependent upon double-precision support in the underlying hardware, even for single-precision variants.

C++ API function	Description	Precision (float)	Precision (double)
float acosf(float x) float acos(float x) double acos(double x)	Returns the arc cosine in radians and the value is mathematically defined to be between 0 and PI (inclusive).	3	2
float acoshf(float x) float acosh(float x) double acosh(double x)	Returns the hyperbolic arccosine.	4	2
float asinf(float x) float asin(float x) double asin(double x)	Returns the arc sine in radians and the value is mathematically defined to be between -PI/2 and PI/2 (inclusive).	4	2
float asinhf(float x) float asinh(float x) double asinh(double x)	Returns the hyperbolic arcsine.	3	2
float atanf(float x) float atan(float x) double atan(double x)	Returns the arc tangent in radians and the value is mathematically defined to be between -PI/2 and PI/2 (inclusive).	2	2
float atanhf(float x) float atanh(float x)	Returns the hyperbolic arctangent.	3	2

double atanh(float x)			
float atan2f(float y, float x) float atan2(float y, float x) double atan2(double y, double x)	Calculates the arc tangent of the two variables x and y. It is similar to calculating the arc tangent of y / x, except that the signs of both arguments are used to determine the quadrant of the result.). Returns the result in radians, which is between -PI and PI (inclusive).	3	2
float cbrtf(float x) float cbrt(float x) double cbrt(double x)	Returns the (real) cube root of x.	1	1
float ceilf(float x) float ceil(float x) double ceil(double x)	Rounds x up to the nearest integer.	0	0
float copysignf(float x, float y) float copysign(float x, float y) double copysign(double x, double y)	Return a value whose absolute value matches that of x, but whose sign matches that of y. If x is a NaN, then a NaN with the sign of y is returned.	N/A	N/A
float cosf(float x) float cos(float x) double cos(double x)	Returns the cosine of x.	2	2
float coshf(float x) float cosh(float x) double cosh(double x)	Returns the hyperbolic cosine of x.	2	2
float cospif(float x) float cospi(float x) double cospi(double x)	Returns the cosine of pi * x.	2	2
float erff(float x) float erf(float x) double erf(double x)	Returns the error function of x; defined as $\text{erf}(x) = 2/\sqrt{\pi} \cdot \int_0^x \exp(-t^2) dt$	3	2
float erfcf(float x) float erfc(float x) double erfc(double x)	Returns the complementary error function of x that is $1.0 - \text{erf}(x)$ .	6	5
float erfinvf(float x) float erfinv(float x) double erfinv(double x)	Returns the inverse error function.	3	8
float erfcinvf(float x) float erfcinv(float x) double erfcinv(double x)	Returns the inverse of the complementary error function.	7	8
float expf(float x) float exp(float x) double exp(double x)	Returns the value of e (the base of natural logarithms) raised to the power of x.	2	1
float exp2f(float x) float exp2(float x) double exp2(double x)	Returns the value of 2 raised to the power of x.	2	1
float exp10f(float x) float exp10(float x) double exp10(double x)	Returns the value of 10 raised to the power of x.	2	1

float expm1f(float x)  float expm1(float x) double expm1(double x)	Returns a value equivalent to 'exp (x) - 1'	1	1
float fabsf(float x)  float fabs(float x) double fabs(double x)	Returns the absolute value of floating-point number	N/A	N/A
float fdimf(float x, float y)  float fdim(float x, float y) double fdim(double x, double y)	These functions return max(x-y,0). If x or y or both are NaN, Nan is returned.	0	0
float floorf(float x)  float floor(float x) double floor(double x)	Rounds x down to the nearest integer.	0	0
float fmaf(float x, float y, float z)  float fma(float x, float y, float z) double fma(double x, double y, double z)	Computes (x * y) + z, rounded as one ternary operation: they compute the value (as if) to infinite precision and round once to the result format, according to the current rounding mode. A range error may occur.	0	0 <sup>6</sup>
float fmaxf(float x, float y)  float fmax(float x, float y) double fmax(double x, double y)	Selects the greater of x and y.	N/A	N/A
float fminf(float x, float y)  float fmin(float x, float y) double fmin(double x, double y)	Selects the lesser of x and y.	N/A	N/A
float fmodf(float x, float y)  float fmod(float x, float y) double fmod(double x, double y)	Computes the remainder of dividing x by y. The return value is x - n * y, where n is the quotient of x / y, rounded towards zero to an integer.	0	0
int fpclassify(float x);  int fpclassify(double x);	Floating point numbers can have special values, such as infinite or NaN. With the macro fpclassify(x) you can find out what type x is. The function takes any floating-point expression as argument. The result is one of the following values: <ul style="list-style-type: none"> <li>• FP_NAN : x is "Not a Number".</li> <li>• FP_INFINITE: x is either plus or minus infinity.</li> <li>• FP_ZERO: x is zero.</li> <li>• FP_SUBNORMAL : x is too small to be represented in normalized format.</li> <li>• FP_NORMAL : if nothing of the above is correct then it must be a normal floating-point number.</li> </ul>	N/A	N/A
float frexpf(float x, int * exp)  float frexp(float x, int * exp) double frexp(double x, int * exp)	Splits the number x into a normalized fraction and an exponent which is stored in exp.	0	0
float hypotf(float x, float y)  float hypot(float x, float y) double hypot(double x, double y)	Returns sqrt(x*x+y*y). This is the length of the hypotenuse of a right-angle triangle with sides of length x and y, or the distance of the point (x,y) from the origin.	3	2
int ilogbf (float x)  int ilogb(float x)	Return the exponent part of their argument as a signed integer. When no error occurs, these functions are equivalent to the corresponding logb() functions, cast to (int). An error will occur	0	0

---

<sup>6</sup> IEEE-754 round to nearest even.

int ilogb(double x)	for zero and infinity and NaN, and possibly for overflow.		
int isfinite(float x)	Determines if x is finite.	N/A	N/A
int isfinite(double x)			
int isinf(float x)	Determines if x is infinite.	N/A	N/A
int isinf(double x)			
int isnan(float x)	Determines if x is NaN.	N/A	N/A
int isnan(double x)			
int isnormal(float x)	Determines if x is normal.	N/A	N/A
int isnormal(double x)			
float ldexp(float x, float exp)	Returns the result of multiplying the floating-point number x by 2 raised to the power exp	0	0
float ldexp(float x, float exp)			
double ldexp(double x, double exp)			
float lgammaf(float x)	Computes the natural logarithm of the absolute value of gamma of x. A range error occurs if x is too large. A range error may occur if x is a negative integer or zero.	6 <sup>7</sup>	4 <sup>8</sup>
float lgamma(float x)			
double lgamma(double x)			
float logf(float x)	Returns the natural logarithm of x.	1	1
float log(float x)			
double log(double x)			
float log10f(float x)	Returns the base 10 logarithm of x.	3	1
float log10(float x)			
double log10(double x)			
float log2f(float x)	Returns the base 2 logarithm of x.	3	1
float log2(float x)			
double log2(double x)			
float log1pf(float x)	Returns a value equivalent to 'log (1 + x)'. It is computed in a way that is accurate even if the value of x is near zero.	2	1
float log1p(float x)			
double log1p(double x)			
float logbf(float x)	These functions extract the exponent of x and return it as a floating-point value. If FLT_RADIX is two, logb(x) is equal to floor(log2(x)), except it's probably faster.	0	0
float logb(float x)			
double logb(double x)	If x is de-normalized, logb() returns the exponent x would have if it were normalized.		
float modff(float x, float * iptr)	Breaks the argument x into an integral part and a fractional part, each of which has the same sign as x. The integral part is stored in iptr.	0	0
float modf(float x, float * iptr)			
double modf(double x, double * iptr)			
float nanf(int tagp)	return a representation (determined by tagp) of a quiet NaN. If the implementation does not support quiet NaNs, these functions return zero.	N/A	N/A
float nanf(int tagp)			
double nan(int tagp)			

<sup>7</sup> Outside interval -10.001 ... -2.264; larger inside.

<sup>8</sup> Outside interval -10.001 ... -2.264; larger inside.

float nearbyintf(float x) float nearbyint(float x) double nearbyint(double x)	Rounds the argument to an integer value in floating point format, using the current rounding direction	0	
float nextafterf(float x, float y) float nextafter(float x, float y) double nextafter(double x, double y)	Returns the next representable neighbor of x in the direction towards y. The size of the step between x and the result depends on the type of the result. If x = y the function simply returns y. If either value is NaN, then NaN is returned. Otherwise a value corresponding to the value of the least significant bit in the mantissa is added or subtracted, depending on the direction.	N/A	N/A
float powf(float x, float y) float pow(float x, float y) double pow(double x, double y)	Returns the value of x raised to the power of y.	8	2
float rcbrtf(float x) float rcbrt(float x) double rcbrt(double x)	Calculates reciprocal of the (real) cube root of x	2	1
float remainderf(float x, float y) float remainder(float x, float y) double remainder(double x, double y)	Computes the remainder of dividing x by y. The return value is x - n * y, where n is the value x / y, rounded to the nearest integer. If this quotient is 1/2 (mod 1), it is rounded to the nearest even number (independent of the current rounding mode). If the return value is 0, it has the sign of x.	0	0
float remquo(float x, float y, int * quo) float remquo(float x, float y, int * quo) double remquo(double x, double y, int * quo)	Computes the remainder and part of the quotient upon division of x by y. A few bits of the quotient are stored via the quo pointer. The remainder is returned.	0	0
float roundf(float x) float round(float x) double round(double x)	Rounds x to the nearest integer.	0	0
float rsqrtf(float x) float rsqrt(float x) double rsqrt(double x)	Returns the reciprocal of the square root of x.	2	1
float sinpif(float x) float sinpi(float x) double sinpi(double x)	Returns the sine of pi * x.	2	2
float scalbf(float x, float exp) float scalb(float x, float exp) double scalb(double x, double exp)	Multiplies their first argument x by FLT_RADIX (probably 2) to the power exp.	0	0
float scalbnf(float x, int exp) float scalbn(float x, int exp) double scalbn(double x, int exp)	Multiplies their first argument x by FLT_RADIX (probably 2) to the power exp. If FLT_RADIX equals 2, then scalbn() is equivalent to ldexp(). The value of FLT_RADIX is found in <float.h>.	0	0
int signbit(float x) int signbit(double x)	Returns a non-zero value if the value of X has its sign bit set.	N/A	N/A
float sinf(float x) float sin(float x) double sin(double x)	Returns the sine of x.	2	2
void sincosf(float x, float * s, float * c) void sincos(float x, float * s, float * c) void sincos(double x, double * s, double * c)	Returns the sine and cosine of x.	2	2
float sinh(float x)	Returns the hyperbolic sine of x.	3	2

float sinh(float x) double sinh(double x)			
float sqrtf(float x)  float sqrt(float x) double sqrt(double x)	Returns the non-negative square root of x	0	0 <sup>9</sup>
float tgammaf(float x)  float tgamma(float x) double tgamma(double x)	This function returns the value of the Gamma function for the argument x.	11	8
float tanf(float x)  float tan(float x) double tan(double x)	Returns the tangent of x.	4	2
float tanhf(float x)  float tanh(float x) double tanh(double x)	Returns the hyperbolic tangent of x.	2	2
float tanpif(float x)  float tanpi(float x) double tanpi(double x)	Returns the tangent of pi * x.	2	2
float truncf(float x)  float trunc(float x) double trunc(double x)	Rounds x to the nearest integer not larger in absolute value.	0	0

3325

3326

The following list of standard math functions from the “std::” namespace shall be imported into the concurrency::precise  
\_math namespace:

3327

3328

3329

```
using std::acosf;
```

3330

```
using std::asinf;
```

3331

```
using std::atanf;
```

3332

```
using std::atan2f;
```

3333

```
using std::ceilf;
```

3334

```
using std::cosf;
```

3335

```
using std::coshf;
```

3336

```
using std::expf;
```

3337

```
using std::fabsf;
```

3338

```
using std::floorf;
```

3339

```
using std::fmodf;
```

3340

```
using std::frexpf;
```

3341

```
using std::ldexpf;
```

3342

```
using std::logf;
```

3343

```
using std::log10f;
```

3344

```
using std::modff;
```

3345

```
using std::powf;
```

3346

```
using std::sinf;
```

3347

```
using std::sinhf;
```

3348

```
using std::sqrtf;
```

3349

```
using std::tanf;
```

3350

```
using std::tanhf;
```

3351

---

<sup>9</sup> IEEE-754 round to nearest even.



```

3352     using std::acos;
3353     using std::asin;
3354     using std::atan;
3355     using std::atan2;
3356     using std::ceil;
3357     using std::cos;
3358     using std::cosh;
3359     using std::exp;
3360     using std::fabs;
3361     using std::floor;
3362     using std::fmod;
3363     using std::frexp;
3364     using std::ldexp;
3365     using std::log;
3366     using std::log10;
3367     using std::modf;
3368     using std::pow;
3369     using std::sin;
3370     using std::sinh;
3371     using std::sqrt;
3372     using std::tan;
3373     using std::tanh;
3374

```

Importing these names into the `precise_math` namespace enables each of them to be called in unqualified syntax from a function that has both “`restrict(cpu,amp)`” restrictions. E.g.,

```

3377
3378 void compute() restrict(cpu,amp) {
3379     ...
3380     float x = cos(y); // resolves to std::cos in “cpu” context; else fast_math::cos in “amp” context
3381     ...
3382 }
3383

```

## 10 Graphics (Optional)

Programming model elements defined in `<amp_graphics.h>` and `<amp_short_vectors.h>` are designed for graphics programming in conjunction with accelerated compute on an accelerator device, and are therefore appropriate only for proper GPU accelerators. Accelerator devices that do not support native graphics functionality need not implement these features.

All types in this section are defined in the `concurrency::graphics` namespace.

### 10.1 texture<T,N>

The `texture` class provides the means to create textures from raw memory or from file. `textures` are similar to `arrays` in that they are containers of data and they behave like STL containers with respect to assignment and copy construction.

`textures` are templated on `T`, the element type, and on `N`, the rank of the texture. `N` can be one of 1, 2 or 3.

The element type of the `texture`, also referred to as the texture’s logical element type, is one of a closed set of short vector types defined in the `concurrency::graphics` namespace and covered elsewhere in this specification. The below table briefly enumerates all supported element types.

Rank of element type, (also	Signed Integer	Unsigned Integer	Single precision floating point number	Single precision	Single precision	Double precision
-----------------------------	----------------	------------------	--	------------------	------------------	------------------

referred to as "number of scalar elements")				singed normalized number	unsigned normalized number	floating point number
1	int	unsigned int	float	norm	unorm	double
2	int_2	uint_2	float_2	norm_2	unorm_2	double_2
3	int_3	uint_3	float_3	norm_3	unorm_3	double_3
4	int_4	uint_4	float_4	norm_4	unorm_4	double_4

Remarks:

1. *norm* and *unorm* vector types are vector of *floats* which are normalized to the range [-1..1] and [0...1], respectively.
2. Grayed-out cells represent vector types which are defined by C++ AMP but which are not necessarily supported as *texture* value types. Implementations can optionally support the types in the grayed-out cells in the above table.

*Microsoft-specific: grayed-out cells in the above table are not supported.*

### 10.1.1 Synopsis

```

template <typename T, int N>
class texture
{
public:
    static const int rank = _Rank;
    typedef typename T value_type;
    typedef short_vectors_traits<T>::scalar_type scalar_type;

    texture(const extent<N>& _Ext);

    texture(int _E0);
    texture(int _E0, int _E1);
    texture(int _E0, int _E1, int _E2);

    texture(const extent<N>& _Ext, const accelerator_view& _Acc_view);

    texture(int _E0, const accelerator_view& _Acc_view);
    texture(int _E0, int _E1, const accelerator_view& _Acc_view);
    texture(int _E0, int _E1, int _E2, const accelerator_view& _Acc_view);

    texture(const extent<N>& _Ext, unsigned int _Bits_per_scalar_element);

    texture(int _E0, unsigned int _Bits_per_scalar_element);
    texture(int _E0, int _E1, unsigned int _Bits_per_scalar_element);
    texture(int _E0, int _E1, int _E2, unsigned int _Bits_per_scalar_element);

    texture(const extent<N>& _Ext, unsigned int _Bits_per_scalar_element,
            const accelerator_view& _Acc_view);

    texture(int _E0, unsigned int _Bits_per_scalar_element, const accelerator_view&
    _Acc_view);
    texture(int _E0, int _E1, unsigned int _Bits_per_scalar_element,
            const accelerator_view& _Acc_view);
    texture(int _E0, int _E1, int _E2, unsigned int _Bits_per_scalar_element,
            const accelerator_view& _Acc_view);

    template <typename TInputIterator>
    texture(const extent<N>&, TInputIterator _Src_first, TInputIterator _Src_last);

```

```

3450 template <typename TInputIterator>
3451     texture(int _E0, TInputIterator _Src_first, TInputIterator _Src_last);
3452 template <typename TInputIterator>
3453     texture(int _E0, int _E1, TInputIterator _Src_first, TInputIterator _Src_last);
3454 template <typename TInputIterator>
3455     texture(int _E0, int _E1, int _E2, TInputIterator _Src_first,
3456             TInputIterator _Src_last);
3457
3458 template <typename TInputIterator>
3459     texture(const extent<N>&, TInputIterator _Src_first, TInputIterator _Src_last,
3460             const accelerator_view& _Acc_view);
3461
3462 template <typename TInputIterator>
3463     texture(int _E0, TInputIterator _Src_first, TInputIterator _Src_last,
3464             const accelerator_view& _Acc_view);
3465 template <typename TInputIterator>
3466     texture(int _E0, int _E1, TInputIterator _Src_first, TInputIterator _Src_last,
3467             const accelerator_view& _Acc_view);
3468 texture(int _E0, int _E1, int _E2, TInputIterator _Src_first, TInputIterator _Src_last,
3469 const accelerator_view& _Acc_view);
3470
3471 texture(const extent<N>&, const void * _Source, unsigned int _Src_byte_size,
3472         unsigned int _Bits_per_scalar_element);
3473
3474 texture(int _E0, const void * _Source, unsigned int _Src_byte_size,
3475         unsigned int _Bits_per_scalar_element);
3476 texture(int _E0, int _E1, const void * _Source, unsigned int _Src_byte_size,
3477         unsigned int _Bits_per_scalar_element);
3478 texture(int _E0, int _E1, int _E2, const void * _Source,
3479         unsigned int _Src_byte_size, unsigned int _Bits_per_scalar_element);
3480
3481 texture(const extent<N>&, const void * _Source, unsigned int _Src_byte_size,
3482         unsigned int _Bits_per_scalar_element, const accelerator_view& _Acc_view);
3483
3484 texture(int _E0, const void * _Source, unsigned int _Src_byte_size,
3485         unsigned int _Bits_per_scalar_element, const accelerator_view& _Acc_view);
3486 texture(int _E0, int _E1, const void * _Source, unsigned int _Src_byte_size,
3487         unsigned int _Bits_per_scalar_element, const accelerator_view& _Acc_view);
3488 texture(int _E0, int _E1, int _E2, const void * _Source, unsigned int _Src_byte_size,
3489         unsigned int _Bits_per_scalar_element, const accelerator_view& _Acc_view);
3490
3491 texture(const texture& _Src);
3492 texture(const texture& _Src, const accelerator_view& _Acc_view);
3493 texture& operator=(const texture& _Src);
3494
3495 texture(texture&& _Other);
3496 texture& operator=(texture&& _Other);
3497
3498 void copy_to(texture& _Dest) const;
3499 void copy_to(const writeonly_texture_view<T,N>& _Dest) const;
3500
3501 unsigned int get_Bits_per_scalar_element() const;
3502 __declspec(property(get= get_Bits_per_scalar_element)) int bits_per_scalar_element;
3503
3504 unsigned int get_data_length() const;
3505 __declspec(property(get=get_data_length)) unsigned int data_length;
3506
3507 extent<N> get_extent() const restrict(cpu,amp);
3508 __declspec(property(get=get_extent)) extent<N> extent;
3509
3510 accelerator_view get_accelerator_view() const;
3511 __declspec(property(get=get_accelerator_view)) accelerator_view accelerator_view;

```

```

3513
3514     const value_type operator[] (const index<N>& _Index) const restrict(amp);
3515     const value_type operator[] (int _I0) const restrict(amp);
3516     const value_type operator() (const index<N>& _Index) const restrict(amp);
3517     const value_type operator() (int _I0) const restrict(amp);
3518     const value_type operator() (int _I0, int _I1) const restrict(amp);
3519     const value_type operator() (int _I0, int _I1, int _I2) const restrict(amp);
3520     const value_type get(const index<N>& _Index) const restrict(amp);
3521
3522     void set(const index<N>& _Index, const value_type& _Val) restrict(amp);
3523 };
3524

```

## 3525 10.1.2 Introduced typedefs

```
typedef ... value_type;
```

The logical value type of the texture. e.g., for texture<float2, 3>, value\_type would be float2.

3526

```
typedef ... scalar_type;
```

The scalar type that serves as the component of the texture's value type. For example, for texture<int2, 3>, the scalar type would be "int".

## 3527 10.1.3 Constructing an uninitialized texture

3528

```

texture(const extent<N>& _Ext);

texture(int _E0);
texture(int _E0, int _E1);
texture(int _E0, int _E1, int _E2);

texture(const extent<N>& _Ext, const accelerator_view& _Acc_view);

texture(int _E0, const accelerator_view& _Acc_view);
texture(int _E0, int _E1, const accelerator_view& _Acc_view);
texture(int _E0, int _E1, int _E2, const accelerator_view& _Acc_view);

texture(const extent<N>& _Ext, unsigned int _Bits_per_scalar_element);

texture(int _E0, unsigned int _Bits_per_scalar_element);
texture(int _E0, int _E1, unsigned int _Bits_per_scalar_element);
texture(int _E0, int _E1, int _E2, unsigned int _Bits_per_scalar_element);

texture(const extent<N>& _Ext, unsigned int _Bits_per_scalar_element, const accelerator_view&
_Acc_view);

texture(int _E0, unsigned int _Bits_per_scalar_element, const accelerator_view& _Acc_view);
texture(int _E0, int _E1, unsigned int _Bits_per_scalar_element, const accelerator_view&
_Acc_view);
texture(int _E0, int _E1, int _E2, unsigned int _Bits_per_scalar_element, const
accelerator_view& _Acc_view);

```

Creates an uninitialized texture with the specified shape, number of bits per scalar element, on the specified accelerator view.

### Parameters:

_Ext	Extents of the texture to create
_E0	Extent of dimension 0
_E1	Extent of dimension 1
_E2	Extent of dimension 2
_Bits_per_scalar_element	Number of bits per each scalar element in the underlying scalar type of the texture.

_Acc_view	Accelerator view where to create the texture
<b>Error condition</b>	<b>Exception thrown</b>
Out of memory	concurrency::runtime_exception
Invalid number of bits per scalar elementspecified	concurrency::runtime_exception
Invalid combination of value_type and bits per scalar element	concurrency::unsupported_feature
accelerator_view doesn't support textures	concurrency::unsupported_feature

The table below summarizes all valid combinations of underlying scalar types (columns), ranks(rows), supported values for bits-per-scalar-element (inside the table cells), and default value of bits-per-scalar-element for each given combination (highlighted in green). Note that unorm and norm have no default value for bits-per-scalar-element. Implementations can optionally support textures of double4, with implementation-specific values of bits-per-scalar-element.

*Microsoft-specific: the current implementation doesn't support textures of double4.*

Rank	int	uint	float	norm	unorm	double
1	8, 16, 32	8, 16, 32	16, 32	8, 16	8, 16	64
2	8, 16, 32	8, 16, 32	16, 32	8, 16	8, 16	64
4	8, 16, 32	8, 16, 32	16, 32	8, 16	8, 16	

#### 10.1.4 Constructing a texture from a host side iterator

```

template <typename TInputIterator>
texture(const extent<N>& _Ext, TInputIterator _Src_first, TInputIterator _Src_last);
texture(int _E0, TInputIterator _Src_first, TInputIterator _Src_last);
texture(int _E0, int _E1, TInputIterator _Src_first, TInputIterator _Src_last);
texture(int _E0, int _E1, int _E2, TInputIterator _Src_first, TInputIterator _Src_last);

template <typename TInputIterator>
texture(const extent<N>&, TInputIterator _Src_first, TInputIterator _Src_last, const
accelerator_view& _Acc_view);

template <typename TInputIterator>
texture(const extent<N>& _Ext, TInputIterator _Src_first, TInputIterator _Src_last, const
accelerator_view& _Acc_view);
texture(int _E0, TInputIterator _Src_first, TInputIterator _Src_last, const accelerator_view&
_Acc_view);
texture(int _E0, int _E1, TInputIterator _Src_first, TInputIterator _Src_last, const
accelerator_view& _Acc_view);
texture(int _E0, int _E1, int _E2, TInputIterator _Src_first, TInputIterator _Src_last, const
accelerator_view& _Acc_view);

```

Creates a texture from a host-side iterator. The data type of the iterator must be the same as the value type of the texture. Textures with element types based on norm or unorm do not support this constructor (usage of it will result in a compile-time error).

##### Parameters:

_Ext	Extents of the texture to create
_E0	Extent of dimension 0

_E1	Extent of dimension 1
_E2	Extent of dimension 2
_Src_first	Iterator pointing to the first element to be copied into the texture
_Src_last	Iterator pointing immediately past the last element to be copied into the texture
_Acc_view	Accelerator view where to create the texture
<b>Error condition</b>	<b>Exception thrown</b>
Out of memory	concurrency::runtime_exception
Inadequate amount of data supplied through the iterators	concurrency::runtime_exception
Accelerator_view doesn't support textures	concurrency::unsupported_feature

3540

3541 **10.1.5 Constructing a texture from a host-side data source**

3542

```

texture(const extent<N>&, const void * _Source, unsigned int _Src_byte_size, unsigned int
_Bits_per_scalar_element);

texture(int _E0, const void * _Source, unsigned int _Src_byte_size, unsigned int
_Bits_per_scalar_element);
texture(int _E0, int _E1, const void * _Source, unsigned int _Src_byte_size, unsigned int
_Bits_per_scalar_element);
texture(int _E0, int _E1, int _E2, const void * _Source, unsigned int _Src_byte_size, unsigned
int _Bits_per_scalar_element);

texture(const extent<N>&, const void * _Source, unsigned int _Src_byte_size, unsigned int
_Bits_per_scalar_element, const accelerator_view& _Acc_view);

texture(int _E0, const void * _Source, unsigned int _Src_byte_size, unsigned int
_Bits_per_scalar_element, const accelerator_view& _Acc_view);
texture(int _E0, int _E1, const void * _Source, unsigned int _Src_byte_size, unsigned int
_Bits_per_scalar_element, const accelerator_view& _Acc_view);
texture(int _E0, int _E1, int _E2, const void * _Source, unsigned int _Src_byte_size, unsigned
int _Bits_per_scalar_element, const accelerator_view& _Acc_view);

```

Creates a texture from a host-side provided buffer. The format of the data source must be compatible with the texture's vector type, and the amount of data in the data source must be exactly the amount necessary to initialize a texture in the specified format, with the given number of bits per scalar element.

For example, a 2D texture of uint2 initialized with the extent of 100x200 and with \_Bits\_per\_scalar\_element equal to 8 will require a total of  $100 * 200 * 2 * 8 = 320,000$  bits available to copy from \_Source, which is equal to 40,000 bytes. (or in other words, one byte, per one scalar element, for each scalar element, and each pixel, in the texture).

**Parameters:**

_Ext	Extents of the texture to create
_E0	Extent of dimension 0
_E1	Extent of dimension 1
_E2	Extent of dimension 2
_Source	Pointer to a host buffer
_Src_byte_size	Number of bytes of the host source buffer
_Bits_per_scalar_element	Number of bits per each scalar element in the underlying scalar type of the texture.
_Acc_view	Accelerator view where to create the texture
<b>Error condition</b>	<b>Exception thrown</b>

Out of memory	concurrency::runtime_exception
Inadequate amount of data supplied through the host buffer (_Src_byte_size < texture.data_length)	concurrency::runtime_exception
Invalid number of bits per scalar elements specified	concurrency::runtime_exception
Invalid combination of value_type and bits per scalar element	concurrency::unsupported_feature
Accelerator_view doesn't support textures	concurrency::unsupported_feature

3543

3544 

### 10.1.6 Constructing a texture by cloning another

3545

```
texture(const texture& _Src);
```

Initializes one texture from another. The texture is created on the same accelerator view as the source.

**Parameters:**

_Src	Source texture or texture_view to copy from
<b>Error condition</b>	<b>Exception thrown</b>
Out of memory	concurrency::runtime_exception

3546

```
texture(const texture& _Src, const accelerator_view& _Acc_view);
```

Initializes one texture from another.

**Parameters:**

_Src	Source texture or texture_view to copy from
_Acc_view	Accelerator view where to create the texture
<b>Error condition</b>	<b>Exception thrown</b>
Out of memory	concurrency::runtime_exception
Accelerator_view doesn't support textures	concurrency::unsupported_feature

3547

3548 

### 10.1.7 Assignment operator

3549

```
texture& operator=(const texture& _Src);
```

Release the resource of this texture, allocate the resource according to \_Src's properties, then deep copy \_Src's content to this texture.

**Parameters:**

_Src	Source texture or texture_view to copy from
<b>Error condition</b>	<b>Exception thrown</b>
Out of memory	concurrency::runtime_exception

3550

3551 

### 10.1.8 Copying textures

```
void copy_to(texture& _Dest) const;
void copy_to(const writeonly_texture_view<T,N>& _Dest) const;
```

Copies the contents of one texture onto the other. The textures must have been created with exactly the same extent and with compatible physical formats; that is, the number of scalar elements and the number of bits per scalar elements must agree. The textures could be from different accelerators.

**Parameters:**

<code>_Dest</code>	Destination texture or <code>writable_texture_view</code> to copy to
<b>Error condition</b>	<b>Exception thrown</b>
Out of memory	<code>concurrency::runtime_exception</code>
Incompatible texture formats	<code>concurrency::runtime_exception</code>
Extents don't match	<code>concurrency::runtime_exception</code>

3552

### 10.1.9 Moving textures

```
texture(texture&& _Other);
texture& operator=(texture&& _Other);
```

“Moves” (in the C++ R-value reference sense) the contents of `_Other` to “this”. The source and destination textures do not have to be necessarily on the same accelerator originally.

As is typical in C++ move constructors, no actual copying or data movement occurs; simply one C++ texture object is vacated of its internal representation, which is moved to the target C++ texture object.

**Parameters:**

<code>_Other</code>	Object whose contents are moved to “this”
<b>Error condition</b>	<b>Exception thrown</b>
None	

3555

### 10.1.10 Querying texture's physical characteristics

```
unsigned int get_Bits_per_scalar_element() const;
__declspec(property(get=get_Bits_per_scalar_element)) unsigned int bits_per_scalar_element;
```

Gets the bits-per-scalar-element of the texture. Returns 0, if the texture is created using Direct3D Interop (10.1.15).

**Error conditions: none**

3557

3558

```
unsigned int get_data_length() const;
__declspec(property(get=get_data_length)) unsigned int data_length;
```

Gets the physical data length (in bytes) that is required in order to represent the texture on the host side with its native format.

**Error conditions: none**

3559

### 10.1.11 Querying texture's logical dimensions

```
extent<N> get_extent() const restrict(cpu,amp);
__declspec(property(get=get_extent)) extent<N> extent;
```

These members have the same meaning as the equivalent ones on the array class

**Error conditions: none**

3561

### 10.1.12 Querying the accelerator\_view where the texture resides

```
accelerator_view get_accelerator_view() const;
__declspec(property(get=get_accelerator_view)) accelerator_view accelerator_view;
```

3562

3563



Retrieves the `accelerator_view` where the texture resides

**Error conditions:** none

### 10.1.13 Reading and writing textures

This is the core function of class `texture` on the accelerator. Unlike `arrays`, the entire value type has to be get/set, and is returned or accepted wholly. `textures` do not support returning a reference to their data internal representation.

Due to platform restrictions, only a limited number of `texture` types support simultaneous reading and writing. Reading is supported on all `texture` types, but writing through a `texture&` is only supported for `textures` of `int`, `uint`, and `float`, and even in those cases, the number of bits used in the physical format must be 32. In case a lower number of bits is used (8 or 16) and a kernel is invoked which contains code that could possibly both write into and read from one of these rank-1 `texture` types, then an implementation is permitted to raise a runtime exception.

*Microsoft-specific: the Microsoft implementation always raises a runtime exception in such a situation.*

Trying to call “set” on a `texture&` of a different element type (i.e., on other than `int`, `uint`, and `float`) results in a static assert. In order to write into `textures` of other value types, the developer must go through a `writeonly_texture_view<T,N>`.

```
const value_type operator[] (const index<N>& _Index) const restrict(amp);
const value_type operator[] (int _I0) const restrict(amp);
const value_type operator() (const index<N>& _Index) const restrict(amp);
const value_type operator() (int _I0) const restrict(amp);
const value_type operator() (int _I0, int _I1) const restrict(amp);
const value_type operator() (int _I0, int _I1, int _I2) const restrict(amp);
const value_type get(const index<N>& _Index) const restrict(amp);
void set(const index<N>& _Index, const value_type& _Value) const restrict(amp);
```

Loads one texel out of the texture. In case the overload where an integer tuple is used, if an overload which doesn't agree with the rank of the matrix is used, then a `static_assert` ensues and the program fails to compile.

In the texture is indexed, at runtime, outside of its logical bounds, behavior is undefined.

Parameters

<code>_Index</code>	An N-dimension logical integer coordinate to read from
<code>_I0, _I1, _I0</code>	Index components, equivalent to providing <code>index&lt;1&gt;(_I0)</code> , or <code>index&lt;2&gt;(_I0,_I1)</code> or <code>index&lt;2&gt;(_I0,_I1,_I2)</code> . The arity of the function used must agree with the rank of the matrix. e.g., the overload which takes <code>(_I0,_I1)</code> is only available on textures of rank 2.
<code>_Value</code>	Value to write into the texture

**Error conditions:** if `set` is called on texture types which are not supported, a `static_assert` ensues.

### 10.1.14 Global texture copy functions

```
template <typename T, int N>
void copy(const texture<T,N>& _Texture, void * _Dst, unsigned int _Dst_byte_size);
```

Copies raw texture data to a host-side buffer. The buffer must be laid out in accordance with the texture format and dimensions.

Parameters

<code>_Texture</code>	Source texture or <code>texture_view</code>
<code>_Dst</code>	Pointer to destination buffer on the host
<code>_Dst_byte_size</code>	Number of bytes in the destination buffer
<b>Error condition</b>	<b>Exception thrown</b>

Out of memory (*)	
Buffer too small	

(\*) Out of memory errors may occur due to the need to allocate temporary buffers in some memory transfer scenarios.

```
template <typename T, int N>
void copy(const void * _Src, unsigned int _Src_byte_size, texture<T,N>& _Texture);
```

Copies raw texture data to a device-side texture. The buffer must be laid out in accordance with the texture format and dimensions.

#### Parameters

_Texture	Destination texture
_Src	Pointer to source buffer on the host
_Src_byte_size	Number of bytes in the destination buffer
<b>Error condition</b>	<b>Exception thrown</b>
Out of memory	
Buffer too small	

#### 10.1.14.1 Global async texture copy functions

For each *copy* function specified above, a *copy\_async* function will also be provided, returning a *completion\_future*.

#### 10.1.15 Direct3d Interop Functions

The following functions are provided in the *direct3d* namespace in order to convert between DX COM interfaces and textures.

```
template <typename T, int N>
texture<T,N> make_texture(const Concurrency::accelerator_view &Av, const IUnknown* pTexture);
```

Creates a texture from the corresponding DX interface

#### Parameters

Av	A D3D accelerator view on which the texture is to be created.
pTexture	A pointer to a suitable texture
<b>Return value</b>	Created texture
<b>Error condition</b>	<b>Exception thrown</b>
Out of memory	
Invalid D3D texture argument	

```
template <typename T, int N>
IUnknown * get_texture(const texture<T, N>& _Texture);
```

Retrieves a DX interface pointer from a C++ AMP texture object. Class texture allows retrieving a texture interface pointer (the exact interface depends on the rank of the class).

#### Parameters

_Texture	Source texture
<b>Return value</b>	Texture interface as IUnknown *
<b>Error condition: no</b>	

## 10.2 writeonly\_texture\_view<T,N>

3596 C++ AMP write-only texture views, coded as `writeonly_texture_view<T, N>`, which provides write-only access into any  
 3597 `texture`.  
 3598

### 3599 10.2.1 Synopsis

```
3600 template <typename T, int N>
3601 class writeonly_texture_view<T,N>
3602 {
3603 public:
3604     static const int rank = _Rank;
3605     typedef typename T value_type;
3606     typedef short_vectors_traits<T>::scalar_type scalar_type;
3607
3608     writeonly_texture_view(texture<T,N>& _Src) restrict(cpu,amp);
3609
3610     writeonly_texture_view(const writeonly_texture_view&) restrict(cpu,amp);
3611
3612     writeonly_texture_view operator=(const writeonly_texture_view&) restrict(cpu,amp);
3613
3614     ~writeonly_texture_view() restrict(cpu,amp);
3615
3616     unsigned int get_Bits_per_scalar_element() const;
3617     __declspec(property(get= get_Bits_per_scalar_element)) int bits_per_scalar_element;
3618
3619     unsigned int get_data_length() const;
3620     __declspec(property(get=get_data_length)) unsigned int data_length;
3621
3622     extent<N> get_extent() const restrict(cpu,amp);
3623     __declspec(property(get=get_extent)) extent<N> extent;
3624
3625     accelerator_view get_accelerator_view() const;
3626     __declspec(property(get=get_accelerator_view)) accelerator_view accelerator_view;
3627
3628     void set(const index<N>& _Index, const value_type& _Val) const restrict(amp);
3629 };
```

### 3630 10.2.2 Introduced typedefs

```
typedef ... value_type;
```

The logical value type of the `writeonly_texture_view`. e.g., for `writeonly_texture_view<float2,3>`, `value_type` would be `float2`.

3631

```
typedef ... scalar_type;
```

The scalar type that serves as the component of the texture's value type. For example, for `writeonly_texture_view<int2,3>`, the scalar type would be "int".

### 3632 10.2.3 Construct a writeonly view over a texture

```
writeonly_texture_view(texture<T,N>& _Src) restrict(cpu);
writeonly_texture_view(texture<T,N>& _Src) restrict(amp);
```

Creates a write-only view to a given texture.

When create the `writeonly_texture_view` in a `direct3d` function, if the number of scalar elements of `T` is larger than 1, a compilation error will be given.

#### Parameters

<code>_Src</code>	Source texture
-------------------	----------------

3633

### 3634 10.2.4 Copy constructors and assignment operators

```
writeonly_texture_view(const writeonly_texture_view& _Other) restrict(cpu,amp);
writeonly_texture_view operator=(const writeonly_texture_view& _Other) restrict(cpu,amp);
```

writeonly\_texture\_views are shallow objects which can be copied and moved both on the CPU and on an accelerator. They are captured by value when passed to parallel\_for\_each

#### Parameters

_Other	Source writeonly_texture view to copy
--------	---------------------------------------

Error condition	Exception thrown
-----------------	------------------

### 10.2.5 Destructor

```
~writeonly_texture_view() restrict(cpu,amp);
```

texture\_view can be destructed on the accelerator.

**Error conditions: none**

### 10.2.6 Querying underlying texture's physical characteristics

```
unsigned int get_Bits_per_scalar_element() const;
__declspec(property(get=get_Bits_per_scalar_element)) unsigned int bits_per_scalar_element;
```

Gets the bits-per-scalar-element of the texture

**Error conditions: none**

```
unsigned int get_data_length() const;
__declspec(property(get=get_data_length)) unsigned int data_length;
```

Gets the physical data length (in bytes) that is required in order to represent the texture on the host side with its native format.

**Error conditions: none**

### 10.2.7 Querying the underlying texture's accelerator\_view

```
accelerator_view get_accelerator_view() const;
__declspec(property(get=get_accelerator_view)) accelerator_view accelerator_view;
```

Retrieves the accelerator\_view where the underlying texture resides.

**Error conditions: none**

#### 10.2.7.1 Querying underlying texture's logical dimensions (through a view)

```
extent<N> get_extent() const restrict(cpu,amp);
__declspec(property(get=get_extent)) extent<N> extent;
```

These members have the same meaning as the equivalent ones on the array class

**Error conditions: none**

#### 10.2.7.2 Writing a write-only texture view

This is the main purpose of this type. All *texture* types can be written through a write-only view.

```
void set(const index<N>& _Index, const value_type& _Val) const restrict(amp);
```

Stores one texel in the texture.

If the texture is indexed, at runtime, outside of its logical bounds, behavior is undefined.

Parameters

_Index	An N-dimension logical integer coordinate to read from
_I0, _I1, _I0	Index components
_Val	Value to store into the texture
<b>Error conditions: none</b>	

3650

3651 

## 10.2.8 Global writeonly\_texture\_view copy functions

3652

```
template <typename T, int N>
void copy(const void * _Src, unsigned int _Src_byte_size, const writeonly_texture_view<T,N>&
_TextureView);
```

Copies raw texture data to a device-side writeonly texture view. The buffer must be laid out in accordance with the texture format and dimensions.

**Parameters**

_TextureView	Destination texture view
_Src	Pointer to source buffer on the host
_Src_byte_size	Number of bytes in the destination buffer
<b>Error condition</b>	<b>Exception thrown</b>
Out of memory	
Buffer too small	

3653 

### 10.2.8.1 Global async writeonly\_texture\_view copy functions

3654 For each *copy* function specified above, a *copy\_async* function will also be provided, returning a completion\_future.3655 

## 10.2.9 Direct3d Interop Functions

3656 The following functions are provided in the *direct3d* namespace in order to convert between DX COM interfaces and  
3657 *writeonly\_texture\_views*.  
3658

```
template <typename T, int N>
IUnknown * get_texture(const writeonly_texture_view<T, N>& _TextureView);
```

Retrieves a DX interface pointer from a C++ AMP writeonly\_texture\_view object.

**Parameters**

_TextureView	Source texture view
<b>Return value</b>	Texture interface as IUnknown *
<b>Error condition: no</b>	

3659

3660 

## 10.3 norm and unorm

3661 The *norm* type is a single-precision floating point value that is normalized to the range [-1.0f, 1.0f]. The *unorm* type is a  
3662 single-precision floating point value that is normalized to the range [0.0f, 1.0f].

3663 

### 10.3.1 Synopsis

3664

```
class norm
{
public:
    norm() restrict(cpu, amp);
    explicit norm(float _V) restrict(cpu, amp);
    explicit norm(unsigned int _V) restrict(cpu, amp);
    explicit norm(int _V) restrict(cpu, amp);
    explicit norm(double _V) restrict(cpu, amp);
```

3672

```

3673     norm(const norm& _Other) restrict(cpu, amp);
3674     norm(const unorm& _Other) restrict(cpu, amp);
3675
3676     norm& operator=(const norm& _Other) restrict(cpu, amp);
3677
3678     operator float(void) const restrict(cpu, amp);
3679
3680     norm& operator+=(const norm& _Other) restrict(cpu, amp);
3681     norm& operator-=(const norm& _Other) restrict(cpu, amp);
3682     norm& operator*=(const norm& _Other) restrict(cpu, amp);
3683     norm& operator/=(const norm& _Other) restrict(cpu, amp);
3684     norm& operator++() restrict(cpu, amp);
3685     norm operator++(int) restrict(cpu, amp);
3686     norm& operator--() restrict(cpu, amp);
3687     norm operator--(int) restrict(cpu, amp);
3688     norm operator-() restrict(cpu, amp);
3689 };
3690
3691 class unorm
3692 {
3693 public:
3694     unorm() restrict(cpu, amp);
3695     explicit unorm(float _V) restrict(cpu, amp);
3696     explicit unorm(unsigned int _V) restrict(cpu, amp);
3697     explicit unorm(int _V) restrict(cpu, amp);
3698     explicit unorm(double _V) restrict(cpu, amp);
3699     unorm(const unorm& _Other) restrict(cpu, amp);
3700     explicit unorm(const norm& _Other) restrict(cpu, amp);
3701
3702     unorm& operator=(const unorm& _Other) restrict(cpu, amp);
3703
3704     operator float() const restrict(cpu, amp);
3705
3706     unorm& operator+=(const unorm& _Other) restrict(cpu, amp);
3707     unorm& operator-=(const unorm& _Other) restrict(cpu, amp);
3708     unorm& operator*=(const unorm& _Other) restrict(cpu, amp);
3709     unorm& operator/=(const unorm& _Other) restrict(cpu, amp);
3710     unorm& operator++() restrict(cpu, amp);
3711     unorm operator++(int) restrict(cpu, amp);
3712     unorm& operator--() restrict(cpu, amp);
3713     unorm operator--(int) restrict(cpu, amp);
3714 };
3715
3716 unorm operator+(const unorm& lhs, const unorm& rhs) restrict(cpu, amp);
3717 norm operator+(const norm& lhs, const norm& rhs) restrict(cpu, amp);
3718
3719 unorm operator-(const unorm& lhs, const unorm& rhs) restrict(cpu, amp);
3720 norm operator-(const norm& lhs, const norm& rhs) restrict(cpu, amp);
3721
3722 unorm operator*(const unorm& lhs, const unorm& rhs) restrict(cpu, amp);
3723 norm operator*(const norm& lhs, const norm& rhs) restrict(cpu, amp);
3724
3725 unorm operator/(const unorm& lhs, const unorm& rhs) restrict(cpu, amp);
3726 norm operator/(const norm& lhs, const norm& rhs) restrict(cpu, amp);
3727
3728 bool operator==(const unorm& lhs, const unorm& rhs) restrict(cpu, amp);
3729 bool operator==(const norm& lhs, const norm& rhs) restrict(cpu, amp);
3730

```

```

3731 bool operator!=(const unorm& lhs, const unorm& rhs) restrict(cpu, amp);
3732 bool operator!=(const norm& lhs, const norm& rhs) restrict(cpu, amp);
3733
3734 bool operator>(const unorm& lhs, const unorm& rhs) restrict(cpu, amp);
3735 bool operator>(const norm& lhs, const norm& rhs) restrict(cpu, amp);
3736
3737 bool operator<(const unorm& lhs, const unorm& rhs) restrict(cpu, amp);
3738 bool operator<(const norm& lhs, const norm& rhs) restrict(cpu, amp);
3739
3740 bool operator>=(const unorm& lhs, const unorm& rhs) restrict(cpu, amp);
3741 bool operator>=(const norm& lhs, const norm& rhs) restrict(cpu, amp);
3742
3743 bool operator<=(const unorm& lhs, const unorm& rhs) restrict(cpu, amp);
3744 bool operator<=(const norm& lhs, const norm& rhs) restrict(cpu, amp);
3745
3746 #define UNORM_MIN ((unorm)0.0f)
3747 #define UNORM_MAX ((unorm)1.0f)
3748 #define UNORM_ZERO ((norm)0.0f)
3749 #define NORM_ZERO ((norm)0.0f)
3750 #define NORM_MIN ((norm)-1.0f)
3751 #define NORM_MAX ((norm)1.0f)
3752

```

### 3753 10.3.2 Constructors and Assignment

3754 An object of type *norm* or *unorm* can be explicitly constructed from one of the following types:

- 3755 • *float*
- 3756 • *double*
- 3757 • *int*
- 3758 • *unsigned int*
- 3759 • *norm*
- 3760 • *unorm*

3761 In all these constructors, the object is initialized by first converting the argument to the *float* data type, and then clamping  
3762 the value into the range defined by the type.

3763

3764 Assignment from *norm* to *norm* is defined, as is assignment from *unorm* to *unorm*. Assignment from other types requires  
3765 an explicit conversion.

### 3766 10.3.3 Operators

3767 All arithmetic operators that are defined for the *float* type are defined for *norm* and *unorm* as well. For each supported  
3768 operator  $\oplus$ , the result is computed in single-precision floating point arithmetic, and if required is then clamped back to the  
3769 appropriate range.

3770

3771 Both *norm* and *unorm* are implicitly convertible to *float*.

## 3772 10.4 Short Vector Types

3773 C++ AMP defines a set of short vector types (of length 2, 3, and 4) which are based on one of the following scalar types: *{int,*  
3774 *unsigned int, float, double, norm, unorm}*, and are named as summarized in the following table:

3775

Scalar Type	Length		
	2	3	4
<b>int</b>	int_2, int2	int_3, int3	int_4, int4
<b>unsigned int</b>	uint_2, uint2	uint_3, uint3	uint_4, uint4
<b>float</b>	float_2, float2	float_3, float3	float_4, float4

<b>double</b>	double_2, double2	double_3, double3	double_4, double4
<b>norm</b>	norm_2, norm2	norm_3, norm3	norm_4, norm4
<b>unorm</b>	unorm_2, unorm2	unorm_3, unorm3	unorm_4, unorm4

There is no functional difference between the type `scalar_N` and `scalarN`. `scalarN` type is available in the `graphics::direct3d` namespace.

Unlike `index<N>` and `extent<N>`, short vector types have no notion of significance or endian-ness, as they are not assumed to be describing the shape of data or compute (even though a user might choose to use them this way). Also unlike extents and indices, short vector types cannot be indexed using the subscript operator.

Components of short vector types can be accessed by name. By convention, short vector type components can use either Cartesian coordinate names ("x", "y", "z", and "w"), or color scalar element names ("r", "g", "b", and "w").

- For length-2 vectors, only the names "x", "y" and "r", "g" are available.
- For length-3 vectors, only the names "x", "y", "z", and "r", "g", "b" are available.
- For length-4 vectors, the full set of names "x", "y", "z", "w", and "r", "g", "b", "a" are available.

Note that the names derived from the color channel space (rgba) are available only as properties, not as getter and setter functions.

#### 10.4.1 Synopsis

Because the full synopsis of all the short vector types is quite large, this section will summarize the basic structure of all the short vector types.

In the summary class definition below the word "scalartype" is one of { `int`, `uint`, `float`, `double`, `norm`, `unorm` }. The value `N` is 2, 3 or 4.

```
class scalartype_N
{
public:
    typedef scalartype value_type;
    static const int size = N;

    scalartype_N() restrict(cpu, amp);
    scalartype_N(scalartype value) restrict(cpu, amp);
    scalartype_N(const scalartype_N& other) restrict(cpu, amp);

    // Component-wise constructor... see 10.4.2.1 Constructors from components

    // Constructors that explicitly convert from other short vector types...
    // See 10.4.2.2 Explicit conversion constructors.

    scalartype_N& operator=(const scalartype_N& other) restrict(cpu, amp);

    // Operators
    scalartype_N& operator++() restrict(cpu, amp);
    scalartype_N operator++(int) restrict(cpu, amp);
    scalartype_N& operator--() restrict(cpu, amp);
    scalartype_N operator--(int) restrict(cpu, amp);
    scalartype_N& operator+=(const scalartype_N& rhs) restrict(cpu, amp);
    scalartype_N& operator-=(const scalartype_N& rhs) restrict(cpu, amp);
    scalartype_N& operator*=(const scalartype_N& rhs) restrict(cpu, amp);
```



```

3824     scalar_type_N& operator/=(const scalar_type_N& rhs) restrict(cpu, amp);
3825
3826     // Unary negation: not for scalar_type == uint or unorm
3827     scalar_type_N operator-() const restrict(cpu, amp);
3828
3829     // More integer operators (only for scalar_type == int or uint)
3830     scalar_type_N operator~() const restrict(cpu, amp);
3831     scalar_type_N& operator%=(const scalar_type_N& rhs) restrict(cpu, amp);
3832     scalar_type_N& operator^=(const scalar_type_N& rhs) restrict(cpu, amp);
3833     scalar_type_N& operator|=(const scalar_type_N& rhs) restrict(cpu, amp);
3834     scalar_type_N& operator&=(const scalar_type_N& rhs) restrict(cpu, amp);
3835     scalar_type_N& operator>>=(const scalar_type_N& rhs) restrict(cpu, amp);
3836     scalar_type_N& operator<<=(const scalar_type_N& rhs) restrict(cpu, amp);
3837
3838     // Component accessors and properties (a.k.a. swizzling):
3839     // See 10.4.3 Component Access (Swizzling)
3840 };
3841
3842     scalar_type_N operator+(const scalar_type_N& lhs, const scalar_type_N& rhs) restrict(cpu, amp);
3843     scalar_type_N operator-(const scalar_type_N& lhs, const scalar_type_N& rhs) restrict(cpu, amp);
3844     scalar_type_N operator*(const scalar_type_N& lhs, const scalar_type_N& rhs) restrict(cpu, amp);
3845     scalar_type_N operator/(const scalar_type_N& lhs, const scalar_type_N& rhs) restrict(cpu, amp);
3846     bool operator==(const scalar_type_N& lhs, const scalar_type_N& rhs) restrict(cpu, amp);
3847     bool operator!=(const scalar_type_N& lhs, const scalar_type_N& rhs) restrict(cpu, amp);
3848
3849     // More integer operators (only for scalar_type == int or uint)
3850     scalar_type_N operator%(const scalar_type_N& lhs, const scalar_type_N& rhs) restrict(cpu, amp);
3851     scalar_type_N operator^(const scalar_type_N& lhs, const scalar_type_N& rhs) restrict(cpu, amp);
3852     scalar_type_N operator|(const scalar_type_N& lhs, const scalar_type_N& rhs) restrict(cpu, amp);
3853     scalar_type_N operator&(const scalar_type_N& lhs, const scalar_type_N& rhs) restrict(cpu, amp);
3854     scalar_type_N operator<<(const scalar_type_N& lhs, const scalar_type_N& rhs) restrict(cpu, amp);
3855     scalar_type_N operator>>(const scalar_type_N& lhs, const scalar_type_N& rhs) restrict(cpu, amp);

```

## 10.4.2 Constructors

**scalar\_type\_N()** restrict(cpu,amp)

Default constructor. Initializes all components to zero.

**scalar\_type\_N(scalar\_type value)** restrict(cpu,amp)

Initializes all components of the short vector to 'value'.

**Parameters:**

<i>value</i>	The value with which to initialize each component of this vector.
--------------	---

**scalar\_type\_N(const scalar\_type\_N& other)** restrict(cpu,amp)

Copy constructor. Copies the contents of 'other' to 'this'.

**Parameters:**

<i>other</i>	The source vector to copy from.
--------------	---------------------------------

### 10.4.2.1 Constructors from components

A short vector type can also be constructed with values for each of its components.

**scalar\_type\_2(scalar\_type v1, scalar\_type v2)** restrict(cpu,amp) // only for length 2

```

scalartype_3(scalartype v1, scalartype v2, scalartype v3) restrict(cpu,amp) // only for length 3
scalartype_4(scalartype v1, scalartype v2,
              scalartype v3, scalartype v4) restrict(cpu,amp) // only for length 4

```

Creates a short vector with the provided initialize values for each component.

**Parameters:**

<i>v1</i>	The value with which to initialize the "x" (or "r") component.
<i>v2</i>	The value with which to initialize the "y" (or "g") component
<i>v3</i>	The value with which to initialize the "z" (or "b") component.
<i>v4</i>	The value with which to initialize the "w" (or "a") component

3864

3865 

### 10.4.2.2 Explicit conversion constructors

3866 A short vector of type *scalartype<sub>1</sub>\_N* can be constructed from an object of type *scalartype<sub>2</sub>\_N*, as long as *N* is the same in  
 3867 both types. For example, a *uint\_4* can be constructed from a *float\_4*.  
 3868

```

explicit scalartype_N(const int_N& other) restrict(cpu,amp)
explicit scalartype_N(const uint_N& other) restrict(cpu,amp)
explicit scalartype_N(const float_N& other) restrict(cpu,amp)
explicit scalartype_N(const double_N& other) restrict(cpu,amp)
explicit scalartype_N(const norm_N& other) restrict(cpu,amp)
explicit scalartype_N(const unorm_N& other) restrict(cpu,amp)

```

Construct a short vector from a differently-typed short vector, performing an explicit conversion. Note that in the above list of 6 constructors, each short vector type will have 5 of these.

**Parameters:**

<i>other</i>	The source vector to copy/convert from.
--------------	---

3869 

### 10.4.3 Component Access (Swizzling)

3870 The components of a short vector may be accessed in a large variety of ways, depending on the length of the short vector.

- 3871
- As single scalar components ( $N \geq 2$ )
  - 3872 • As pairs of components, in any permutation ( $N \geq 2$ )
  - 3873 • As triplets of components, in any permutation ( $N \geq 3$ )
  - 3874 • As quadruplets of components, in any permutation ( $N = 4$ ).

3875

3876 Because the permutations of such component accessors are so large, they are described here using symmetric group  
 3877 notation. In such notation,  $S_{xy}$  represents all permutations of the letters *x* and *y*, namely *xy* and *yx*. Similarly,  $S_{xyz}$  represents  
 3878 all  $3! = 6$  permutations of the letters *x*, *y*, and *z*, namely *xy*, *xz*, *yx*, *yz*, *zx*, and *zy*.  
 3879

3880 Recall that the *z* (or *b*) component of a short vector is only available for vector lengths 3 and 4. The *w* (or *a*) component of a  
 3881 short vector is only available for vector length 4.  
 3882

3883 

#### 10.4.3.1 Single-component access

```

scalartype get_x() const restrict(cpu,amp)
scalartype get_y() const restrict(cpu,amp)
scalartype get_z() const restrict(cpu,amp)
scalartype get_w() const restrict(cpu,amp)

```

```

void set_x(scalartype v) restrict(cpu,amp)
void set_y(scalartype v) restrict(cpu,amp)
void set_z(scalartype v) restrict(cpu,amp)
void set_w(scalartype v) restrict(cpu,amp)

__declspec(property(get=get_x, put=set_x)) scalartype x
__declspec(property(get=get_y, put=set_y)) scalartype y
__declspec(property(get=get_z, put=set_z)) scalartype z
__declspec(property(get=get_w, put=set_w)) scalartype w
__declspec(property(get=get_x, put=set_x)) scalartype r
__declspec(property(get=get_y, put=set_y)) scalartype g
__declspec(property(get=get_z, put=set_z)) scalartype b
__declspec(property(get=get_w, put=set_w)) scalartype a

```

These functions (and properties) allow access to individual components of a short vector type. Note that the properties in the “rgba” space map to functions in the “xyzw” space.

3884

### 3885 10.4.3.2 Two-component access

```

scalartype_2 get_Sxy() const restrict(cpu,amp)
scalartype_2 get_Sxz() const restrict(cpu,amp)
scalartype_2 get_Sxw() const restrict(cpu,amp)
scalartype_2 get_Syz() const restrict(cpu,amp)
scalartype_2 get_Syw() const restrict(cpu,amp)
scalartype_2 get_Szw() const restrict(cpu,amp)

void set_Sxy(scalartype_2 v) restrict(cpu,amp)
void set_Sxz(scalartype_2 v) restrict(cpu,amp)
void set_Sxw(scalartype_2 v) restrict(cpu,amp)
void set_Syz(scalartype_2 v) restrict(cpu,amp)
void set_Syw(scalartype_2 v) restrict(cpu,amp)
void set_Szw(scalartype_2 v) restrict(cpu,amp)

__declspec(property(get=get_Sxy, put=set_Sxy)) scalartype_2 Sxy
__declspec(property(get=get_Sxz, put=set_Sxz)) scalartype_2 Sxz
__declspec(property(get=get_Sxw, put=set_Sxw)) scalartype_2 Sxw
__declspec(property(get=get_Syz, put=set_Syz)) scalartype_2 Syz
__declspec(property(get=get_Syw, put=set_Syw)) scalartype_2 Syw
__declspec(property(get=get_Szw, put=set_Szw)) scalartype_2 Szw
__declspec(property(get=get_Sxy, put=set_Sxy)) scalartype_2 Srg
__declspec(property(get=get_Sxz, put=set_Sxz)) scalartype_2 Srb
__declspec(property(get=get_Sxw, put=set_Sxw)) scalartype_2 Sra
__declspec(property(get=get_Syz, put=set_Syz)) scalartype_2 Sgb
__declspec(property(get=get_Syw, put=set_Syw)) scalartype_2 Sga
__declspec(property(get=get_Szw, put=set_Szw)) scalartype_2 Sba

```

These functions (and properties) allow access to pairs of components. For example:

```

int_3 f3(1,2,3);
int_2 yz = f3.yz; // yz = (2,3)

```

3886

### 3887 10.4.3.3 Three-component access

```

scalartype_3 get_Sxyz() const restrict(cpu,amp)
scalartype_3 get_Sxyzw() const restrict(cpu,amp)

```

```

scalartype_3 get_Sxzw() const restrict(cpu,amp)
scalartype_3 get_Syzw() const restrict(cpu,amp)

void set_Sxyz(scalartype_3 v) restrict(cpu,amp)
void set_Sxyw(scalartype_3 v) restrict(cpu,amp)
void set_Sxzw(scalartype_3 v) restrict(cpu,amp)
void set_Syzw(scalartype_3 v) restrict(cpu,amp)

__declspec(property(get=get_Sxyz, put=set_Sxyz)) scalartype_3 Sxyz
__declspec(property(get=get_Sxyw, put=set_Sxyw)) scalartype_3 Sxyw
__declspec(property(get=get_Sxzw, put=set_Sxzw)) scalartype_3 Sxzw
__declspec(property(get=get_Syzw, put=set_Syzw)) scalartype_3 Syzw
__declspec(property(get=get_Sxyz, put=set_Sxyz)) scalartype_3 Srgb
__declspec(property(get=get_Sxyw, put=set_Sxyw)) scalartype_3 Srga
__declspec(property(get=get_Sxzw, put=set_Sxzw)) scalartype_3 Srba
__declspec(property(get=get_Syzw, put=set_Syzw)) scalartype_3 Sgba

```

These functions (and properties) allow access to triplets of components (for vectors of length 3 or 4). For example:

```

int_4 f3(1,2,3,4);
int_3 wzy = f3.wzy; // wzy = (4,3,2)

```

3888

#### 3889 10.4.3.4 Four-component access

```

scalartype_4 get_Sxyzw() const restrict(cpu,amp)

void set_Sxyzw(scalartype_4 v) restrict(cpu,amp)

__declspec(property(get=get_Sxyzw, put=set_Sxyzw)) scalartype_4 Sxyzw
__declspec(property(get=get_Sxyzw, put=set_Sxyzw)) scalartype_4 Srgba

```

These functions (and properties) allow access to all four components (obviously, only for vectors of length 4). For example:

```

int_4 f3(1,2,3,4);
int_4 wzyx = f3.wzyw; // wzyx = (4,3,2,1)

```

3890

### 3891 10.5 Template class short\_vector\_traits

3892 The template class short\_vector\_traits provides the ability to reflect on the supported short vector types and obtain the  
 3893 length of the vector and the underlying scalar type.

#### 3894 10.5.1 Synopsis

```

3895
3896 template<typename _Type> struct short_vector_traits
3897 {
3898     short_vector_traits()
3899     {
3900         static_assert(false, "short_vector_traits is not supported for this type (_Type)");
3901     }
3902 };
3903
3904 template<>
3905 struct short_vector_traits<unsigned int>
3906 {
3907     typedef unsigned int value_type;
3908     static int const size = 1;
3909 };

```

```

3910
3911 template<>
3912 struct short_vector_traits<uint_2>
3913 {
3914     typedef unsigned int value_type;
3915     static int const size = 2;
3916 };
3917
3918 template<>
3919 struct short_vector_traits<uint_3>
3920 {
3921     typedef unsigned int value_type;
3922     static int const size = 3;
3923 };
3924
3925 template<>
3926 struct short_vector_traits<uint_4>
3927 {
3928     typedef unsigned int value_type;
3929     static int const size = 4;
3930 };
3931
3932 template<>
3933 struct short_vector_traits<int>
3934 {
3935     typedef int value_type;
3936     static int const size = 1;
3937 };
3938
3939 template<>
3940 struct short_vector_traits<int_2>
3941 {
3942     typedef int value_type;
3943     static int const size = 2;
3944 };
3945
3946 template<>
3947 struct short_vector_traits<int_3>
3948 {
3949     typedef int value_type;
3950     static int const size = 3;
3951 };
3952
3953 template<>
3954 struct short_vector_traits<int_4>
3955 {
3956     typedef int value_type;
3957     static int const size = 4;
3958 };
3959
3960 template<>
3961 struct short_vector_traits<float>
3962 {
3963     typedef float value_type;
3964     static int const size = 1;
3965 };
3966
3967 template<>

```

```

3968 struct short_vector_traits<float_2>
3969 {
3970     typedef float value_type;
3971     static int const size = 2;
3972 };
3973
3974 template<>
3975 struct short_vector_traits<float_3>
3976 {
3977     typedef float value_type;
3978     static int const size = 3;
3979 };
3980
3981 template<>
3982 struct short_vector_traits<float_4>
3983 {
3984     typedef float value_type;
3985     static int const size = 4;
3986 };
3987
3988 template<>
3989 struct short_vector_traits<unorm>
3990 {
3991     typedef unorm value_type;
3992     static int const size = 1;
3993 };
3994
3995 template<>
3996 struct short_vector_traits<unorm_2>
3997 {
3998     typedef unorm value_type;
3999     static int const size = 2;
4000 };
4001
4002 template<>
4003 struct short_vector_traits<unorm_3>
4004 {
4005     typedef unorm value_type;
4006     static int const size = 3;
4007 };
4008
4009 template<>
4010 struct short_vector_traits<unorm_4>
4011 {
4012     typedef unorm value_type;
4013     static int const size = 4;
4014 };
4015
4016 template<>
4017 struct short_vector_traits<norm>
4018 {
4019     typedef norm value_type;
4020     static int const size = 1;
4021 };
4022
4023 template<>
4024 struct short_vector_traits<norm_2>
4025 {

```

```

4026     typedef norm value_type;
4027     static int const size = 2;
4028 };
4029
4030 template<>
4031 struct short_vector_traits<norm_3>
4032 {
4033     typedef norm value_type;
4034     static int const size = 3;
4035 };
4036
4037 template<>
4038 struct short_vector_traits<norm_4>
4039 {
4040     typedef norm value_type;
4041     static int const size = 4;
4042 };
4043
4044 template<>
4045 struct short_vector_traits<double>
4046 {
4047     typedef double value_type;
4048     static int const size = 1;
4049 };
4050
4051 template<>
4052 struct short_vector_traits<double_2>
4053 {
4054     typedef double value_type;
4055     static int const size = 2;
4056 };
4057
4058 template<>
4059 struct short_vector_traits<double_3>
4060 {
4061     typedef double value_type;
4062     static int const size = 3;
4063 };
4064
4065 template<>
4066 struct short_vector_traits<double_4>
4067 {
4068     typedef double value_type;
4069     static int const size = 4;
4070 };

```

## 10.5.2 Typedefs

### *typedef scalar\_type value\_type*

Introduces a typedef identifying the underlying scalar type of the vector type. `scalar_type` depends on the instantiation of class `short_vector_types` used. This is summarized in the list below

Instantiated Type	Scalar Type
<code>short_vector_type&lt;unsigned int&gt;</code>	unsigned int
<code>short_vector_type&lt;uint_2&gt;</code>	unsigned int
<code>short_vector_type&lt;uint_3&gt;</code>	unsigned int

short_vector_type<uint_4>	unsigned int
short_vector_type<int>	int
short_vector_type<int_2>	int
short_vector_type<int_3>	int
short_vector_type<int_4>	int
short_vector_type<float>	float
short_vector_type<float_2>	float
short_vector_type<float_3>	float
short_vector_type<float_4>	float
short_vector_type<unorm>	norm
short_vector_type<unorm_2>	norm
short_vector_type<unorm_3>	norm
short_vector_type<unorm_4>	norm
short_vector_type<norm>	norm
short_vector_type<norm_2>	norm
short_vector_type<norm_3>	norm
short_vector_type<norm_4>	norm
short_vector_type<double>	double
short_vector_type<double_2>	double
short_vector_type<double_3>	double
short_vector_type<double_4>	double

4073  
4074

### 10.5.3 Members

*static int const size;*

Introduces a static constant integer specifying the number of elements in the short vector type, based on the table below:

Instantiated Type	Scalar Type
short_vector_type<unsigned int>	1
short_vector_type<uint_2>	2
short_vector_type<uint_3>	3
short_vector_type<uint_4>	4
short_vector_type<int>	1
short_vector_type<int_2>	2
short_vector_type<int_3>	3
short_vector_type<int_4>	4
short_vector_type<float>	1
short_vector_type<float_2>	2
short_vector_type<float_3>	3
short_vector_type<float_4>	4
short_vector_type<unorm>	1
short_vector_type<unorm_2>	2
short_vector_type<unorm_3>	3
short_vector_type<unorm_4>	4
short_vector_type<norm>	1



short_vector_type<norm_2>	2	
short_vector_type<norm_3>	3	
short_vector_type<norm_4>	4	
short_vector_type<double>	1	
short_vector_type<double_2>	2	
short_vector_type<double_3>	3	
short_vector_type<double_4>	4	

4075

## 4076 11 D3D interoperability (Optional)

4077

4078 The C++ AMP runtime provides functions for D3D interoperability, enabling seamless use of D3D resources for compute in  
 4079 C++ AMP code as well as allow use of resources created in C++ AMP in D3D code, without the creation of redundant  
 4080 intermediate copies. These features allow users to incrementally accelerate the compute intensive portions of their DirectX  
 4081 applications using C++ AMP and use the D3D API on data produced from C++ AMP computations.

4082

4083 The following D3D interoperability functions are available in the *direct3d* namespace:

4084

```
accelerator_view create_accelerator_view(IUnknown *_D3d_device_interface)
```

Creates a new *accelerator\_view* from an existing Direct3D device interface pointer. On failure the function throws a *runtime\_exception* exception. On success, the reference count of the parameter is incremented by making a *AddRef* call on the interface to record the C++ AMP reference to the interface, and users can safely *Release* the object when no longer required in their DirectX code.

The *accelerator\_view* created using this function is thread-safe just as any C++ AMP created *accelerator\_view*, allowing concurrent submission of commands to it from multiple host threads. However, concurrent use of the *accelerator\_view* and the raw *ID3D11Device* interface from multiple host threads must be properly synchronized by users to ensure mutual exclusion. Unsynchronized concurrent usage of the *accelerator\_view* and the raw *ID3D11Device* interface will result in undefined behavior.

The C++ AMP runtime provides detailed error information in debug mode using the Direct3D Debug layer. However, if the Direct3D device passed to the above function was not created with the *D3D11\_CREATE\_DEVICE\_DEBUG* flag, the C++ AMP debug mode detailed error information support will be unavailable.

### Parameters:

*\_D3d\_device\_interface*

An AMP supported D3D device interface pointer to be used to create the *accelerator\_view*. The parameter must meet all of the following conditions for successful creation of a *accelerator\_view*:

- 1) Must be a supported D3D device interface. For this release, only *ID3D11Device* interface is supported.
- 2) The device must have an AMP supported feature level. For this release this means a *D3D\_FEATURE\_LEVEL\_11\_0*.
- 3) The D3D Device should not have been created with the "D3D11\_CREATE\_DEVICE\_SINGLETHREADED" flag.

### Return Value:

The newly created *accelerator\_view* object.

### Exceptions:

*runtime\_exception*

- 1) "Failed to create *accelerator\_view* from D3D device.", *E\_INVALIDARG*
- 2) "NULL D3D device pointer.", *E\_INVALIDARG*

4085  
4086

IUnknown * get_device(const accelerator_view & _Rv)	
<p>Returns a D3D device interface pointer underlying the passed accelerator_view. Fails with a “runtime_exception” exception of the passed accelerator_view is not a D3D device resource view. On success, it increments the reference count of the D3D device interface by calling “AddRef” on the interface. Users must call “Release” on the returned interface after they are finished using it, for proper reclamation of the resources associated with the object.</p> <p>Concurrent use of the accelerator_view and the raw ID3D11Device interface from multiple host threads must be properly synchronized by users to ensure mutual exclusion. Unsynchronized concurrent usage of the accelerator_view and the raw ID3D11Device interface will result in undefined behavior.</p>	
<b>Parameters:</b>	
_Rv	The accelerator_view object for which the D3D device interface is needed.
<b>Return Value:</b>	
A IUnknown interface pointer corresponding to the D3D device underlying the passed accelerator_view. Users must use the <a href="#">QueryInterface</a> member function on the returned interface to obtain the correct D3D device interface pointer.	
<b>Exceptions:</b>	
runtime_exception	<ol style="list-style-type: none"> <li>1) "Uninitialized resource view argument.", E_INVALIDARG</li> <li>2) "Cannot get D3D device from a non-D3D accelerator_view.", E_INVALIDARG</li> </ol>

4087  
4088

<pre>template &lt;typename T, int N&gt; array&lt;T,N&gt; make_array(const extent&lt;N&gt; &amp; _Extent,                      const accelerator_view &amp; _Rv,                      IUnknown * _D3d_buffer_interface)</pre>	
<p>Creates an array with the specified extents on the specified accelerator_view from an existing Direct3D buffer interface pointer. On failure the member function throws a <a href="#">runtime_exception</a> exception. On success, the reference count of the Direct3D buffer object is incremented by making an <a href="#">AddRef</a> call on the interface to record the C++ AMP reference to the interface, and users can safely <a href="#">Release</a> the object when no longer required in their DirectX code.</p>	
<b>Parameters:</b>	
_Extent	The extent of the array to be created.
_Rv	The accelerator_view that the array is to be created on.
_D3d_buffer_interface	<p>AN AMP supported D3D device buffer pointer to be used to create the array. The parameter must meet all of the following conditions for successful creation of a accelerator_view:</p> <ol style="list-style-type: none"> <li>1) Must be a supported D3D buffer interface. For this release, only ID3D11Buffer interface is supported.</li> <li>2) The D3D device on which the buffer was created must be the same as that underlying the accelerator_view parameter rv.</li> <li>3) The D3D buffer must additionally satisfy the following conditions: <ol style="list-style-type: none"> <li>a. The buffer size in bytes must be equal to the size in bytes of the field to be created (g.get_size() * sizeof(_Elem_type)).</li> <li>b. Must have been create with DEFAULT_USAGE.</li> <li>c. SHADER_RESOURCE and UNORDERED_ACCESS bindings should be allowed for the buffer.</li> </ol> </li> <li>4) The D3D buffer must be a STRUCTURED_BUFFER with a structure byte stride of 4.</li> </ol>

<b>Return Value:</b>	
The newly created array object.	
<b>Exceptions:</b>	
runtime_exception	1) "Invalid extents argument.", E_INVALIDARG 2) "Uninitialized resource view argument.", E_INVALIDARG 3) "NULL D3D buffer pointer.", E_INVALIDARG 4) "Invalid D3D buffer argument.", E_INVALIDARG 5) "Cannot create D3D buffer on a non-D3D accelerator_view.", E_INVALIDARG

4089  
4090

<pre>template &lt;size_t RANK, typename _Elem_type&gt; IUnknown * get_d3d_buffer_interface(const array&lt;_Elem_type, RANK&gt; &amp;_F)</pre>	
Returns a D3D buffer interface pointer underlying the passed array. Fails with a "runtime_exception" exception if the passed array is not on a D3D device resource view. On success, it increments the reference count of the D3D buffer interface by calling "AddRef" on the interface. Users must call "Release" on the returned interface after they are finished using it, for proper reclamation of the resources associated with the object.	
<b>Parameters:</b>	
<code>_F</code>	The array for which the underlying D3D buffer interface is needed.
<b>Return Value:</b>	
A IUnknown interface pointer corresponding to the D3D buffer underlying the passed array. Users must use the QueryInterface member function on the returned interface to obtain the correct D3D buffer interface pointer.	
<b>Exceptions:</b>	
runtime_exception	"Cannot get D3D buffer from a non-D3D array.", E_INVALIDARG

4091  
4092

## 12 Error Handling

4093  
4094

### 12.1 static\_assert

4095  
4096

The C++ intrinsic `static_assert` is often used to handle error states that are detectable at compile time. In this way `static_assert` is a technique for conveying static semantic errors and as such they will be categorized similar to exception types.

4097  
4098  
4099  
4100

### 12.2 Runtime errors

4101  
4102

On encountering an irrecoverable error, C++ AMP runtime throws a C++ exception to communicate/propagate the error to client code. (Note: exceptions are not thrown from `restrict(amp)` code.) The actual exceptions thrown by each API are listed in the API descriptions. Following are the exception types thrown by C++ AMP runtime:

4103  
4104  
4105  
4106

#### 12.2.1 runtime\_exception

4107  
4108

A `runtime_exception` instance comprises a textual description of the error and a `HRESULT` error code to indicate the cause of the error.

4109  
4110  
4111

**class runtime\_exception**

The exception type that all AMP runtime exceptions derive from. A *runtime\_exception* instance comprises of a textual description of the error and a HRESULT error code to indicate the cause of the error.

```
runtime_exception(const char * _Message, HRESULT _Hresult) throw()
```

Construct a runtime\_exception exception with the specified message and HRESULT error code.

**Parameters:**

<i>_Message</i>	Descriptive message of error
<i>_Hresult</i>	HRESULT error code that caused this exception

```
runtime_exception (HRESULT _Hresult) throw()
```

Construct a runtime\_exception exception with the specified HRESULT error code.

**Parameters:**

<i>_Hresult</i>	HRESULT error code that caused this exception
-----------------	---

```
HRESULT get_error_code() const throw()
```

Returns the error code that caused **this** exception.

**Return Value:**

Returns the HRESULT error code that caused **this** exception.

**12.2.1.1 Specific Runtime Exceptions**

Exception String	Source	Explanation
No supported accelerator available.	Accelerator constructor, array constructor	No device available at runtime supports C++ AMP.
Failed to create buffer	Array constructor	Couldn't create buffer on accelerator, likely due to lack of resource availability.

**12.2.2 out\_of\_memory**

An instance of this exception type is thrown when an underlying OS/DirectX API call fails due to failure to allocate system or device memory (*E\_OUTOFMEMORY HRESULT* error code). Note that if the runtime fails to allocate memory from the heap using the C++ *new* operator, a *std::bad\_alloc* exception is thrown and not the C++ AMP *out\_of\_memory* exception.

**class out\_of\_memory : public runtime\_exception**

Exception thrown when an underlying OS/DirectX call fails due to lack of system or device memory.

```
explicit out_of_memory(const char * _Message) throw()
```

Construct a `out_of_memory` exception with the specified message.

**Parameters:**

<code>_Message</code>	Descriptive message of error
-----------------------	------------------------------

`out_of_memory() throw()`

Construct a `out_of_memory` exception.

**Parameters:**

None.

### 12.2.3 `invalid_compute_domain`

An instance of this exception type is thrown when the runtime fails to devise a dispatch for the compute domain specified at a `parallel_for_each` call site.

`class invalid_compute_domain : public runtime_exception`

Exception thrown when the runtime fails to launch a kernel using the compute domain specified at the `parallel_for_each` call site.

`explicit invalid_compute_domain(const char * _Message) throw()`

Construct an `invalid_compute_domain` exception with the specified message.

**Parameters:**

<code>_Message</code>	Descriptive message of error
-----------------------	------------------------------

`invalid_compute_domain() throw()`

Construct an `invalid_compute_domain` exception.

**Parameters:**

None.

### 12.2.4 `unsupported_feature`

An instance of this exception type is thrown on executing a `restrict(amp)` function on the host which uses an intrinsic unsupported on the host (such as `tilde_index<>::barrier.wait()`) or when invoking a `parallel_for_each` or allocating an object on an accelerator which doesn't support certain features which are required for the execution to proceed, such as, but not limited to:

1. The accelerator is not capable of executing code, but serves as a memory allocation arena only
2. The accelerator doesn't support the allocation of textures
3. A texture object is created with an invalid combination of `bits_per_scalar_element` and short-vector type
4. Read and write operations are both requested on a texture object with `bits_per_scalar != 32`

`class unsupported_feature : public runtime_exception`

Exception thrown when an unsupported feature is used.

```
explicit unsupported_feature (const char * _Message) throw()
```

Construct an unsupported\_feature exception with the specified message.

**Parameters:**

<code>_Message</code>	Descriptive message of error
-----------------------	------------------------------

```
unsupported_feature () throw()
```

Construct an unsupported\_feature exception.

**Parameters:**

None.

### 12.2.5 accelerator\_view\_removed

An instance of this exception type is thrown when the C++ AMP runtime detects that a connection with a particular accelerator, represented by an instance of class `accelerator_view`, has been lost. When such an incident happens, all data allocated through the accelerator view and all in-progress computations on the accelerator view may be lost. This exception may be thrown by *parallel\_for\_each*, as well as any other copying and/or synchronization method.

```
class accelerator_view_removed : public runtime_exception
```

HRESULT error code indicating the cause of removal of the accelerator\_view

```
explicit accelerator_view_removed(const char * _Message, HRESULT _View_removed_reason) throw();
explicit accelerator_view_removed(HRESULT _View_removed_reason) throw();
```

Construct an accelerator\_view\_removed exception with the specified message and HRESULT

**Parameters:**

<code>_Message</code>	Descriptive message of error
<code>_HRESULT</code>	HRESULT error code indicating the cause of removal of the accelerator_view

```
HRESULT get_view_removed_reason() const throw();
```

Provides the HRESULT error code indicating the cause of removal of the accelerator\_view

**Return Value:**

The HRESULT error code indicating the cause of removal of the accelerator\_view

## 12.3 Error handling in device code (amp-restricted functions) (Optional)

The use of the *throw* C++ keyword is disallowed in C++ AMP vector functions (*amp* restricted) and will result in a compilation error. C++ AMP offers the following intrinsics in vector code for error handling.

**Microsoft-specific:** the Microsoft implementation of C++ AMP provides the methods specified in this section, provided all of the following conditions are met.

1. The debug version of the runtime is being used (i.e. the code is compiled with the `_DEBUG` preprocessor definition).
2. The debug layer is available on the system. This, in turn requires DirectX SDK to be installed on the system on Windows 7. On Windows 8 no SDK installation is necessary..

3. *The accelerator\_view on which the kernel is invoked must be on a device which supports the printf and abort intrinsics. As of the date of writing this document, only the REF device supports these intrinsics.*

*When the debug version of the runtime is not used or the debug layer is unavailable, executing a kernel that using these intrinsics through a parallel\_for\_each call will result in a runtime exception. On devices that do not support these intrinsics, these intrinsics will behave as no-ops.*

**void** direct3d\_printf(const char \*\_Format\_string, ...) restrict(amp)

Prints formatted output from a kernel to the debug output. The formatting semantics are same as the C Library printf function. Also, this function is executed as any other device-side function: per-thread, and in the context of the calling thread. Due to the asynchronous nature of kernel execution, the output from this call may appear anytime between the launch of the kernel containing the printf call and completion of the kernel's execution.

**Parameters:**

<i>_Format_string</i>	The format string.
...	An optional list of parameters of variable count.

**Return Value:**

None.

**void** direct3d\_errorf(char \*\_Format\_string, ...) restrict(amp)

This intrinsic prints formatted error messages from a kernel to the debug output. This function is executed as any other device-side function: per-thread, and in the context of the calling thread. Note that due to the asynchronous nature of kernel execution, the actual error messages may appear in the debug output asynchronously, any time between the dispatch of the kernel and the completion of the kernel's execution. When these error messages are detected by the runtime, it raises a "runtime\_exception" exception on the host with the formatted error message output as the exception message.

**Parameters:**

<i>_Format_string</i>	The format string.
...	An optional list of parameters of variable count.

**void** direct3d\_abort() restrict(amp)

This intrinsic aborts the execution of threads in the compute domain of a kernel invocation, that execute this instruction. This function is executed as any other device-side function: per-thread, and in the context of the calling thread. Also the thread is terminated without executing any destructors for local variables. When the abort is detected by the runtime, it raises a "runtime\_exception" exception on the host with the abort output as the exception message. Note that due to the asynchronous nature of kernel execution, the actual abort may be detected any time between the dispatch of the kernel and the completion of the kernel's execution.

Due to the asynchronous nature of kernel execution, the [direct3d\\_printf](#), [direct3d\\_errorf](#) and [direct3d\\_abort](#) messages from kernels executing on a device appear asynchronously during the execution of the shader or after its completion and not immediately after the async launch of the kernel. Thus these messages from a kernel may be interleaved with messages from other kernels executing concurrently or error messages from other runtime calls in the debug output. It is the programmer's responsibility to include appropriate information in the messages originating from kernels to indicate the origin of the messages.

## 13 Appendix: C++ AMP Future Directions (Informative)

It is likely that C++ AMP will evolve over time. The set of features allowed inside *amp*-restricted functions will grow. However, compilers will have to continue to support older hardware targets which only support the previous, smaller feature set. This section outlines possible such evolution of the language syntax and associated feature set.

### 13.1 Versioning Restrictions

This section contains an informative description of additional language syntax and rules to allow the versioning of C++ AMP code. If an implementation desires to extend C++ AMP in a manner not covered by this version of the specification, it is recommended that it follows the syntax and rules specified here.

#### 13.1.1 *auto* restriction

The *restriction* production (section 2.1) of the C++ grammar is amended to allow the contextual keyword *auto*.

```
restriction:
    amp-restriction
    cpu
    auto
```

A function or lambda which is annotated with *restrict(auto)* directs the compiler to check all known restrictions and automatically deduce the set of restrictions that a function complies with. *restrict(auto)* is only allowed for functions where the function declaration is also a function definition, and no other declaration of the same function occurs.

A function may be simultaneously explicitly and *auto* restricted, e.g., *restrict(cpu,auto)*. In such case, it will be explicitly checked for compulsory conformance with the set of explicitly specified (non-*auto*) restrictions, and implicitly checked for possible conformance with all other restrictions that the compiler supports.

Consider the following example:

```
int f1() restrict(amp);
int f2() restrict(cpu,auto)
{
    f1();
}
```

In this example, *f2* is verified for compulsory adherence to the *restrict(cpu)* restriction. This results in an error, since *f2* calls *f1*, which is not *cpu*-restricted. Had we changed *f1*'s restriction to *restrict(cpu)*, then *f2* will pass the adherence test to the explicitly specified *restrict(cpu)*. Now with respect to the *auto* restriction, the compiler has to check whether *f2* conforms to *restrict(amp)*, which is the only other restriction not explicitly specified. In the context of verifying the plausibility of inferring an *amp*-restriction for *f2*, the compiler notices that *f2* calls *f1*, which is, in our modified example, not *amp*-restricted, and therefore *f2* is also inferred to be not *amp*-restricted. Thus the total inferred restriction for *f2* is *restrict(cpu)*. If we now change the restriction for *f1* into *restrict(cpu,amp)*, then the inference for *f2* would reach the conclusion that *f2* is *restrict(cpu,amp)* too.

When two overloads are available to call from a given restriction context, and they differ only by the fact that one is explicitly restricted while the other is implicitly inferred to be restricted, the explicitly restricted overload shall be chosen.

#### 13.1.2 Automatic restriction deduction

Implementations are encouraged to support a mode in which functions that have their definitions accompany their declarations, and where no other declarations occur for such functions, have their restriction set automatically deduced.



In such a mode, when the compiler encounters a function declaration which is also a definition, and a previous declaration for the function hasn't been encountered before, then the compiler analyses the function as if it was restricted with `restrict(cpu,auto)`. This allows easy reuse of existing code in `amp`-restricted code, at the cost of prolonged compilation times.

### 13.1.3 `amp` Version

The `amp`-restriction production of the C++ grammar is amended thus:

*amp-restriction:*

**amp** *amp-version*<sub>opt</sub>

*amp-version:*

: *integer-constant*

: *integer-constant* . *integer-constant*

An `amp` version specifies the lowest version of `amp` that this function supports. In other words, if a function is decorated with `restrict(amp:1)`, then that function also supports any version greater or equal to 1. When the `amp` version is elided, the implied version is implementation-defined. Implementations are encouraged to support a compiler flag controlling the default version assumed. When versioning is used in conjunction with `restrict(auto)` and/or automatic restriction deduction, the compiler shall infer the maximal version of the `amp` restriction that the function adheres to.

Section 2.3.2 specifies that restriction specifiers of a function shall not overlap with any restriction specifiers in another function within the same overload set.

```
int func(int x) restrict(cpu,amp);
int func(int x) restrict(cpu); // error, overlaps with previous declaration
```

This rule is relaxed in the case of versioning: functions overloaded with `amp` versions are not considered to overlap:

```
int func(int x) restrict(cpu);
int func(int x) restrict(amp:1);
int func(int x) restrict(amp:2);
```

When an overload set contains multiple versions of the `amp` specifier, the function with the highest version number that is not higher than the callee is chosen:

```
void glorp() restrict(amp:1) { }
void glorp() restrict(amp:2) { }

void glorp_caller() restrict(amp:2) {
    glorp(); // okay; resolves to call "glorp() restrict(amp:2)"
}
```

## 13.2 Projected Evolution of `amp`-Restricted Code

Based on the nascent availability of features in advanced GPUs and corresponding hardware-vendor-specific programming models, it is apparent that the limitations associated with `restrict(amp)` will be gradually lifted. The table below captures one possible path for future `amp` versions to follow. If implementers need to (non-normatively) extend the `amp`-restricted language subset, it is recommended that they consult the table below and try to conform to its style.

Implementations may not define an `amp` version greater or equal to 2.0. All non-normative extensions shall be restricted to the patterns 1.x (where  $x > 0$ ). Version number 1.0 is reserved to implementations strictly adhering to this version of the specification, while version number 2.0 is reserved for the next major version of this specification.

Area	Feature	amp:1	amp:1.1	amp:1.2	amp:2	cpu
------	---------	-------	---------	---------	-------	-----

Local/Param/Function Return	char (8 - signed/unsigned/plain)	No	Yes	Yes	Yes	Yes
Local/Param/Function Return	short (16 - signed/unsigned)	No	Yes	Yes	Yes	Yes
Local/Param/Function Return	int (32 - signed/unsigned)	Yes	Yes	Yes	Yes	Yes
Local/Param/Function Return	long (32 - signed/unsigned)	Yes	Yes	Yes	Yes	Yes
Local/Param/Function Return	long long (64 - signed/unsigned)	No	No	Yes	Yes	Yes
Local/Param/Function Return	half-precision float (16)	No	No	No	No	No
Local/Param/Function Return	float (32)	Yes	Yes	Yes	Yes	Yes
Local/Param/Function Return	double (64)	Yes <sup>10</sup>	Yes	Yes	Yes	Yes
Local/Param/Function Return	long double (?)	No	No	No	No	Yes
Local/Param/Function Return	bool (8)	Yes	Yes	Yes	Yes	Yes
Local/Param/Function Return	wchar_t (16)	No	Yes	Yes	Yes	Yes
Local/Param/Function Return	Pointer (single-indirection)	Yes	Yes	Yes	Yes	Yes
Local/Param/Function Return	Pointer (multiple-indirection)	No	No	Yes	Yes	Yes
Local/Param/Function Return	Reference	Yes	Yes	Yes	Yes	Yes
Local/Param/Function Return	Reference to pointer	Yes	Yes	Yes	Yes	Yes
Local/Param/Function Return	Reference/pointer to function	No	No	Yes	Yes	Yes
Local/Param/Function Return	static local	No	No	Yes	Yes	Yes
Struct/class/union members	char (8 - signed/unsigned/plain)	No	Yes	Yes	Yes	Yes
Struct/class/union members	short (16 - signed/unsigned)	No	Yes	Yes	Yes	Yes
Struct/class/union members	int (32 - signed/unsigned)	Yes	Yes	Yes	Yes	Yes
Struct/class/union members	long (32 - signed/unsigned)	Yes	Yes	Yes	Yes	Yes
Struct/class/union members	long long (64 - signed/unsigned)	No	No	Yes	Yes	Yes
Struct/class/union members	half-precision float (16)	No	No	No	No	No
Struct/class/union members	float (32)	Yes	Yes	Yes	Yes	Yes
Struct/class/union members	double (64)	Yes	Yes	Yes	Yes	Yes
Struct/class/union members	long double (?)	No	No	No	No	Yes
Struct/class/union members	bool (8)	No	Yes	Yes	Yes	Yes
Struct/class/union members	wchar_t (16)	No	Yes	Yes	Yes	Yes
Struct/class/union members	Pointer	No	No	Yes	Yes	Yes
Struct/class/union members	Reference	No	No	Yes	Yes	Yes
Struct/class/union members	Reference/pointer to function	No	No	No	Yes	Yes
Struct/class/union members	bitfields	No	No	No	Yes	Yes
Struct/class/union members	unaligned members	No	No	No	No	Yes
Struct/class/union members	pointer-to-member (data)	No	No	Yes	Yes	Yes
Struct/class/union members	pointer-to-member (function)	No	No	Yes	Yes	Yes
Struct/class/union members	static data members	No	No	No	Yes	Yes
Struct/class/union members	static member functions	Yes	Yes	Yes	Yes	Yes
Struct/class/union members	non-static member functions	Yes	Yes	Yes	Yes	Yes
Struct/class/union members	Virtual member functions	No	No	Yes	Yes	Yes
Struct/class/union members	Constructors	Yes	Yes	Yes	Yes	Yes

<sup>10</sup> Double precision support is an optional feature on some amp:1-compliant hardware.

Struct/class/union members	Destructors	Yes	Yes	Yes	Yes	Yes
Enums	char (8 - signed/unsigned/plain)	No	Yes	Yes	Yes	Yes
Enums	short (16 - signed/unsigned)	No	Yes	Yes	Yes	Yes
Enums	int (32 - signed/unsigned)	Yes	Yes	Yes	Yes	Yes
Enums	long (32 - signed/unsigned)	Yes	Yes	Yes	Yes	Yes
Enums	long long (64 - signed/unsigned)	No	No	No	No	Yes
Structs/Classes	Non-virtual base classes	Yes	Yes	Yes	Yes	Yes
Structs/Classes	Virtual base classes	No	Yes	Yes	Yes	Yes
Arrays	of pointers	No	No	Yes	Yes	Yes
Arrays	of arrays	Yes	Yes	Yes	Yes	Yes
Declarations	__static	Yes	Yes	Yes	Yes	No
Function Declarators	Varargs (...)	No	No	No	No	Yes
Function Declarators	throw() specification	No	No	No	No	Yes
Statements	global variables	No	No	No	Yes	Yes
Statements	static class members	No	No	No	Yes	Yes
Statements	Lambda capture-by-reference (on gpu)	No	No	Yes	Yes	Yes
Statements	Lambda capture-by-reference (in p_f_e)	No	No	No	Yes	Yes
Statements	Recursive function call	No	No	Yes	Yes	Yes
Statements	conversion between pointer and integral	No	Yes	Yes	Yes	Yes
Statements	new	No	No	Yes	Yes	Yes
Statements	delete	No	No	Yes	Yes	Yes
Statements	dynamic_cast	No	No	No	No	Yes
Statements	typeid	No	No	No	No	Yes
Statements	goto	No	No	No	No	Yes
Statements	labels	No	No	No	No	Yes
Statements	asm	No	No	No	No	Yes
Statements	throw	No	No	No	No	Yes
Statements	try/catch	No	No	No	No	Yes
Statements	__try/__except	No	No	No	No	Yes
Statements	__leave	No	No	No	No	Yes

4299

4300